

# Representation as a Computational Resource in Quantum Compilation: Approximate Recoverability Tradeoffs and Model-Specific Fault-Tolerant Consequences

Jinze Yang,<sup>1</sup> Yangyang Li,<sup>2,\*</sup> and Xiu-Hao Deng<sup>3,4,†</sup>

<sup>1</sup>*School of Physics, Xidian University, Xi'an 710071, China*

<sup>2</sup>*Key Laboratory of Intelligent Perception and Image Understanding of Ministry of Education of China, Joint International Research Laboratory of Intelligent Perception and Computation,*

*International Research Center for Intelligent Perception and Computation, Collaborative Innovation Center of Quantum Information of Shaanxi Province, School of Artificial Intelligence, Xidian University, Xi'an 710071, China*

<sup>3</sup>*Shenzhen International Quantum Academy, Shenzhen 518048, China*

<sup>4</sup>*Shenzhen Branch, Hefei National Laboratory, Shenzhen, 518048, China*

(Dated: August 10, 2026)

Fault-tolerant compilation deliberately destroys structure. Time-dependent and variational schedules split one accumulated phase across many rounds, gate synthesis replaces an angle by a Clifford+ $T$  word in which no individual gate carries the angle at all, and randomized compiling twirls each rotation into sign-randomized fragments. We show that this dispersal has a price, and that the price is payable in exactly two currencies: bits of compiler state, or magic states. To make the price precise we formalize a downstream compiler—one positioned after randomization and lowering, as a cloud transpiler or a control-stack stream optimizer is—and charge every restart channel in bits: control, store, live window, parameters, and intermediate representation. For a commuting phase layer with  $m$  generators compiled to projective accuracy  $\epsilon$ , the physically realized stream is an  $r$ -round modular accumulation on the Ross-Selinger synthesis grid  $Q_\epsilon = K_\epsilon + 1 = \Theta(1/\epsilon)$ , whose exact packing has separation  $2\sin(\pi/(2Q_\epsilon))$ . We prove that any such compiler making  $p$  forward passes and failing with probability at most  $\delta$  obeys  $\bar{A}_p + (2p-1)S \geq (1-\delta)m \log_2 K_\epsilon - h_2(\delta)$ , where  $\bar{A}_p$  is committed output and  $S$  the serialized crossing state. The bound is carried by the schedule alone: sign normalization is a bijection of the residue alphabet, so the standard folded (compile-time) presentation of randomized compiling inherits it verbatim, and no claim depends on tracking frames. Under the additional disclosed frame-tracked presentation the state  $S$  is the Pauli frame itself, and the bound expresses a direct tension: the randomization that averages coherent noise is the same randomization that forces a re-aggregating compiler to carry  $\Omega(m \log(1/\epsilon))$  bits. A memory-capped block compiler attains the envelope within a constant factor; direct semantic aggregation attains the compact-output corner with logarithmic state.

We then measure the predicted channels rather than infer them. Instrumented reference compilers expose the state, IR, and commitment coordinates directly. Propagating one fixed total error budget through actual Clifford+ $T$  synthesis and surface-code scheduling turns the representational gap into a fault-tolerant one: the same unitary, compiled from a dispersed rather than an aggregate representation, costs orders of magnitude more logical  $T$  states and spacetime volume under the disclosed model. The two currencies are not, however, worth the same. Once a committed record is executed rather than buffered, the committed channel pays a toll the state channel never pays: the entropy  $\log_2 \binom{m}{k}$  of the compiler's own decision about which  $k$  coordinates to commit early. A compiler that defers is handed that information for free by the stream that follows. Under a preregistration frozen before the run, the toll is measured to match the predicted entropy, to be independent of  $m$ , and to vanish exactly when the commitment schedule is public. The operational conclusion is a design rule with a proof: randomize and lower late, because the information a compiler discards before aggregating cannot be recovered for free—it must be bought back in memory or in magic states, and buying it back early is strictly the more expensive of the two.

## I. INTRODUCTION

A fault-tolerant compiler is not only a translator; it is also, by design, a destroyer of structure. Time-dependent schedules and variational ansatz layers split a single accumulated phase across many rounds [1–4]. Gate synthesis dissolves an angle into a Clifford+ $T$  word whose individual letters carry none of it [5–8]. Randomized compiling


replaces each rotation by a sign-randomized fragment so that coherent errors average away [9]. Each of these transformations is correct, standard, and desirable for its own reason—and each of them takes a quantity that was once written down in one place and scatters it across an instruction stream.

This paper asks what it costs to put such a quantity back together, and shows that the cost is not a matter of engineering skill [10–12]. It is an information-theoretic quantity with two admissible payment channels. A compiler positioned downstream of the dispersal—a cloud transpiler receiving a flat program, or a stream optimizer

\* yyli@xidian.edu.cn

† dengxiuhao@iqasz.cn

(a) One QAOA cost layer, as a downstream compiler sees it



(b) Two currencies, at one cut




FIG. 1. Dispersal, and the two currencies that undo it. (a) A QAOA cost layer  $e^{-i\gamma C}$  on  $m$  edges is one commuting object carrying  $m$  records. Lowered under a digitized schedule it reaches a downstream compiler as  $r$  contiguous rounds of grid residues, optionally sign-randomized and separated by Clifford records; the edge total  $x_j$  is then a sum that no single record and no proper prefix reveals. The mixer layers bound the block rather than sitting inside it: at generic angles they are non-Clifford  $R_X$  rotations (remark 10). The controlled-power ladder beneath is the complementary case—its rounds are contiguous, but round one already fixes the aggregate  $\theta_j(2^r - 1)$ , so it supplies contiguity without dispersal (remark 7). (b) At any cut the scattered information is either carried across in compiler state or committed to output and repaid downstream in magic states. Theorem 5 states that the two cannot both be small, and proposition 7 charges the committed channel a further  $\log_2 \binom{m}{k}$  when the commitment schedule is private.

inside a control stack—must either carry the scattered information in internal state as it goes, or commit output before the information has arrived and pay later in fault-tolerant resources. There is no third option, and the exchange rate between the two is fixed by the accuracy  $\epsilon$  and the number of phase terms  $m$ .

A concrete instance fixes the scale (figure 1). Take a single QAOA [1] cost layer on a 3-regular MaxCut instance over 22 vertices, so  $m = 33$  edges, compiled to projective accuracy  $\epsilon = 10^{-3}$ . The synthesis grid at that tolerance has  $L_\epsilon = \lfloor \pi / (8 \arcsin 2\epsilon) \rfloor = 196$  and hence  $K_\epsilon = 197$  distinguishable settings per edge (proposition 3), so each edge total is worth  $\log_2 197 = 7.62$  bits and the layer as a whole is worth 252 bits, or about 31 bytes. A compiler that defers aggregation carries those 31 bytes across every cut. A compiler that commits early does not, and buys the same information back in the only other currency available:

each coordinate it must afterwards correct costs a fresh Ross–Selinger synthesis of roughly  $3 \log_2(1/\epsilon) \approx 30$   $T$  gates (remark 6), each of which is supplied by a distilled magic state rather than by the code itself [13, 14], so repairing the layer runs to order  $10^3$   $T$  gates. Thirty-one bytes against a thousand magic states is the exchange rate in miniature; section X measures where it ends, at a  $2814 \times$  spacetime-volume ratio between the aggregate and materialized paths under a common error budget.

The same picture shows where the statement stops, and it is worth saying early because a figure is easier to over-read than a theorem. Two hypotheses are required, and different families supply them. The rounds must be contiguous in the stream, and they must carry coefficients that vary freely from round to round. A variational cost layer supplies the second within one block of consecutive commuting rounds, the mixer closing the block rather than lying inside it; controlled-power ladders and phase estimation supply the first and not the second, since a ladder’s aggregate is already determined by its first round. Clifford records—routing, SWAP, and the Pauli-frame insertions of randomized compiling itself—are interior and normalize away (lemma 6). The boundaries are stated in full in remarks 8 and 10.

Underlying all three dispersal mechanisms is a single representation problem. A high-level circuit describing a commuting diagonal phase layer with  $m$  generators can be compiled in at least two qualitatively different ways. A *semantic-first* pipeline preserves the generator-coefficient structure, aggregates repeated coefficients, and lowers the aggregate representative once. A *materialize-first* pipeline expands the same layer into a flat target-basis instruction stream and then attempts to recover aggregate structure from local patterns. Both pipelines produce correct unitaries, but they expose very different resource profiles to downstream synthesis and fault-tolerant estimation.

This representation dependence is well known anecdotally. Phase-polynomial and Pauli-gadget methods [15–19] exploit algebraic structure before returning to a gate circuit. ZX-calculus methods [20–23] provide a global representation in which phases can propagate and merge beyond local syntax. Architecture-aware compilers [24, 25] and full-stack toolchains [11] make representation choices that propagate into physical resource estimates [12]. What is missing is a quantitative framework that isolates the effect of representation from other compilation variables and connects it to model-specific fault-tolerant cost.

We address this gap by formalizing *approximate recoverability*: given a target unitary  $U$  and a fixed approximation tolerance  $\epsilon$  in projective operator norm, what compiler resources are required to emit an output within  $\epsilon$  of  $U$  from a lowered representation? The contributions of this paper are:

1. An approximate recoverability framework with a budget-explicit compiler model charging a unified bit-level cut budget  $B_{\text{cut}} = B_{\text{control}} + B_{\text{store}} + B_{\text{window}} + B_{\text{parameter}} + B_{\text{IR}}$ , together with a self-delimiting growing-width token codec and explicit charges for

RAM/DAG/cache layout, dynamic passes, random seeds, output length, and approximation error (section III).

2. A cyclic  $K_\epsilon$ -ary packing with exact wrap margin,  $K_\epsilon = \Theta(1/\epsilon)$ , and an  $r$ -round modular update representation in which no token determines the aggregate. For randomized failure  $\delta$ , separate output-code bounds  $\bar{B}_{\text{rel}}^{\text{out}}, \bar{B}_{\text{self}}^{\text{out}} \geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$ , and the explicit forward-pass frontier  $\bar{A}_p + (2p - 1)S \geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$ . A memory-capped block compiler reaches compact relative output with  $S = O((m/p) \log(K_\epsilon + 1))$ , giving a constant-factor upper envelope; the older two-round  $K = 2, Q = 3$  result retains its stronger zero-error one-pass constant (sections IV to VI).
3. A refinement showing that the two payment channels are not interchangeable at par. When a committed record is executed rather than buffered, the committed channel additionally pays  $\log_2 \binom{m}{k}$  bits—the entropy of the compiler’s own decision about which  $k$  coordinates to commit early—while the state channel pays none, because the stream that follows discloses that decision for free. The toll vanishes exactly when the commitment schedule is public, is independent of  $m$ , and is measured on an instrumented device under a preregistration frozen before the run (proposition 7 and figure 3).
4. A representation-aware recovery pipeline with bounded recognition, aggregation certificates, and conservative fallback, evaluated on a complete nine-representation by five-compiler matrix and a separately versioned campaign of local, semantic-capable, and authors-reference baselines. Bridge loss and non-completion statuses are reported separately (sections VII and VIII).
5. Direct event-level measurement of committed output, all five cross-cut state fields, serialized IR, RSS, and time for echo, aggregation, memory-capped hybrid, and materialized-IR reference compilers. A fixed-total-error synthesis/QRE campaign reports logical  $T$ , data and factory qubits, cycles, factories, runtime, and spacetime volume under two versioned QEC models (sections VIII D and X).
6. A natural-workload campaign covering structured algorithm families, two negative controls, and the complete six-factor ablation, together with a normalized frozen-data registry, numeric provenance map, and independent-directory reconstruction of theorem, benchmark, and QRE subsets (sections IX and XIII).

The prior version of this work established a representation-dependent recoverability boundary for exact semantics on fixed-width witness families. Those results—the final-output lower bound for constant-budget finite-window transducers, the four-qubit irrational-angle

independent-Pauli witness, and the rational CP witness with period  $T = 131040$ —are retained here as an exact-semantics corollary (corollary 1). The primary asymptotic axes are now the number of independent phase terms  $m$  and the requested approximation tolerance  $\epsilon$ ; the output lower bound scales as  $m \log(1/\epsilon)$ .

## II. OPERATIONAL SETUP

We define the operational pipeline in four stages: *semantic input*, *basis lowering*, *recovery*, and *synthesis/resource estimation*.

**Definition 1** (Semantic input). *A semantic input is a circuit description  $C_{\text{sem}}$  in a representation that preserves generator-coefficient structure. For a commuting diagonal phase layer with  $m$  independent generators  $P_1, \dots, P_m$  and coefficients  $\theta_1, \dots, \theta_m$ , the semantic input has  $O(m)$  records of the form  $(P_j, \theta_j)$ .*

**Definition 2** (Basis lowering). *Let  $\mathcal{L}_B(\cdot)$  denote exact lowering into a fixed target basis  $B$ . The lowered circuit  $\mathcal{L}_B(C_{\text{sem}})$  is a flat instruction stream over  $B$  implementing the same unitary up to global phase. Lowering preserves unitary semantics but need not preserve the generator-coefficient structure.*

**Definition 3** (Recovery). *A recovery procedure  $\mathcal{R}$  takes a lowered circuit  $\mathcal{L}_B(C_{\text{sem}})$  and produces an output circuit  $C_{\text{out}}$  such that  $d_{\text{proj}}(U(C_{\text{out}}), U(C_{\text{sem}})) \leq \epsilon$  under the projective operator norm  $d_{\text{proj}}$ . The recovery is representation-aware if it reconstructs aggregate structure; it is materialize-first if it operates only on local patterns.*

**Definition 4** (Synthesis and resource estimation). *A fault-tolerant synthesis procedure maps  $C_{\text{out}}$  to a discrete-gate circuit over a fault-tolerant gate set (e.g., Clifford+T) subject to a total error budget  $\epsilon_{\text{total}}$ . A quantum resource estimator (QRE) then reports logical  $T$  count, magic-state demand, physical qubits, cycles, and spacetime volume under specified architecture assumptions.*

The central question is: for a fixed target unitary  $U$  and fixed  $\epsilon$ , how do the compiler resources required by  $\mathcal{R}$  depend on whether  $\mathcal{R}$  receives  $C_{\text{sem}}$  or  $\mathcal{L}_B(C_{\text{sem}})$ ?

## III. APPROXIMATE RECOVERABILITY FRAMEWORK

### A. Approximate Distance and Recoverability

**Definition 5** (Projective operator norm distance). *For unitaries  $U, V$  on  $n$  qubits,*

$$d_{\text{proj}}(U, V) = \inf_{\phi \in \mathbb{R}} \|U - e^{i\phi} V\|_{\text{op}}. \quad (1)$$

*The infimum is attained because the circle  $\{e^{i\phi}\}$  is compact and  $\|U - e^{i\phi} V\|_{\text{op}}$  is continuous in  $\phi$ .*

The projective distance directly eliminates global phase and is used as the primary correctness criterion below, (1).

**Lemma 1** (Unitary-channel diamond comparison). *For the unitary channels  $\mathcal{U}(X) = UXU^\dagger$  and  $\mathcal{V}(X) = V XV^\dagger$ , define the unnormalized diamond distance  $d_\diamond(\mathcal{U}, \mathcal{V}) := \|\mathcal{U} - \mathcal{V}\|_\diamond \in [0, 2]$ . Then*

$$d_{\text{proj}}(U, V) \leq d_\diamond(\mathcal{U}, \mathcal{V}) \leq 2d_{\text{proj}}(U, V).$$

If the normalized convention  $\tilde{d}_\diamond := \frac{1}{2}\|\mathcal{U} - \mathcal{V}\|_\diamond$  is used, the corresponding explicit relation is  $\frac{1}{2}d_{\text{proj}} \leq \tilde{d}_\diamond \leq d_{\text{proj}}$ .

*Proof.* Let  $\Theta \in [0, 2\pi]$  be the length of a shortest closed arc containing the spectrum of  $U^\dagger V$ . Centering that arc by a global phase gives

$$d_{\text{proj}}(U, V) = 2 \sin(\Theta/4).$$

The numerical-range formula for two unitary channels gives

$$d_\diamond(\mathcal{U}, \mathcal{V}) = \begin{cases} 2 \sin(\Theta/2), & 0 \leq \Theta \leq \pi, \\ 2, & \pi \leq \Theta \leq 2\pi. \end{cases}$$

For  $\Theta \leq \pi$ , the ratio is  $2 \cos(\Theta/4) \in [\sqrt{2}, 2]$ ; for  $\Theta \geq \pi$ , one has  $d_{\text{proj}} \in [\sqrt{2}, 2]$  and  $d_\diamond = 2$ . The two inequalities follow.  $\square$

Consequently, a projective packing with separation  $s$  is also a diamond packing with separation at least  $s$ . Every lower bound below therefore holds with the *same numerical accuracy threshold* when correctness is instead required as  $d_\diamond(\mathcal{U}(C_{\text{out}}), \mathcal{U}) \leq \epsilon$ . Under the normalized convention, the same argument uses packing separation  $s/2$  and correctness threshold  $\epsilon/2$ . Conversely, a projective-error upper bound  $\epsilon$  implies diamond error at most  $2\epsilon$ ; the exact constructive upper bounds used here have zero error in either metric. No diamond packing conclusion is inferred from the upper inequality alone.

**Definition 6** ( $\epsilon$ -approximate recoverability). *Fix  $\epsilon > 0$ . Given a target unitary  $U$  presented as a lowered circuit  $\mathcal{L}_B(C_{\text{sem}})$ , the  $\epsilon$ -approximate recoverability problem asks for an output  $C_{\text{out}}$  with  $d_{\text{proj}}(U(C_{\text{out}}), U) \leq \epsilon$ , using compiler resources charged by the model below.*

## B. Budget-Explicit Compiler Model

We retain the fixed-width finite-window model only for the exact corollary and use the following more general randomized, random-access model for the growing-width approximate results.

**Definition 7** (Sparse formal-parameter serialization). *Fix formal angle symbols  $\theta_1, \dots, \theta_M$ . A parameter is*

*stored in normalized sparse form*

$$\lambda = \sum_{\ell=1}^s c_\ell \theta_{j_\ell} + \frac{a}{d} \pi, \\ j_1 < \dots < j_s, \quad c_\ell, a \in \mathbb{Z}, \quad d \geq 1,$$

*with zero coefficients omitted and  $\gcd(a, d) = 1$ . Let  $\mathbf{u}(z)$  and  $\mathbf{s}(z)$  be fixed self-delimiting codes for nonnegative and signed integers. The executable parameter serialization is*

$$\begin{aligned} \text{ser}_{\text{par}}(\lambda) &= \mathbf{u}(s) \prod_{\ell=1}^s (\mathbf{u}(j_\ell) \mathbf{s}(c_\ell)) \\ &\quad \cdot \mathbf{s}(a) \mathbf{u}(d), \\ \text{prec}(\lambda) &:= |\text{ser}_{\text{par}}(\lambda)|. \end{aligned} \tag{2}$$

*Thus one nonzero formal coefficient costs  $O(\log M)$  rather than  $\Omega(M)$  bits. Qubit and generator locations are not part of this parameter payload; they are serialized by the complete-token codec below.*

**Definition 8** (Width-aware complete-token codec). *Fix a self-delimiting integer code  $\mathbf{u}$ , a signed variant  $\mathbf{s}$ , a two-bit Pauli code  $\mathbf{p}$ , and a reserved record terminator  $\mathbf{d}$ . A token  $\tau$  with gate-type ID  $g$ , arity  $r$ , ordered operand qubits  $q_1, \dots, q_r$ , sparse Pauli support  $\{(z_\ell, \sigma_\ell)\}_{\ell=1}^s$ , and  $k$  parameter slots is encoded as*

$$\begin{aligned} \text{ser}_{\text{tok}}(\tau) &= \mathbf{u}(g) \mathbf{u}(r) \prod_{i=1}^r \mathbf{u}(q_i) \mathbf{u}(s) \prod_{\ell=1}^s (\mathbf{u}(z_\ell) \mathbf{p}(\sigma_\ell)) \\ &\quad \cdot \mathbf{u}(k) \prod_{a=1}^k (\mathbf{u}(\text{id}_a) \text{ser}_{\text{par}}(\lambda_a)) \mathbf{d}. \end{aligned}$$

*The zero values of  $r, s$ , or  $k$  are explicitly encoded, so the record is executable and uniquely parseable without out-of-band arity, support, or delimiter information. Update, IR-node, and output records use the same field order, with an opcode specifying which fields are active.*

*At growing width  $n$ , an arbitrary qubit or generator choice alone requires  $\lceil \log_2 n \rceil$  bits in any injective codec; the concrete codec above uses  $O(\log n)$  bits per listed address and  $O(s \log n)$  bits for an  $s$ -sparse Pauli support. Consequently there is no width-independent “structural cost per token.” The literal bit strings  $\text{ser}_{\text{tok}}$  are charged wherever a token is live or emitted. Only when  $n = n_0$  and a complete-token precision bound  $q$  are both fixed is the token set  $\Gamma_{B, n_0}^{\leq q}$  finite; the fixed-width exact-semantics statements abbreviate it as  $\Gamma_B(q)$ .*

**Definition 9** (Randomized dynamic-pass compiler and serialized cut budget). *A compiler  $\mathcal{A}$  may use random-access RAM, stacks, hash tables, caches, DAG/graph IRs, and intermediate tapes; it may also select a finite number  $p = p(I, \rho)$  of scans or passes dynamically from the input  $I$  and a finite random seed/tape  $\rho$ . The external input is append-only on its first exposure. Any consumed prefix that remains available for a later scan—whether through*

a file, memory mapping, spool, or cache—is live compiler storage and is charged below. The final output is write-once. At every execution cut  $c$ , serialize a complete restart snapshot in five disjoint fields:

- $\text{ser}_{\text{control}}(c)$ : format version, pass and finite-control IDs, dynamic pass count and termination/scheduling state, all input/output heads and emitted-prefix lengths, end flags, framing lengths, and the random seed/tape with its cursor;
- $\text{ser}_{\text{store}}(c)$ : the persistent and ephemeral RAM, stacks, hash-table keys, bucket layout, cache contents and metadata, with parameter payloads and IR objects replaced by location tags;
- $\text{ser}_{\text{window}}(c)$ : the ordered live read-window token structures, including the structural fields of definition 8 but excluding parameter payloads;
- $\text{ser}_{\text{parameter}}(c)$ : the location-tagged concatenation of every live parameter payload in control, store, window, and IR, each encoded by (2);
- $\text{ser}_{\text{IR}}(c)$ : every other live intermediate tape/DAG/table object and every rereadable consumed-input block that can affect future execution, serialized by node opcode, the structural fields of definition 8, edge endpoints, multiplicity, and object boundaries, but with parameter payloads omitted.

The corresponding bit-level budget is

$$B_{\text{cut}}(c) = B_{\text{control}}(c) + B_{\text{store}}(c) + B_{\text{window}}(c) + B_{\text{parameter}}(c) + B_{\text{IR}}(c). \quad (3)$$

where each summand is the literal bit length of the named serialization. For a fixed run,  $B_{\text{cut}}$  is the maximum over its cuts. For randomized compilers we use the worst-input expected budget

$$\overline{B}_{\text{cut}}(\mathcal{A}, \mathcal{I}) := \max_{I \in \mathcal{I}} \mathbb{E}_{\rho} [\max_c B_{\text{cut}}(c; I, \rho)]. \quad (4)$$

The previously named quantities are now unambiguous bit measures:  $Q := \max_c B_{\text{parameter}}(c)$  and  $V := \max_c B_{\text{IR}}(c)$ ; no theorem below substitutes either one for the total  $B_{\text{cut}}$ . The final output is write-once and is charged separately by the output codes below.

**Remark 1** (Model scope and escape channels). Equation (3) charges all restart information, not only the finite-control state. In particular, a ZX graph, phase-polynomial table, spilled update log, hash table, cache, random seed, or materialized intermediate tape is serialized in the indicated field, and all coefficients are serialized in  $B_{\text{parameter}}$ . Dynamic rescanning is allowed, but the pass schedule, heads, and any rereadable old input are part of the snapshot. The five fields are disjoint by construction, so parameter bits are not counted again as store, window, or IR structure. A compiler may escape a RAM-only workspace lower bound by materializing a large IR, but it then pays the same bits in  $B_{\text{cut}}$ .

**Definition 10** (Three noninterchangeable output measures). Fix two prefix-free, streaming binary codes. The self-contained code  $E_{\text{self}}$  serializes a target-basis circuit using gate opcodes, all qubit/generator addresses, record delimiters, and sparse parameters using definition 8; it assumes no family dictionary. For a public packing-family dictionary  $\mathcal{F}_m = (m, P_1, \dots, P_m, K_{\text{pack}}, \Delta, B)$ , the conditional code  $E_{\text{rel}}(\cdot | \mathcal{F}_m)$  has (i) an aggregate mode consisting of a self-delimiting header and a packed base- $K_{\text{pack}}$  block of  $m$  digits, (ii) a literal mode containing complete target-basis records, and (iii) an ordered-share mode whose header declares a split point and whose fixed public schedule permits generator addresses to be omitted from its residue slots. Ordered-share slots still include their round boundary and packed-block delimiters, and (iv) a  $p$ -block hybrid mode whose header declares the pass blocks and passthrough coordinates and whose data section contains packed aggregate blocks and round-major residue blocks in that public order. Define

$$\begin{aligned} B_{\text{rel}}(C | \mathcal{F}_m) &:= |E_{\text{rel}}(C | \mathcal{F}_m)|, \\ B_{\text{self}}(C) &:= |E_{\text{self}}(C)|, \\ G(C) &:= \text{number of target-basis gates in } C. \end{aligned} \quad (5)$$

For a compiler, a superscript “out” and a maximum over valid inputs denote the worst-case output value of the corresponding measure; a barred superscript “out” denotes the worst-input expectation over  $p$ . These quantities are not converted into one another: conditional description bits, self-contained circuit bits, and gate count are reported and bounded separately.

### C. Executable Model–Implementation Correspondence

The formal fields above are implemented by one canonical binary codec rather than by dimension-changing estimates. The executable specification is `encoding/FORMAT.md`; its reference implementation is `encoding/codec.py`, and the only bit-accounting entry point is `instrumentation/resource_accounting.py`. A complete codec frame has the form

$$\text{UCCACCT0} \parallel \text{u(version)} \parallel \text{u(type)} \parallel \text{u}(|z|) \parallel z, \quad (6)$$

where the displayed terminal zero denotes a null byte,  $\text{u}$  is the canonical base-128 unsigned varint, and  $z$  is a typed canonical payload. Signed integers use zigzag coding; rational coefficients are reduced signed numerator/positive-denominator pairs. Maps are ordered by their UTF-8 key bytes. Qubit, generator, symbol, location, node, edge-endpoint, port, pass, and head addresses are literal varints in  $z$ .

**Lemma 2** (Canonical prefix-free implementation). Version-one complete frames are injective and prefix-free. Concatenations of frames are uniquely parseable without an out-of-band arity or record count. The semantic and

| Formal object                        | Codec object/file                                         | Included information                                                                                            | Measured field            |
|--------------------------------------|-----------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|---------------------------|
| semantic input                       | <code>SemanticStream</code>                               | width and address-complete aggregate records                                                                    | input artifact bits       |
| dispersed flat input                 | <code>FlatUpdateStream</code>                             | width, round count, and every additive up-date                                                                  | input artifact bits       |
| <code>ser<sub>control</sub></code>   | <code>ControlState/control.uccbin</code>                  | pass count/index, heads, scheduling, status, framing and seed state when present                                | $B_{\text{control}}$      |
| <code>ser<sub>store</sub></code>     | <code>StoreState/store.uccbin</code>                      | aggregate tables, maps, caches, stacks and their layout, excluding tagged parameter/IR payloads                 | $B_{\text{store}}$        |
| <code>ser<sub>window</sub></code>    | <code>LiveWindow/window.uccbin</code>                     | ordered live complete-token structure and explicit parameter-slot count                                         | $B_{\text{window}}$       |
| <code>ser<sub>parameter</sub></code> | <code>ParameterLedger/parameter.uccbin</code>             | sparse symbolic parameters under complete location tuples                                                       | $B_{\text{parameter}}$    |
| <code>ser<sub>IR</sub></code>        | <code>IRField/ir.uccbin</code>                            | serialized DAG width, opcodes, operands, nodes, edges, ports, boundaries and meta-data, with parameters removed | $B_{\text{IR}}$           |
| committed prefix                     | <code>committed-prefix.bin</code>                         | literal immutable prefix of a codec output frame                                                                | $B_{\text{com}}$          |
| external output                      | <code>ExternalToolOutput</code>                           | tool, format and literal API payload                                                                            | output artifact bits      |
| certificate input                    | <code>CertificateInput</code>                             | certificate kind, certified-artifact SHA-256 and literal sidecar                                                | certificate artifact bits |
| pass count                           | <code>manifest passes.count</code> and row <code>p</code> | actual declared scans of the instrumented compiler                                                              | $p$                       |

TABLE I. One-to-one map from formal serialized objects to executable artifacts. No row permits a gate, token, support, or IR-node count to stand in for bits.

*flat decoders preserve their exact aggregate coefficient table.*

*Proof.* Canonical varints have a unique terminating byte, and the version, type, and payload length are therefore uniquely decoded. The declared payload length locates the end of the frame. A complete valid frame cannot be a proper prefix of another complete valid frame: the shorter declared payload ends at its frame boundary, whereas any following byte belongs to a subsequent frame. Typed payload values have distinct one-byte tags; lists and maps carry counts, maps have strictly increasing keys, and reduced rational pairs are unique. This proves injectivity and unique concatenated parsing. Finally, `semantic_signature` sums exact rational sparse parameters by the literal pair (generator address, Pauli support), so encode-decode leaves the aggregate table unchanged for both input representations.  $\square$

**Proposition 1** (File-backed accounting soundness). *Suppose the version-one verifier accepts a cut manifest  $M$ . For every  $f \in \{\text{control, store, window, parameter, IR, com}\}$ , let  $A_f$  be the file named by  $M$ . Then the implementation reports*

$$\begin{aligned}
 B_f &= 8|A_f|_{\text{byte}}, \\
 B_{\text{cut}} &= B_{\text{control}} + B_{\text{store}} + B_{\text{window}} \\
 &\quad + B_{\text{parameter}} + B_{\text{IR}},
 \end{aligned} \tag{7}$$

*and reports the manifest pass count as  $p$ . The reported values change when a serialized address crosses a varint-width boundary. They cannot be obtained from a node or gate count alone.*

*Proof.* The accountant first writes the five complete frames and the literal committed prefix, then obtains every byte count by `stat`, stores the SHA-256 of each

file, writes  $M$ , and invokes the verifier before returning a row. The verifier reopens all six files, recomputes sizes and digests, decodes the five restart fields, and rejects any disagreement. It then forms the sum in (7); there is no count-to-bit code path. Since an address is a literal canonical varint, for example 127 and 128 occupy different address lengths. Conversely, two one-node DAGs may contain those different addresses and hence have unequal file lengths despite equal node count.  $\square$

**Remark 2** (Observable boundary versus restart cut). *The masked-share tracer writes the five restart fields after every update, destroys the live Python state, and resumes from only those files, the committed prefix, and unread input; its maximum is therefore an implementation of  $B_{\text{cut}}$ . External compilers do not expose their private in-process DAGs, caches, or tapes. Their matched-matrix and natural-factorial rows consequently report a file-backed API-boundary snapshot and are diagnostic only; they are not promoted to measurements of the maximum private  $B_{\text{cut}}$ . Peak RSS remains a separate byte dimension and is not substituted for any field in (3).*

**Remark 3** (Self-contained and conditional descriptions). *Both input-stream objects declare either **self\_contained** or **conditional**. The former forbids an external dictionary. The latter requires the SHA-256 of the exact public family dictionary, and its manifest records the dictionary's separate file size and whether it is charged at the cut. Thus a conditional description cannot silently be reported as a self-contained circuit code.*

**Remark 4** (Physical scope of the bounded-memory streaming model). *The charged streaming model of definition 9—bounded live state, explicit pass accounting, append-only first exposure—is the operating regime of a real-time fault-tolerant control stack, not a convenience*

*abstraction. Between rounds, the classical electronics adjacent to the quantum device holds decoder state, Pauli-frame records, and instruction buffers under hard latency and buffer budgets: corrections must be resolved before the next non-Clifford operation, and frame tracking exists precisely because unbounded buffering is unavailable at that boundary [11, 26]. Distributed and modular architectures motivate the same scope from the interconnect side: a cross-module interaction splits the support of a gate across modules while its coefficient stays within one module, so modularity bounds the state that crosses a partition without by itself dispersing any coefficient. These architectures therefore ground the model's memory and pass constraints only; no claim is made that modular partitioning instantiates the dispersed streams of section IV C, whose instances come from scheduling and twirling alone.*

#### D. Exact-Semantics Corollary from Prior Work

We record the exact-semantics ( $\epsilon = 0$ ) results as corollaries; the full transducer proofs are restated in section A 3. The witness family is fixed-width.

**Remark 5** (Separate fixed-width and growing-width encoding regimes). *The exact corollary in this subsection fixes  $n = n_0$ , a precision bound  $q$ , a deterministic constant pass count  $p_0$ , and one finite alphabet  $\Gamma_{B,n_0}^{\leq q}$ ; its unit is therefore a token over that fixed alphabet. No such constant-cost token assumption is made in the growing-width results. There  $n = m+1$  (or  $n \geq m$ ), the pass count may be  $p = p(m, I, \rho)$ , and every input, output, RAM, and IR record is charged in bits by definition 8 and (3).*

**Definition 11** (Fixed-width witness family). *Let  $H_\Lambda$  be a fixed invertible boundary layer (a Hadamard layer in the experiments) and let  $D = \prod_{j=1}^{m_0} \exp(-i\theta_j P_j/2)$  be a fixed commuting diagonal layer with Pauli-Z characters  $P_j = Z^{a_j}$ ,  $a_j \in \mathbb{F}_2^n$ , and angles  $\theta_j$ . The witness at repetition count  $r$  is  $C_r = H_\Lambda D^r H_\Lambda^{-1}$ , with exact lowering  $\mathcal{L}_B(C_r) = \Pi M^r \Sigma$  for fixed token strings  $\Pi, M, \Sigma$ .*

Two instances are certified and reused here. The *rational CP witness* ( $n = 4$ ; ten  $rz/cp$  terms;  $\text{rank}_{\mathbb{F}_2}(A) = 4$ ) satisfies finite-range noncollision with exact period  $T = 131040$ , certified by the self-contained rational angle, single-term order, relative-phase, and LCM derivation in section A 2. The *irrational independent-Pauli witness* ( $n = 4$ ; supports 1100, 0110, 0011, 1000;  $\theta_1/\pi = \sqrt{2}/11$ ) satisfies unbounded noncollision by the following uniqueness lemma.

**Lemma 3** (Coefficient uniqueness for independent characters). *If  $a_1, \dots, a_{m_0} \in \mathbb{F}_2^n$  are linearly independent and two diagonal layers with coefficient vectors  $\alpha, \alpha'$  satisfy  $D(\alpha) = e^{i\varphi} D(\alpha')$ , then  $\alpha_j \equiv \alpha'_j \pmod{2\pi}$  for every  $j$ . In particular, if some  $\theta_j/\pi$  is irrational, the powers  $D^r$  are pairwise distinct up to global phase for all  $r \geq 1$ .*

*Proof.* Independence makes  $z \mapsto (\chi_{a_1}(z), \dots, \chi_{a_{m_0}}(z))$  surjective onto  $\{\pm 1\}^{m_0}$ . Comparing the phase constraint at two sign patterns that differ only in coordinate  $j$  gives  $2(\alpha_j - \alpha'_j) \equiv 0 \pmod{4\pi}$ . For the second claim,  $(r - r')\theta_j \in 2\pi\mathbb{Z}$  with  $\theta_j/\pi$  irrational forces  $r = r'$ .  $\square$

**Assumption 1** (Finite-range noncollision). *On a range  $1 \leq r \leq R$ , the powers  $D^1, \dots, D^R$  are pairwise distinct up to global phase.*

**Assumption 2** (Unbounded noncollision). *The powers  $D^r$ ,  $r \geq 1$ , are pairwise distinct up to global phase.*

**Corollary 1** (Exact-semantics final-output lower bound). *Assume exact lowering ( $\epsilon = 0$ ), a fixed-width witness  $C_r = H_\Lambda D^r H_\Lambda^{-1}$  with  $\mathcal{L}_B(C_r) = \Pi M^r \Sigma$ , and unbounded noncollision (assumption 2). Let  $\mathcal{A}$  be a uniform deterministic finite-window, finite-burst transducer with constant budget (constant  $p, w, \mu$  and  $B_{\text{cut}}$ ) that is correct for every  $r \geq 1$ . Then there exist constants  $h \geq 1$ ,  $D \geq 1$  (depending on  $\mathcal{A}$  and the witness but not  $r$ ) such that for every  $r \geq h + D$ ,*

$$|T_p(r)| \geq \frac{r - h - D + 1}{D} = \Omega_{\mathcal{A}}(r). \quad (8)$$

*Proof.* This is theorem 6, proved in full in section A 3. Constant budget gives one finite charged alphabet; per-tape pumpability propagation gives pumpability of the final output; noncollision prevents collapse of distinct pump residues.  $\square$

**Corollary 2** (Finite-range tail form). *Under finite-range noncollision (assumption 1) with  $N(R) = R$  on  $1 \leq r \leq R$ ,  $|T_p(r)| \geq (r - h - D + 1)/D$  for  $h + D \leq r \leq R$ . For the rational CP witness, this applies only for  $R < T = 131040$ .*

**Corollary 3** (Counting toll for short outputs). *If finite-range noncollision supplies  $N(R)$  distinct targets, then every deterministic exact compiler satisfies*

$$\max_{1 \leq r \leq R} B_{\text{self}}(T_p(r)) \geq \lceil \log_2 N(R) \rceil.$$

*The same statement holds for  $B_{\text{rel}}(\cdot | \mathcal{F})$  when one fixed public dictionary  $\mathcal{F}$  is used on the entire range.*

The irrational independent-Pauli witness instantiates corollary 1 for all  $r$  via lemma 3; the rational CP witness instantiates corollary 2 on ranges  $R < T$ . These exact-semantics results concern  $\epsilon = 0$  and fixed width  $n = 4$ ; the results below generalize to growing  $m$  and expose the dependence on  $\epsilon > 0$ .

## IV. PACKING CONSTRUCTION

### A. Growing-Width Family

**Definition 12** (Packing family). *For each  $m$  and accuracy scale  $\epsilon$ , a packing family  $\mathcal{F}_{m,\epsilon} = \{U_x : x \in \mathcal{X}_{m,\epsilon}\}$  has*

pairwise projective separation greater than  $2\epsilon$ , a semantic description of  $O(m)$  addressed records, and a lowered streaming description whose aggregate information may be dispersed across updates. The accuracy-dependent construction below has  $|\mathcal{X}_{m,\epsilon}| = K_\epsilon^m$  with  $K_\epsilon = \Theta(1/\epsilon)$ .

Every growing-width family in this paper—the codebook that carries the separation, the streams that carry the lower bounds, the two-round device that is executed in section VIII D, and the schedules of section IV C—is one scheme under four parameters. We state it once and list the specializations in table II.

**Definition 13** (Phase-accumulation family). *Let  $P_1, \dots, P_m$  be  $\mathbb{F}_2$ -independent Pauli-Z characters on  $n \geq m$  qubits (i.e.  $P_j = Z^{a_j}$  with  $\text{rank}_{\mathbb{F}_2}(a_1, \dots, a_m) = m$ ), and write  $R_{P_j}(\theta) = \exp(-i\theta P_j/2)$  and  $[K]_0 = \{0, \dots, K-1\}$ . A family is fixed by a step  $\Delta > 0$ , an aggregate alphabet  $[K]_0$ , a number of rounds  $r \geq 1$ , and a sign set  $\Sigma \subseteq \{\pm 1\}$ . The semantic input is the addressed vector  $x \in [K]_0^m$ , realizing*

$$U_x = \exp\left(-\frac{i\Delta}{2} \sum_j x_j P_j\right) = \prod_{j=1}^m R_{P_j}(\Delta x_j). \quad (9)$$

The flat input is the round-major record sequence  $\text{Update}(t, j, a_j^{(t)}, s_j^{(t)})$  with residues  $a_j^{(t)}$  and signs  $s_j^{(t)} \in \Sigma$ , realizing

$$W = \prod_{t=1}^r \prod_{j=1}^m R_{P_j}(\Delta s_j^{(t)} a_j^{(t)}), \quad (10)$$

whose aggregate is  $x_j = \sum_t s_j^{(t)} a_j^{(t)}$ . Two arithmetic regimes are used, and no result mixes them:

- no-wrap: residues and aggregate live in  $\mathbb{Z}$  and  $(K-1)\Delta \leq \pi$ , so distinct aggregates are distinct angles;
- cyclic:  $\Delta = 2\pi/Q$  for a modulus  $Q \geq 3$ , residues lie in  $\mathbb{Z}_Q$ , the aggregate is taken mod  $Q$ , and  $K = Q-1$  unless stated otherwise.

Since the  $P_j$  commute,  $W = U_x$  up to global phase in either regime (proposition 2).

**Proposition 2** (Flat inputs aggregate projectively). *Every valid flat input of definition 13 implements  $U_x$  up to global phase, with no promise about any individual round.*

*Proof.* The  $P_j$  commute, so  $W = \prod_j R_{P_j}(\Delta \sum_t s_j^{(t)} a_j^{(t)})$ . In the no-wrap regime the exponent is  $\Delta x_j$  by definition. In the cyclic regime write  $\sum_t s_j^{(t)} a_j^{(t)} = x_j + Qk_j$  for integers  $k_j$ ; since  $Q\Delta = 2\pi$ ,  $R_{P_j}(\Delta x_j + 2\pi k_j) = (-1)^{k_j} R_{P_j}(\Delta x_j)$ , and multiplying over  $j$  changes only a global sign.  $\square$

**Definition 14** ( $K$ -ary no-wrap construction). *The no-wrap specialization of definition 13 with  $r = 1$ ,  $\Sigma = \{+1\}$ , integer  $K \geq 2$ , and  $(K-1)\Delta \leq \pi$ . Its codebook is  $\{U_x : x \in [K]_0^m\}$  as in (9).*

| Def. | Specialization   | Regime        | $r$      | $\Sigma$  | Role         |
|------|------------------|---------------|----------|-----------|--------------|
| 14   | no-wrap packing  | no-wrap       | 1        | $\{+\}$   | codebook     |
| 20   | dispersed update | cyclic        | $\geq 2$ | $\{+\}$   | lower bounds |
| 16   | masked share     | cyclic, $Q=3$ | 2        | $\{+\}$   | executed     |
| 17   | scheduled stream | cyclic        | $\geq 2$ | $\{\pm\}$ | physical     |
| 15   | balanced         | no-wrap       | $r$      | $\{+\}$   | diagnostic   |

TABLE II. The growing-width families are specializations of definition 13. The balanced encoding is listed for completeness: its tokens contain  $x_j/r$ , so it proves no lower bound (remark 7). The fixed-width witness family of definition 11 belongs to the separate exact-semantics regime of remark 5 and is not an instance of this scheme.

**Lemma 4** (Packing separation). *For the construction in definition 14, all  $x \neq y \in [K]_0^m$  satisfy*

$$d_{\text{proj}}(U_x, U_y) \geq 2 \sin(\Delta/4). \quad (11)$$

The constant is tight: pairs differing by one in a single coordinate attain it.

*Proof.* Let  $s = x - y$  and choose  $j^*$  with  $s_{j^*} \neq 0$ . Since  $P_j$  commute,  $U_x U_y^\dagger = \exp(-\frac{i\Delta}{2} \sum_j s_j P_j)$  is diagonal in the computational basis. Its eigenvalue at basis state  $z$  is  $\lambda(z) = \exp(-\frac{i\Delta}{2} \sum_j s_j \chi_{P_j}(z))$ , where  $\chi_{P_j}(z) = (-1)^{a_j \cdot z} \in \{\pm 1\}$ .

Because  $a_1, \dots, a_m$  are  $\mathbb{F}_2$ -independent, the map  $z \mapsto (\chi_{P_1}(z), \dots, \chi_{P_m}(z))$  is surjective onto  $\{\pm 1\}^m$ . Hence two basis states can be chosen whose character vectors agree everywhere except at  $j^*$ . Their eigenvalue arguments differ by  $\Delta s_{j^*}$  modulo  $2\pi$ . The no-wrap condition gives

$$\Delta \leq |\Delta s_{j^*}| \leq (K-1)\Delta \leq \pi.$$

Now include the global-phase infimum required by definition 5. For every  $\phi \in \mathbb{R}$ ,

$$\|U_x - e^{i\phi} U_y\|_{\text{op}} = \max_z |\lambda(z) e^{-i\phi} - 1| \geq 2 \sin(\Delta/4), \quad (12)$$

because two points on the unit circle whose arguments differ by  $|\Delta s_{j^*}| \in [\Delta, \pi]$  cannot both lie within circular distance less than  $\Delta/2$  of 1 after any common rotation, and  $|e^{i\gamma} - 1| = 2|\sin(\gamma/2)|$ . Taking the infimum over  $\phi$  preserves the bound. A one-coordinate difference of one has exactly two spectral values separated by  $\Delta$ , and the centered choice of  $\phi$  attains the lower bound.  $\square$

**Proposition 3** ( $\epsilon$ -tuned  $K$ -ary packing). *For  $0 < \epsilon \leq 1/8$ , set*

$$L_\epsilon = \left\lfloor \frac{\pi}{8 \arcsin(2\epsilon)} \right\rfloor, \quad M_\epsilon = 4L_\epsilon, \quad (13)$$

$$K_\epsilon = L_\epsilon + 1, \quad \Delta_\epsilon = \frac{2\pi}{M_\epsilon}.$$

Then  $K_\epsilon = \Theta(1/\epsilon)$ ,  $(K_\epsilon - 1)\Delta_\epsilon = \pi/2$  (so no coefficient wraps around  $2\pi$ ), and the  $K_\epsilon^m$  targets in definition 14 obey

$$d_{\text{proj}}(U_x, U_y) \geq 2 \sin(\Delta_\epsilon/4) \geq 4\epsilon \quad (x \neq y). \quad (14)$$



*Proof.* The restriction  $\epsilon \leq 1/8$  gives  $L_\epsilon \geq 1$ . Since  $\arcsin(2\epsilon) = \Theta(\epsilon)$ , the asserted order of  $K_\epsilon$  follows. The equality in the no-wrap statement is direct. Finally  $L_\epsilon \leq \pi/(8 \arcsin(2\epsilon))$ , hence  $\Delta_\epsilon/4 = \pi/(8L_\epsilon) \geq \arcsin(2\epsilon)$ . Apply lemma 4 and monotonicity of sine on this range.  $\square$

The next lemma removes the alphabet size from the designer's hands: up to an explicit factor paid for the no-wrap property,  $K_\epsilon$  is the metric-entropy count of pairwise distinguishable  $Z$ -rotations at the working separation, an intrinsic function of the tolerance rather than a construction parameter.

**Lemma 5** (Synthesis-grid calibration). *For  $0 < \epsilon \leq 1/8$ , let  $N(\epsilon)$  be the maximum cardinality of a set of single-axis rotations  $\{R_P(\theta) : \theta \in \mathbb{R}/2\pi\mathbb{Z}\}$  that is pairwise  $4\epsilon$ -separated in  $d_{\text{proj}}$ , and let  $g_\epsilon := 4 \arcsin(2\epsilon)$ . Then*

$$N(\epsilon) = \left\lfloor \frac{2\pi}{g_\epsilon} \right\rfloor, \quad \frac{1}{2\epsilon} - 1 < N(\epsilon) \leq \frac{\pi}{4\epsilon}, \quad (15)$$

and the alphabet of proposition 3 satisfies

$$\frac{N(\epsilon)}{4} < K_\epsilon < \frac{N(\epsilon)}{4} + \frac{3}{2}. \quad (16)$$

Moreover the grid step is separation-tight,  $g_\epsilon \leq \Delta_\epsilon \leq g_\epsilon(1 + O(\epsilon))$ , so the factor 4 in (16) is exactly the restriction of the coefficient range to the quarter period  $[0, \pi/2]$  imposed by the no-wrap condition.

*Proof.* For a fixed non-identity Pauli  $P$ , both characters  $\pm 1$  are realized, so  $R_P(\theta)R_P(\theta')^\dagger$  has eigenvalue arguments  $\mp\delta/2$  with  $\delta = \theta - \theta'$ ; projectively  $\theta$  is defined modulo  $2\pi$ . Writing  $g \in [0, \pi]$  for the circular distance between  $\theta$  and  $\theta'$  in  $\mathbb{R}/2\pi\mathbb{Z}$ , the global-phase infimum of definition 5 is attained at the centered rotation, giving  $d_{\text{proj}} = 2 \sin(g/4)$ , increasing in  $g$ ; the computation is the one in the proof of lemma 4. Hence pairwise  $4\epsilon$ -separation is equivalent to pairwise circular distance at least  $g_\epsilon$ , because  $2 \sin(g_\epsilon/4) = 4\epsilon$  exactly. The maximum number of points of a circle of circumference  $2\pi$  with pairwise circular distance at least  $g_\epsilon$  is  $\lfloor 2\pi/g_\epsilon \rfloor$ , attained by equal spacing. The two-sided bound follows from  $x \leq \arcsin x \leq (\pi/2)x$  on  $[0, 1]$ . For (16), put  $y = \pi/(2 \arcsin(2\epsilon))$ , so that  $N(\epsilon) = \lfloor y \rfloor$  and, by (13),  $K_\epsilon = \lfloor y/4 \rfloor + 1$ . Then  $K_\epsilon > y/4 \geq N(\epsilon)/4$  and  $K_\epsilon \leq y/4 + 1 < (N(\epsilon) + 1)/4 + 1 = N(\epsilon)/4 + 5/4$ . Finally  $\Delta_\epsilon = \pi/(2L_\epsilon)$  with  $\pi/(8 \arcsin(2\epsilon)) - 1 \leq L_\epsilon \leq \pi/(8 \arcsin(2\epsilon))$  gives the step bounds.  $\square$

**Remark 6** (Ross–Selinger realization of the grid). *The synthesis layer realizes exactly this resolution. Ancilla-free Clifford+T approximation of an arbitrary  $z$ -rotation to within  $\epsilon$ —in the up-to-phase metric matched to  $d_{\text{proj}}$  [8, Def. 9.1, Cor. 9.5, Alg. 9.8]—is achieved by the algorithm of [8, Alg. 7.6] with  $T$ -count  $3 \log_2(1/\epsilon) + O(\log \log(1/\epsilon))$  in the typical case and  $c + 4 \log_2(1/\epsilon)$  for a constant  $c \approx 10$  in the worst case [8, Prop. 8.8 and §8.3]; given a factoring oracle the returned  $T$ -count is exactly optimal per instance*

[8, Prop. 8.2]. Every grid angle therefore costs the same  $\Theta(\log(1/\epsilon))$  synthesis budget, uniformly over the alphabet—the uniformity used by the fixed-total-error campaign of section X, whose rotations are synthesized by the `staq_grid_synth` implementation of this algorithm. Consequently the lower-bound term  $m \log_2 K_\epsilon$  is  $m \log_2 \Theta(1/\epsilon)$  with the constant pinned by (16), not selected; both directions of the calibration are stated as two-sided bounds and no tighter constant is claimed. For the cyclic grid used by the dispersed streams (corollary 7) the identification is exact:  $Q_\epsilon = \lfloor \pi/(2 \arcsin(2\epsilon)) \rfloor = N(\epsilon)$ , so the sharing modulus equals the metric-entropy count itself, and the factor 4 above concerns only the no-wrap subfamily.

## B. Streaming Encoding

**Definition 15** (Balanced streaming encoding). *The no-wrap specialization of definition 13 in which every round carries the same residue  $x_j/r$ , emitted on the round-robin schedule  $\sigma(i) = ((i-1) \bmod m) + 1$  over  $N = rm$  slots. Each token therefore contains  $x_j/r$ ; remark 7 records why this encoding proves no lower bound.*

**Definition 16** (Two-round masked-share lowering). *The cyclic specialization of definition 13 with  $Q = 3$ ,  $\alpha := \Delta = 2\pi/3$ ,  $r = 2$ , and  $\Sigma = \{+1\}$ , in which the aggregate alphabet is restricted to the binary packing  $K_{\text{pack}} = 2$ , so  $x \in \{0, 1\}^m$  while shares range over  $\mathbb{Z}_3$ . A record carries the round tag, the full generator address  $j$ , and one residue in  $\mathbb{Z}_3$ ; it does not carry  $x_j$ , since for every fixed first-round residue both  $x_j = 0$  and  $x_j = 1$  admit a compatible second-round residue. Each residue has  $O(1)$  parameter payload, while a worst-case complete token has  $\Theta(\log m)$  structural address bits as required by definition 8.*

**Remark 7** (Roles of the two flat schedules). *The balanced encoding of definition 15 is retained only as the configured experimental input in section VIII; because each of its tokens contains  $x_j/r$ , it is not used to prove a representation lower bound. The theorem uses the masked-share stream instead. Before the second round arrives, every first-round update has two valid binary aggregate completions, whereas the semantic representation supplies  $x$  itself. The separation below therefore charges recovery of an aggregate hidden across rounds, not a voluntary delay after a full coefficient was already visible.*

## C. Physically Realized Dispersed Representations

The masked-share and  $r$ -round streams above are stated adversarially: shares are chosen so that no token determines the aggregate. This subsection replaces the adversary by three standard fault-tolerant mechanisms. Time-dependent scheduling supplies the additive decomposition, randomized compiling supplies sign symmetry, and the synthesis grid of lemma 5 supplies the modulus. The hypotheses of the lower bounds proved in sections V and VI

below are then met by streams that a fault-tolerant stack emits anyway, not by a constructed secret sharing.

**Definition 17** (Scheduled phase-accumulation stream). *The cyclic specialization of definition 13 with modulus  $Q \geq 3$ , step  $\Delta = 2\pi/Q$ ,  $r \geq 2$  rounds, and the full sign set  $\Sigma = \{\pm 1\}$ . The residue  $a_j^{(t)} \in \mathbb{Z}_Q$  is a round- $t$  physical coefficient rounded to the grid and the sign  $s_j^{(t)}$  is a twirl variable ( $s_j^{(t)} \equiv +1$  for untwirled streams), so that the aggregate of definition 13 reads*

$$x_j = \sum_{t=1}^r s_j^{(t)} a_j^{(t)} \bmod Q. \quad (17)$$

*A twirled stream presents its signs in one of two modes. In folded mode each sign is absorbed at emission,  $a \mapsto sa \bmod Q$ , into the residue and its dressing Clifford records, as in compile-time randomized compiling [9, 27]. In tracked mode the signs accumulate in a classical Pauli-frame register carried by the control stack and are settled downstream of the emission point [26, 28].*

**Lemma 6** (Clifford interleaving). *Let the  $r$  rounds of a scheduled stream be separated by Clifford records, and let  $D_{t-1}$  denote the accumulated Clifford emitted before round  $t$  (so  $D_0 = I$ ). Suppose each  $D_{t-1}$  carries the generating set into itself up to sign and permutation,*

$$D_{t-1}^\dagger P_j D_{t-1} = \varepsilon_j^{(t)} P_{\pi_t(j)}, \quad \varepsilon_j^{(t)} \in \{\pm 1\}, \quad (18)$$

*with  $\pi_t$  a permutation of  $[m]$ . Then commuting every Clifford record to the end of the stream yields a scheduled stream of definition 17 on the same generators, in which the round- $t$  record of generator  $j$  contributes residue  $a_j^{(t)}$  with sign  $\varepsilon_j^{(t)} s_j^{(t)}$  to coordinate  $\pi_t(j)$ , followed by the fixed Clifford suffix  $D_r$ . Both  $\pi_t$  and  $\varepsilon^{(t)}$  are functions of the prefix, so the normalization is prefix-measurable, and every statement about definition 17 applies to the normalized image.*

*Proof.* Write the emitted unitary as  $D_r' L_r \cdots D_1' L_1$  with  $L_t$  the round- $t$  layer and  $D_t'$  the Clifford record following it, so  $D_t = D_t' \cdots D_1'$ . Moving each Clifford leftwards replaces  $L_t$  by  $D_{t-1}^\dagger L_t D_{t-1}$  and leaves  $D_r$  as a suffix. Since  $L_t = \prod_j R_{P_j}(\Delta s_j^{(t)} a_j^{(t)})$  and conjugation is an automorphism, (18) gives  $D_{t-1}^\dagger L_t D_{t-1} = \prod_j R_{P_{\pi_t(j)}}(\Delta \varepsilon_j^{(t)} s_j^{(t)} a_j^{(t)})$ , which is a round of definition 17 after the stated re-indexing. The Clifford records lie in the stream, so  $D_{t-1}$ , and hence  $\pi_t$  and  $\varepsilon^{(t)}$ , are computable from the prefix. Finally  $D_r$  is common to every member of the family and unitary, so it preserves projective distance and the separation of theorem 3 is unaffected.  $\square$

**Remark 8** (What the interleaving lemma does and does not cover). *Condition (18) is a real restriction, not a formality. It holds for Pauli-frame layers ( $\pi_t = \text{id}$ ,  $\varepsilon$  the frame signs), hence for the twirl insertions of randomized*

*compiling itself; and for routing or SWAP layers whose induced qubit relabeling permutes the generating set, such as the cyclic shifts of the chain characters  $a_j = e_j + e_{j+1}$ . It fails for a general Clifford: if  $D_{t-1}$  normalizes the diagonal group but sends some  $P_j$  to a product of generators rather than to a generator, the rounds accumulate in different bases and the coordinatewise structure of definition 17 is lost. It also does not rescue the non-commuting interleaving of remark 10(i): the layers separating the diagonal rounds of a generic Trotter step are  $R_X$  rotations, which are not Clifford at generic angles.*

**Remark 9** (Downstream compiler scope). *The bounds obtained from scheduled streams constrain compilers positioned downstream of the schedule, the twirl, and the lowering: the input is the emitted record stream, and the upstream context—the pre-randomization circuit, the schedule generator, and the twirl seed—is not part of it. This position is not hypothetical; it is where much practical compilation runs: (1) a cloud transpiler receiving a flat QASM/QIR program without access to the user’s high-level intent [29]; (2) a real-time stream optimizer inside the classical control stack [26]; (3) a vendor-side pass that post-processes an already emitted instruction stream. If the compiler can see the circuit before randomization and lowering, the theorems do not apply, and this boundary is itself informative. The fooling families range over every valid schedule and sign pattern; a compiler coupled to the upstream generator—for example, one holding the  $\lambda$ -bit seed that pseudorandomly generated the twirl signs—faces only  $2^\lambda$  of the  $2^m$  sign patterns, and the frame lower bound of proposition 5(ii) degrades from  $m$  to  $\lambda$  accordingly. The theorems therefore quantify exactly what is forfeited by discarding upstream context before compilation: they are the price attached to the design rule randomize and lower late.*

**Proposition 4** (Physical instantiation). *Each of the following mechanisms emits scheduled streams of definition 17, and each generated stream family satisfies the hypotheses of definition 20: every aggregate  $x \in [K]_0^m$ ,  $K = Q - 1$ , is realized by some valid stream, and no token or proper prefix determines it.*

- (a) Time-dependent Trotterization. *For a commuting time-dependent Hamiltonian  $H(t) = \sum_j c_j(t) P_j$  with step  $\tau$ , round  $t$  applies  $R_{P_j}(2c_j(t)\tau)$ ; grid rounding gives  $a_j^{(t)}$  and  $s \equiv +1$ .*
- (b) Layerwise-parameterized phase layers. *A depth- $r$  ansatz whose round- $t$  commuting layer carries free coefficients  $\theta_j^{(t)}$  gives  $a_j^{(t)}$  by grid rounding of  $\theta_j^{(t)}$ .*
- (c) Randomized compiling of (a) or (b). *Twirling multiplies each residue by  $s_j^{(t)}$ , presented in folded or tracked mode. If the twirl bits are independent and uniform, then under the canonical sharing distribution of proposition 9 every signed residue is uniform on  $\mathbb{Z}_Q$  and independent of  $x_j$ .*

In contrast, the constant schedule  $a_j^{(t)} \equiv a_j$  is degenerate: its first round determines the aggregate. This is the balanced-encoding degeneracy of remark 7.

*Proof.* In (a) and (b) the round coefficients are unconstrained across rounds, so every residue sequence—in particular every valid share decomposition of definition 20—arises from some schedule, and proposition 9 supplies realizability and token privacy verbatim. For prefixes: after any proper prefix, a final-round residue choice reaches every aggregate in  $\mathbb{Z}_Q^m$ , hence every  $x \in [K]_0^m$ . In (c), sign normalization  $a \mapsto sa$  is a bijection of  $\mathbb{Z}_Q$  for either sign, so normalizing every record maps the signed family onto the unsigned family record-for-record, and the same properties transfer; the distributional clause is the one proved in proposition 9, applied after normalization. For the constant schedule,  $x_j = ra_j \bmod Q$  is a function of the first round. An exhaustive enumeration at  $(m, r, Q) = (2, 2, 5)$ , frozen at `data/frozen/fooling_set_m2_r2_Q5.json`, independently confirms all three behaviors: every time-varying and every signed prefix class reaches all  $Q^m$  aggregates, and each constant-schedule prefix class reaches exactly one.  $\square$

**Remark 10** (Contiguity of the rounds, and Clifford interleaving). *Definition 17 is round-major over one commuting layer: the  $r$  rounds that accumulate into  $x_j$  reach the compiler as records of a single layer. Two consequences should be stated plainly, since they delimit proposition 4.*

(i) Non-commuting interleaving is outside the scope, and is not claimed. In a split-operator product formula such as  $H = H_{ZZ} + H_X$ , consecutive diagonal rounds are separated by non-commuting layers. There the rounds of the diagonal part may not be aggregated at all: no compact output exists, the premise of the recoverability question fails, and neither the hypotheses nor the conclusions of theorems 4 and 5 are asserted. Proposition 4(a) is a statement about layers whose  $r$  rounds are genuinely contiguous—a time-dependent or shaped schedule of one commuting layer, a commuting-block Trotterization, a contiguous diagonal ladder—and not about the interleaved structure of a generic Trotter step. The same qualification applies to (b): variational cost layers separated by mixers are aggregatable only within a block of consecutive commuting rounds, and it is such a block that (b) describes. Note that for a fully commuting  $H$  the Trotter error vanishes and the round structure is a scheduling choice rather than an accuracy requirement; what the bounds charge is precisely that choice, made upstream and visible to the downstream compiler only as  $r$  separate records.

(ii) Clifford interleaving is inside the scope. Rounds separated by Clifford records normalize back to definition 17 whenever the accumulated Clifford permutes the generating set up to sign; this is lemma 6, and remark 8 delimits it. Contiguity in definition 17 should therefore be read modulo such Clifford records.

Contiguity and coefficient freedom are separate hypotheses, and different families supply them. Con-

tiguous diagonal rounds arise structurally in controlled-power ladders and phase estimation, where  $c\text{-}U^{2^k}$  contributes a sequence of diagonal layers with no intervening non-commuting structure (the `qpe_controlled_powers` and `controlled_power_ladder` families of section IX). Round coefficients that vary freely across rounds—what the fooling families of proposition 4 require—are supplied by time-dependent schedules and by layerwise-parameterized ansatz blocks.

**Proposition 5** (Frame as cross-cut state). *Let a compiler in the charged model of definition 9 be correct within projective error  $\epsilon < \sin(\pi/(2Q))$  on every valid twirled stream of definition 17, and fix any inter-round cut  $c$  with committed prefix output of length  $\ell_{\text{pre}}$  and five-field snapshot budget  $B_{\text{cut}}(c)$ .*

(i) *In both presentation modes the normalized stream is a valid dispersed stream of definition 20, so theorems 4 and 5 apply verbatim to the physical stream.*

(ii) *In tracked mode,*

$$\ell_{\text{pre}} + B_{\text{cut}}(c) \geq m \quad (19)$$

*already from the pending-frame classes alone, and the frame register is live classical restart state, serialized in `ser_store` (or `ser_control`) by definition 9. For a deferred compiler with  $\ell_{\text{pre}} = O(\log m)$ , the crossing snapshot therefore carries  $m - O(\log m)$  bits that determine the frame action on the  $m$  generators: the Pauli frame is a sub-register of the cross-cut state.*

(iii) *In folded mode the flush map is a prefix-measurable sign normalization, injective on sign components; the same  $m$  bits appear as the sign components of the signed residues, serialized in `ser_parameter` or `ser_store`, and the bounds of (i) are unchanged.*

*Proof.* (i) Sign normalization is a bijection of  $\mathbb{Z}_Q$ , so normalizing every record maps the twirled family bijectively onto the family of definition 20 while preserving aggregates, record boundaries, and cut positions; the cited theorems apply to the image and pull back verbatim.

(ii) Take prefixes whose residue fields emit the residue 1 exactly once on each generator and are otherwise zero, and whose sign fields realize the frame class  $\sigma \in \{\pm 1\}^m$ ; the residue fields are identical across classes. Choose one common suffix adding a residue  $a_j$  on generator  $j$  with both  $a_j + 1 \bmod Q$  and  $a_j - 1 \bmod Q$  in  $[K]_0$ ; such  $a_j$  exists for every  $Q \geq 3$  ( $a_j = 2$  for  $Q = 3$ ,  $a_j = 1$  otherwise). The completed aggregates  $x_j = \sigma_j + a_j \bmod Q$  then differ between two classes in every coordinate where the frames differ, and both completions are valid family instances. By theorem 3 the corresponding targets are separated by  $2 \sin(\pi/(2Q)) > 2\epsilon$ , so a correct compiler must emit different final outputs; but a deterministic suffix computation is a function of the snapshot and the common suffix, so two runs with equal committed prefix and equal snapshot emit identical final outputs. The pair

(prefix output, snapshot) is therefore injective on the  $2^m$  frame classes; prefix-freeness of the declared serializers and Kraft’s inequality give (19). The frame register is classical restart state, so its bits lie in the named fields; subtracting the deferred-output cap bounds the snapshot share.

(iii) In folded mode the sign is consumed at emission, so the crossing information consists of the signed residues themselves; injectivity on sign components is the bijection of (i) restricted to fixed unsigned residues, and (19) holds with the same fooling family presented in folded records.  $\square$

**Corollary 4** (Physically realized lower bounds). *For the streams of proposition 4(a)–(c) with the tuned modulus  $Q_\epsilon$  of corollary 7, the hypotheses of definitions 20 and 21 are met, so theorems 4 and 5 (and, for the two-round binary case, theorem 2) hold verbatim on the physical stream: any compiler positioned downstream of the schedule and of the twirl obeys  $\bar{A}_p + (2p - 1)S \geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$ , and under tracked randomized compiling the crossing state contains the Pauli frame in the sense of proposition 5(ii).*

*Proof.* Immediate from propositions 4 and 5; no separate argument is required.  $\square$

A single knob interpolates between the two regimes, and the interpolation is itself a theorem: the fraction of coordinates that reach the compiler in dispersed rather than pre-aggregated form and the compiler state that survives the cut trade along one line. We call that fraction the *dispersal fraction*  $\eta$ . The name is deliberate:  $\eta$  is not the twirl strength of randomized compiling (the probability with which a random Pauli is applied), and the proof below never uses the twirl—it conditions on the dispersal mask and applies theorem 5 to the dispersed coordinates. What  $\eta$  measures is how much of the phase content upstream chose to present in late-aggregated form, which is exactly the axis of the design rule.

**Proposition 6** (Dispersal–memory Pareto bound). *Fix a dispersal fraction  $\eta \in [0, 1]$ ,  $Q \geq 3$ ,  $K = Q - 1$ , and  $0 \leq \epsilon < \sin(\pi/(2Q))$ . Let each generator independently with probability  $\eta$  enter the stream in dispersed twirled form (canonical sharing of proposition 9 with uniform signs), and otherwise in aggregate-exposed form. Let  $\mathcal{A}$  be correct within  $\epsilon$  with probability at least  $1 - \delta$  on every valid stream and mask, making  $p = p(m)$  forward passes in the charged model of definition 21. Then, in expectation over the dispersal mask  $T$ ,*

$$\mathbb{E}_T[\bar{A}_p + (2p - 1)S] \geq (1 - \delta) \eta m \log_2 K - h_2(\delta). \quad (20)$$

*Under tracked presentation, each twirled generator additionally contributes its frame bit as in proposition 5(ii).*

*Proof.* Condition on the mask  $T$ . Fixing the exposed coordinates and dispersing the coordinates in  $T$  yields a subfamily satisfying the hypotheses of definition 20 with  $|T|$  effective generators, so theorem 5 applies verbatim




FIG. 2. Measured dispersal–memory tradeoff on the executable two-round masked-share device ( $Q = 3$ , binary aggregates, tracked signs; every point restart-verified: the serialized snapshot alone plus the round-two suffix reproduces the exact output). Left: a deferred compact-output compiler pays the dispersal in crossing state; measured slopes 10.3/21.1/43.3 bytes per unit  $\eta$  at  $m = 32/64/128$  sit on the zero-error floor  $\eta m(\log_2 3 + 1)/8$  bytes (dashed; ratios 0.99–1.05), the residue-plus-frame cost of theorem 4 and proposition 5. Right: a streaming compiler pays the same dispersal in committed prefix output ( $A_{\text{pre}}$  slopes 60/129/260 bytes per unit  $\eta$ , larger by the public addresses) while its crossing state stays flat at  $\leq 8$  bytes. The horizontal axis is the dispersal fraction  $\eta$  of proposition 6. Frozen artifact: `data/frozen/eta_sweep_pareto.json`, generated by `PaperDraft/scripts/run_eta_sweep.py`.

and gives  $\bar{A}_p + (2p - 1)S \geq (1 - \delta)|T| \log_2 K - h_2(\delta)$  for every fixed  $T$ . Take the expectation over the Bernoulli mask and use  $\mathbb{E}|T| = \eta m$ .  $\square$

Figure 2 measures (20) on the executable masked-share device: the two admissible payment channels of the bound appear as the two panels, the state channel sits on the predicted floor to within 5%, and the exchange between them is linear in the dispersal fraction with slope proportional to  $m$ .

The two panels also expose an asymmetry that the frontier itself does not predict. Per dispersed generator the state channel costs about 2.6 bits—the floor  $\log_2 3 + 1 = 2.585$  of residue plus sign—while the committed channel costs about 16: an exchange rate near 6:1 between two currencies the frontier charges at par. The next proposition isolates the part of that gap which is not an artifact of the serializer.

**Proposition 7** (Commitment toll). *Fix  $Q$ ,  $K = Q - 1$ ,  $\epsilon < \sin(\pi/(2Q))$ , and a dispersal mask  $T \subseteq [m]$  with  $|T| = k$ . Let  $\mathcal{A}$  be correct on every valid stream of definition 17 and require its committed prefix to be independently parseable—the append-only first-exposure regime of remark 4, in which an emitted record is executed and cannot be reinterpreted once the suffix arrives. Then at the inter-round cut*

$$A_{\text{pre}} \geq \underbrace{k \log_2(2Q)}_{\text{payload}} + \underbrace{\log_2 \binom{m}{k}}_{\text{toll}}, \quad (21)$$

while a deferred compiler attains  $A_{\text{pre}} = 0$  with  $S \leq k \log_2(2Q) + O(\log m)$ . The toll is the entropy of the commitment schedule; it vanishes iff that schedule is determined by public data, in particular at  $k = m$ .

*Proof.* Two valid streams whose round-one content differs in the mask, in a residue, or in a sign require different round-one rotations, and by theorem 3 their completions are separated by  $2 \sin(\pi/(2Q)) > 2\epsilon$ ; a prefix that is executed as emitted must therefore differ between them. The map from  $(T, \{a_j^{(1)}\}_{j \in T}, \{s_j^{(1)}\}_{j \in T})$  to committed prefixes is thus injective on  $\binom{m}{k} Q^{k2^k}$  values, and prefix-freeness of the declared serializer with Kraft’s inequality gives (21). For the deferred compiler, the round-major order of definition 17 is public, so stored residues are addressed by position and the mask is disclosed by the round-two records at no charge; the  $O(\log m)$  term is the header.  $\square$

The two channels are therefore *not* at par once commitment is irreversible: committing early costs the entropy of one’s own decision of when to commit, and deferring is handed that same information by the stream that follows. Because  $\log_2 \binom{m}{k} / k \rightarrow h_2(\eta) / \eta$ , the toll is bounded by a constant—about 2 bits per generator at  $\eta = \frac{1}{2}$  and 0 at  $\eta = 1$ —and does not grow with  $m$ .

Figure 3 tests this on the same device under a preregistration frozen before the run (`PaperDraft/research/E8_preregistration.md`), with every cell restart-verified. Four serializers are compared at  $m \in \{128, 512, 2048\}$  (and  $m$  up to 32768 for two of them): the deferred snapshot, the address-per-record serializer of figure 2, an identity-optimal self-identifying prefix, and a control whose prefix omits the mask and recovers it from the suffix. The measured toll matches  $\log_2 \binom{m}{k} / k$  within 0.17 bits at the worst cell and within 0.02 bits typically; it is flat in  $m$  at fixed  $\eta$  (2.11/1.99/1.97 bits at  $\eta = \frac{1}{2}$  across a  $16\times$  range of  $m$ ); it is exactly zero at  $\eta = 1$ ; and it is exactly zero at every  $\eta$  in the control, confirming that irreversibility of commitment, not addressing, is what buys the asymmetry.

This also settles what the 6:1 of figure 2 is. Under the identity-optimal serializer the ratio is 1.00 at  $\eta = 1$  and at most 2.64 anywhere in the measured range; the remainder is codec slack, and it is the slack—not the toll—that grows with address width, from 5.4:1 at  $m = 128$  to 10.8:1 at  $m = 32768$ . The preregistration recorded this as a prediction *against* the natural reading of figure 2, and we report it as such: there is no address-width law here.

Two attempts to carry the toll onto compilers we do not control were preregistered, and **both failed**; we report them rather than repair them. The first asked only whether real toolchains spend committed bytes on addresses, and its design placed every rotation on the low-index qubits, so it never varied the address width it was meant to vary (a corrected rerun, labelled non-preregistered in the frozen artifact, finds 2.95 bytes of growth per rotation over a  $1000\times$  range of register width, itself below the preregistered threshold). The second tested the mechanism directly: two circuits identical in




FIG. 3. The commitment toll, preregistered and measured on the two-round masked-share device; every cell restart-verified. Left: bits per committed generator paid by a self-identifying committed prefix in excess of the deferred snapshot, against the predicted schedule entropy  $\log_2 \binom{m}{k} / k$  (dashed). The toll is flat in  $m$  and vanishes at  $\eta = 1$ . Right: the per-generator exchange rate between the committed and state channels. Solid: irreducible, under an identity-optimal serializer—at most 2.64, and exactly 1 when the commitment schedule is public. Dotted: the same rate as actually serialized by the address-per-record codec of figure 2, which is where the 6:1 comes from and which is the only part that grows with address width. Frozen artifact: `generated/e8_commit_toll.json`, generated by `PaperDraft/scripts/run_commit_toll.py`.

$m$ , in  $k$ , in rotation count and in angle multiset, differing only in whether the dispersed subset is an arithmetic progression or a random  $k$ -subset, emitted through Qiskit and TKET and compared by compressed size. Its textual control passed—the variants differ by under 0.05% of raw bytes, so the comparison is not confounded by text length—and on TKET the excess reaches 86% of  $\log_2 \binom{m}{k}$  at  $k = m/2$ , comfortably above the seed noise. But the excess peaks at  $k = m/4$  rather than  $k = m/2$ , and on Qiskit it is 8% of prediction and smaller than the seed-to-seed noise. Under the frozen rule this is a failure, and we make no third-party claim: a general-purpose compressor applied to a 100 kB instruction text is evidently not a sharp enough entropy estimator for a few hundred bytes of schedule information. Proposition 7 rests on the instrumented device of figure 3 alone, where every cell is restart-verified; whether the toll is measurable at a third-party API boundary remains open.

## V. INFORMATION LOWER BOUNDS

We prove two lower bounds by one-way message arguments [30]. The first uses the  $K_\epsilon^m$ -point no-wrap packing and Fano’s inequality to obtain an explicit accuracy–description tradeoff, including randomized compilers. The second injects the  $3^m$  possible first-round masks into the message crossing the inter-round cut. Its deterministic zero-error bound is matched by a memory-capped hybrid compiler, and a conditional-Fano extension covers randomized failure probability  $\delta$ .

### A. Packing Decoders

**Lemma 7** (Nearest-packing decoder). *Let  $\mathcal{F} = \{U_x : x \in \mathcal{X}\}$  have pairwise distance at least  $4\epsilon$ . From any output circuit  $C$ , define*

$$\hat{x}(C) \in \arg \min_{y \in \mathcal{X}} d_{\text{proj}}(U(C), U_y),$$

*with a fixed tie rule. If  $d_{\text{proj}}(U(C), U_x) \leq \epsilon$ , then  $\hat{x}(C) = x$ .*

*Proof.* The true target is at distance at most  $\epsilon$ . Every other target is at distance at least  $4\epsilon - \epsilon = 3\epsilon$  by the triangle inequality, so it cannot minimize.  $\square$

The following binary-coordinate version will be used for the masked-share tradeoff, where  $K = 3$  denotes the sharing modulus but the aggregate alphabet remains  $\{0, 1\}$ .

**Lemma 8** (Bit-class separation). *For the binary subfamily of definition 14, put  $\epsilon_0 = \sin(\Delta/4)$ . For each  $i \in [m]$ , partition the family into  $\mathcal{F}_i^{(0)} = \{U_x : x_i = 0\}$  and  $\mathcal{F}_i^{(1)} = \{U_x : x_i = 1\}$ . Then the inter-class distance satisfies*

$$\text{dist}(\mathcal{F}_i^{(0)}, \mathcal{F}_i^{(1)}) := \inf_{\substack{x_i=0 \\ y_i=1}} d_{\text{proj}}(U_x, U_y) \geq 2\epsilon_0. \quad (22)$$

*Proof.* Any pair with  $x_i \neq y_i$  has  $x \neq y$ , so lemma 4 gives  $d_{\text{proj}}(U_x, U_y) \geq 2\epsilon_0$ . Taking the infimum preserves the inequality.  $\square$

**Lemma 9** (Representation-invariant decoder). *For each  $i \in [m]$  and any circuit  $C$ , define*

$$\delta_i(C) := \arg \min_{b \in \{0,1\}} \min_{y: y_i=b} d_{\text{proj}}(U(C), U_y). \quad (23)$$

*(Bob's local computation is unbounded, so he may exhaustively search the  $2^{m-1}$  candidates of each class.) If  $d_{\text{proj}}(U(C_x), U_x) \leq \epsilon < \epsilon_0$ , then  $\delta_i(C_x) = x_i$ .*

*Proof.* The correct class  $b = x_i$  contains  $U_x$  itself, so its inner minimum is at most  $\epsilon$ . Suppose the wrong class  $b \neq x_i$  also attained an inner minimum  $\leq \epsilon$ , realized by some  $U_y$  with  $y_i = b$ . Then the triangle inequality gives  $d_{\text{proj}}(U_x, U_y) \leq 2\epsilon < 2\epsilon_0$ , contradicting lemma 8. Hence only the correct class attains the minimum, i.e.  $\delta_i(C_x) = x_i$ .  $\square$

### B. Output-Description Lower Bound

**Theorem 1** (Randomized  $\epsilon$ -description tradeoff). *Let  $0 < \epsilon \leq 1/8$ , let  $\mathcal{F}_{m,\epsilon}$  be the  $K_\epsilon$ -ary family of proposition 3, and let a possibly randomized compiler  $\mathcal{A}$  output  $C_{x,\rho}$  satisfying, for every  $x$ ,*

$$\Pr[d_{\text{proj}}(U(C_{x,\rho}), U_x) \leq \epsilon] \geq 1 - \delta, \quad 0 \leq \delta < 1/2. \quad (24)$$

*The compiler may use random-access RAM/DAG IR, hash tables, caches, and a dynamic  $p = p(m, x, \rho)$  as in definition 9. Writing  $h_2(\delta) = -\delta \log_2 \delta - (1 - \delta) \log_2 (1 - \delta)$  with  $0 \log 0 := 0$ , each output code separately obeys*

$$\begin{aligned} \max_x \mathbb{E}_\rho B_{\text{rel}}(C_{x,\rho} \mid \mathcal{F}_{m,\epsilon}) &\geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta), \\ \max_x \mathbb{E}_\rho B_{\text{self}}(C_{x,\rho}) &\geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta). \end{aligned} \quad (25)$$

*In particular, for fixed  $\delta < 1/2$  both are  $\Omega(m \log(1/\epsilon))$ .*

*At  $\delta = 0$ , the family-relative bound is tight to lower-order framing terms: a deterministic semantic compiler uses*

$$B_{\text{rel}}(C_x \mid \mathcal{F}_{m,\epsilon}) \leq \lceil m \log_2 K_\epsilon \rceil + O(\log m + \log K_\epsilon). \quad (26)$$

*Separately, for  $P_j = Z_j$  and  $RZ \in B$ , the explicit self-contained circuit satisfies  $B_{\text{self}}(C_x) = O(m(\log m + \log K_\epsilon))$  and  $G(C_x) \leq m$ ; no self-contained-bit or gate-count tightness claim is made.*

*Proof.* Let  $X$  be uniform on  $[K_\epsilon]_0^m$  and let  $C = C_{X,\rho}$ . By (14), the nearest-packing decoder of lemma 7 has error probability at most  $\delta$ . Fano's inequality gives

$$H(X \mid C) \leq h_2(\delta) + \delta \log_2(K_\epsilon^m - 1),$$

and hence  $I(X; C) \geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$ . For either prefix-free code  $E$ , source coding gives  $\mathbb{E}|E(C)| \geq H(C) \geq I(X; C)$ . The maximum conditional expected length over  $x$  is at least this uniform average, proving each line of (25). This argument depends only on the output distribution, so random access, dynamic passes, and internal IR do not weaken it.

For (26), the public dictionary supplies  $(P_1, \dots, P_m, K_\epsilon, \Delta_\epsilon, B)$  and aggregate mode packs the base- $K_\epsilon$  vector into one fixed-length integer block. Its header serializes  $m, K_\epsilon$ , the mode, and the block boundary. In the canonical self-contained construction, each nonzero  $Z_j$  term has an  $O(\log m)$  address and an  $O(\log K_\epsilon)$  rational- $\pi$  parameter under (2); there are at most  $m$  terms.  $\square$

**Corollary 5** (Deterministic fixed-grid specialization). *For any no-wrap  $(K, \Delta)$  family of definition 14, a deterministic compiler accurate to  $\epsilon < \sin(\Delta/4)$  has, separately,*

$$B_{\text{rel}}^{\text{out}} \geq \lceil m \log_2 K \rceil, \quad B_{\text{self}}^{\text{out}} \geq \lceil m \log_2 K \rceil.$$

*In particular the binary fixed-step family has the  $m$  line used in the configured scaling panel.*

*Proof.* Accuracy balls are disjoint by lemma 4, so the deterministic output map is injective on  $K^m$  inputs. If  $\ell_x$  are their lengths under either prefix-free code, Kraft's inequality gives  $\sum_x 2^{-\ell_x} \leq 1$ . Thus  $K^m 2^{-\max_x \ell_x} \leq 1$ , proving the two rounded bounds.  $\square$

**Remark 11** (What the counting bound does and does not say). *Equation (25) holds regardless of internal resources, but the two inequalities must not be cross-compared. The*



packed  $K_\epsilon$ -ary block in (26) is an upper bound only for  $B_{\text{rel}}(\cdot \mid \mathcal{F}_{m,\epsilon})$ ; it is not a self-contained target-basis circuit string. Likewise  $G(C)$  is never multiplied by an implicit “bits per gate” constant. The representation-dependent question is addressed by the masked-share cut below under one fixed family-relative output code.

### C. The Masked-Share Compact-Output-Cut Tradeoff

The semantic input supplies the aggregate vector  $x$ . The flat input of definition 16 instead supplies an arbitrary first share  $u \in \mathbb{Z}_3^m$  and only later a second share  $v$  satisfying  $u + v = x \pmod{3}$ . The cut between these rounds exposes information that is specific to the representation: a correct compiler must preserve enough about  $u$  to handle every still-valid continuation  $v$ .

**Definition 18** (Masked-share cut and precommitted output). *Run  $\mathcal{A}$  on a masked-share stream and let  $c_{\text{MS}}$  be the arrival cut after all  $m$  first-round updates have been delivered and before any second-round update is available. Before this cut the compiler may use arbitrary RAM/DAG operations and a dynamic number of internal scans. A buffered or rereadable first-round record is included in  $B_{\text{store}}$  or  $B_{\text{IR}}$ ; the only excluded prefix is one irreversibly committed to the final write-once output. Let*

$$\begin{aligned} B_{\text{pre}}^{\text{rel}}(\mathcal{A}, m \mid \mathcal{F}_m) &:= \max_{u \in \mathbb{Z}_3^m} \ell_{\text{pre}}(u), \\ B_{\text{cut}}^{\text{MS}}(\mathcal{A}, m) &:= \max_{u \in \mathbb{Z}_3^m} B_{\text{cut}}(c_{\text{MS}}(u)). \end{aligned} \quad (27)$$

Here  $\ell_{\text{pre}}(u)$  is the literal bit length already committed under the prefix-monotone stream  $E_{\text{rel}}$ ; it is zero if final output has not started. Thus  $B_{\text{pre}}^{\text{rel}} \leq B_{\text{rel}}^{\text{out}}$ . For randomized  $\mathcal{A}$ , barred quantities replace the maxima in (27) by  $\max_u \mathbb{E}_\rho[\ell_{\text{pre}}(u, \rho)]$  and  $\max_u \mathbb{E}_\rho[B_{\text{cut}}(c_{\text{MS}}(u, \rho))]$ .

**Theorem 2** (Representation-dependent masked-share tradeoff). *Let the sharing modulus be  $K_{\text{share}} = 3$ , let  $\alpha = 2\pi/3$ , and  $\epsilon_0 = \sin(\pi/6) = 1/2$ . Suppose first that a deterministic compiler  $\mathcal{A}$  is correct for every valid two-round decomposition in definition 16:*

$$\begin{aligned} d_{\text{proj}}(U(C(u, v)), U_{u+v \bmod 3}) &\leq \epsilon < \epsilon_0, \\ u + v \bmod 3 &\in \{0, 1\}^m. \end{aligned}$$

Then, under the single family-relative output code of definition 10,

$$B_{\text{pre}}^{\text{rel}}(\mathcal{A}, m \mid \mathcal{F}_m) + B_{\text{cut}}^{\text{MS}}(\mathcal{A}, m) \geq \lceil m \log_2 3 \rceil. \quad (28)$$

Consequently, if the final output is family-relative compact,

$$B_{\text{rel}}^{\text{out}}(\mathcal{A}, m \mid \mathcal{F}_m) \leq m + c \log_2(m + 2),$$

then

$$B_{\text{cut}}^{\text{MS}}(\mathcal{A}, m) \geq \lceil m \log_2 3 \rceil - m - c \log_2(m + 2) = \Omega(m). \quad (29)$$

In contrast, a one-pass compiler receiving the semantic aggregate vector  $x$  streams the same aggregate-mode code with  $B_{\text{rel}}^{\text{out}} = m + O(\log m)$  and  $B_{\text{cut}} = O(\log m)$ .

More generally, let  $0 \leq \delta < 1/2$  and let  $\mathcal{A}$  be randomized and satisfy the displayed accuracy condition on every valid  $(u, v)$  with probability at least  $1 - \delta$ . It may use arbitrary random access and a dynamic  $p = p(m, u, v, \rho)$ . Then

$$\overline{B}_{\text{pre}}^{\text{rel}} + \overline{B}_{\text{cut}}^{\text{MS}} \geq (1 - \delta)m - h_2(\delta). \quad (30)$$

In particular, for a deferred aggregate-mode compiler with  $\overline{B}_{\text{pre}}^{\text{rel}} = O(\log m)$  and fixed  $\delta < 1/2$ , the flat representation still requires  $\overline{B}_{\text{cut}}^{\text{MS}} = \Omega(m)$ .

*Proof.* Let Alice hold an arbitrary first share  $u \in \mathbb{Z}_3^m$ . She simulates  $\mathcal{A}$  through  $c_{\text{MS}}$  and sends the already committed family-relative output prefix together with the complete five-field restart serialization of (3). All head positions, lengths, addresses, parameters, stores, windows, and materialized IR prefixes needed to resume are present by definition, including any rereadable old input. Bob, given a second share  $v$ , resumes the computation and runs its dynamically chosen remaining operations locally.

The message must distinguish every pair  $u \neq u'$ . To see this, choose a single continuation  $v$  coordinatewise. If  $u_j = u'_j$ , set  $v_j = -u_j$ . If  $u'_j = u_j + 1 \pmod{3}$ , again set  $v_j = -u_j$ ; the two aggregates are then 0 and 1. In the remaining case  $u_j = u'_j + 1 \pmod{3}$ , set  $v_j = -u'_j$ ; the aggregates are 1 and 0. Thus both  $(u, v)$  and  $(u', v)$  are valid, binary aggregate instances, and their aggregate vectors differ whenever  $u \neq u'$ .

If Alice’s two messages were identical, determinism would make Bob emit the same final circuit on these two valid completions. That circuit cannot be within  $\epsilon$  of both targets, because lemma 4 gives target distance at least  $2\epsilon_0 > 2\epsilon$ . Hence the message map is injective on the  $3^m$  choices of  $u$ . Serialize the restart snapshot first; its control field gives the length of the committed prefix that follows, so the complete messages are prefix-free. Kraft’s inequality therefore gives maximum message length at least  $\lceil m \log_2 3 \rceil$ . The sum of the two separate worst-case field lengths is at least that maximum, which proves (28). Prefix monotonicity and the assumed upper bound on the final output give (29).

For the semantic upper bound, the compiler writes the aggregate-mode header, copies the arriving bits  $x_j$  into the support-mask field, and retains only a self-delimiting index and finite control. It stores no coefficient table or intermediate graph. The sparse parameter payload is fixed and the live address/index costs  $O(\log m)$ , so every cut has  $B_{\text{cut}} = O(\log m)$  under the same serializer.

For the randomized extension, draw  $U$  uniformly from  $\mathbb{Z}_3^m$  and  $X$  uniformly from  $\{0, 1\}^m$ , independently, and give Bob  $V = X - U \pmod{3}$ . Then  $X$  and  $V$  are independent and  $H(X \mid V) = m$ . Alice’s message  $M$  consists of the committed prefix and the full restart snapshot after processing  $U$ . The snapshot includes the seed/tape and its cursor, RAM, caches, old-input buffers, and IR, so

Bob can resume the exact randomized run upon receiving  $V$ . The binary decoder of lemma 9 recovers  $X$  with error at most  $\delta$ . Conditional Fano therefore gives

$$\begin{aligned} H(X \mid M, V) &\leq h_2(\delta) + \delta m, \\ I(X; M \mid V) &\geq (1 - \delta)m - h_2(\delta). \end{aligned}$$

The two message fields are self-delimiting, so source coding lower-bounds their expected total length by this conditional information, yielding (30). This is a conditional-Fano (equivalently, distributional/Yao) argument and does not impose a fixed pass count or sequential internal memory.  $\square$

**Proposition 8** (Memory-capped hybrid compilers). *For every  $\beta \in [0, 1]$ , let  $h = \lfloor \beta m \rfloor$ . There is a deterministic zero-error compiler for the masked-share stream, using the ordered-share mode of  $E_{\text{rel}}$ , such that at the inter-round cut*

$$\begin{aligned} B_{\text{cut}}^{\text{MS}} &\leq h \log_2 3 + O(\log m), \\ B_{\text{pre}}^{\text{rel}} &\leq (m - h) \log_2 3 + O(\log m). \end{aligned} \quad (31)$$

Thus (28) is matched to leading order for every memory fraction  $\beta$  under this single family-relative encoding. This statement does not assert a matching curve for  $B_{\text{self}}$  or  $G$ .

*Proof.* The ordered-share header records  $h$ . During round one the compiler packs  $u_1, \dots, u_h$  into a base-three RAM block and immediately commits  $u_{h+1}, \dots, u_m$  in their public generator order. It needs only the packed block, an index, and framing state at the cut. During round two it emits the aggregates  $u_j + v_j \pmod{3}$  for  $j \leq h$  and the separate residues  $v_j$  for  $j > h$ . The latter combine with the already emitted first shares; all rotations commute, and proposition 2 gives the target up to global phase. Fixed-order packing gives the two bounds in (31); no qubit address is hidden, because the ordered-share mode declares its public schedule and split in its header.  $\square$

**Corollary 6** (Compactness versus flat-share recovery). *On the masked-share representation, any compiler with  $B_{\text{cut}}^{\text{MS}} = o(m)$  must have*

$$B_{\text{rel}}^{\text{out}}(\mathcal{A}, m \mid \mathcal{F}_m) \geq m \log_2 3 - o(m),$$

whereas family-relative optimal output  $m + O(\log m)$  requires at least the linear cut budget of (29). This is a representation tradeoff: the semantic aggregate stream attains both compact output and  $O(\log m)$  cut budget.

*Proof.* Use  $B_{\text{pre}}^{\text{rel}} \leq B_{\text{rel}}^{\text{out}}$  in (28), then compare with the semantic construction in theorem 2.  $\square$

**Remark 12** (Three points under one encoding). *All points in (28) use  $E_{\text{rel}}(\cdot \mid \mathcal{F}_m)$ . The semantic pass-through has compact relative output and logarithmic cut budget. A flat compiler that materializes the first round in a table or intermediate graph pays its full serialized*

*structure and parameters in  $B_{\text{cut}}$ . A one-pass flat compiler may instead commit update residues before seeing their mates, but those records enlarge  $B_{\text{pre}}^{\text{rel}}$  and hence the final relative code. The hybrids of proposition 8 interpolate between these endpoints and attain the deterministic lower envelope to leading order for every  $\beta$ . No claim about  $B_{\text{self}}$  or  $G(C)$  is inferred from this curve.*

**Remark 13** (Relation to the exact-semantics results and model scope). *The exact-semantics corollary (corollary 1) bounds final output length under exactness and noncollision, via pumpability. Theorem 1 bounds two explicitly named output-description codes via counting. Theorem 2 is different: it uses the two-round completion adversary and the complete bit-level restart snapshot  $B_{\text{cut}}$ . A global ZX graph, phase-polynomial table, spilled store, large-precision parameter, addressed window, RAM/DAG, cache, hash table, dynamic pass schedule, or random seed is therefore an accounted recovery channel, not an exception. The deterministic theorem allows arbitrary  $p = p(m, I)$  in this charged model, and (30) allows input-dependent randomized  $p = p(m, I, \rho)$ . The balanced experimental schedule is not claimed to instantiate this lower bound.*

## VI. UNIFIED $K$ -ARY RANDOMIZED MULTIPASS FRONTIER

The preceding results isolate output information and the binary two-round cut. The following package aligns the packing modulus with a genuinely dispersed additive-update representation and makes the number of passes and the full serialized state cap explicit. It also separates the proved constant-factor upper envelope from the remaining model-dependent gaps.

### A. Exact cyclic $K$ -ary packing

The no-wrap construction of definition 14 is convenient for output counting, but modular additive shares require a grid whose period is also the sharing modulus. The following formulation records exactly when wraparound creates a collision.

**Definition 19** (Wrap margin). *For an alphabet  $[K]_0$ , step  $\Delta > 0$ , and  $K \geq 2$ , define*

$$\text{dist}_{2\pi}(a, 0) := \min_{z \in \mathbb{Z}} |a - 2\pi z|, \quad (32)$$

$$\gamma(K, \Delta) := \min_{1 \leq d \leq K-1} \text{dist}_{2\pi}(d\Delta, 0). \quad (33)$$

Let  $P_1, \dots, P_m$  be independent Pauli-Z characters as in definition 14, and set

$$U_x^{(K, \Delta)} := \exp \left[ -\frac{i}{2} \sum_{j=1}^m x_j \Delta P_j \right], \quad x \in [K]_0^m. \quad (34)$$



**Theorem 3** (Exact worst-pair separation). *For every  $m \geq 1$ ,*

$$\min_{x \neq y} d_{\text{proj}}(U_x^{(K,\Delta)}, U_y^{(K,\Delta)}) = 2 \sin\left(\frac{\gamma(K,\Delta)}{4}\right). \quad (35)$$

*Consequently the family is projectively injective exactly when  $\gamma(K,\Delta) > 0$ . If  $(K-1)\Delta \leq \pi$ , then  $\gamma(K,\Delta) = \Delta$ . If  $Q \geq K+1$  and  $\Delta = 2\pi/Q$ , then again  $\gamma(K,\Delta) = \Delta$ , even though the coefficient interval  $[0, (K-1)\Delta]$  may cross the principal-argument cut at  $\pi$ .*

*Proof.* Fix  $x \neq y$ , put  $s = x - y$ , and choose  $j$  with  $s_j \neq 0$ . Independence of the characters makes their joint sign map surjective. Hence the relative unitary  $U_x^{(K,\Delta)}(U_y^{(K,\Delta)})^\dagger$  has two eigenvalues obtained from sign patterns that agree off  $j$  and whose circular angular separation is  $\text{dist}_{2\pi}(|s_j|\Delta, 0) \geq \gamma(K,\Delta)$ . After any common phase rotation, one of two circle points separated by  $\eta \in [0, \pi]$  is at chordal distance at least  $2 \sin(\eta/4)$  from 1. This proves the lower bound in (35) after taking the global-phase infimum.

Choose  $d \in \{1, \dots, K-1\}$  attaining  $\gamma(K,\Delta)$  and take a pair that differs by  $d$  in one coordinate only. Its relative spectrum consists of two points with circular separation  $\gamma(K,\Delta)$ , so centering the shorter arc attains  $2 \sin(\gamma/4)$ . This proves equality. Under the no-wrap hypothesis, every  $d\Delta$  lies in  $[\Delta, \pi]$ , whose minimum is  $\Delta$ . For  $\Delta = 2\pi/Q$  and  $1 \leq d \leq K-1 \leq Q-2$ , every nonzero grid distance is at least  $2\pi/Q$ , with equality at  $d = 1$ .  $\square$

**Corollary 7** ( $\epsilon$ -tuned cyclic family). *For  $0 < \epsilon \leq 1/8$ , let*

$$Q_\epsilon := \left\lfloor \frac{\pi}{2 \arcsin(2\epsilon)} \right\rfloor, \quad K_\epsilon := Q_\epsilon - 1, \quad \Delta_\epsilon := \frac{2\pi}{Q_\epsilon}. \quad (36)$$

*Then*

$$\frac{1}{4\epsilon} \leq K_\epsilon \leq \frac{\pi}{4\epsilon}, \quad \min_{x \neq y} d_{\text{proj}}(U_x, U_y) = 2 \sin\left(\frac{\pi}{2Q_\epsilon}\right) \geq 4\epsilon. \quad (37)$$

*Thus, with  $\mathcal{G}_{m,\epsilon} := \{U_x : x \in [K_\epsilon]_0^m\}$ ,*

$$\log_2 |\mathcal{G}_{m,\epsilon}| = m \log_2 K_\epsilon = \Theta\left(m \log \frac{1}{\epsilon}\right). \quad (38)$$

*The allowed coefficient interval is  $[0, (Q_\epsilon - 2)2\pi/Q_\epsilon]$ . It can extend beyond  $\pi$  but omits the residue  $Q_\epsilon - 1$ ; equation (35), rather than a principal-interval argument, controls its wraparound.*

*Proof.* The bound  $\arcsin(2\epsilon) \leq \pi\epsilon$  and  $\lfloor a \rfloor \geq a - 1$  give  $K_\epsilon \geq 1/(2\epsilon) - 2 \geq 1/(4\epsilon)$ ; the reverse bound uses  $\arcsin(2\epsilon) \geq 2\epsilon$ . Since  $Q_\epsilon \leq \pi/[2 \arcsin(2\epsilon)]$ ,  $\pi/(2Q_\epsilon) \geq \arcsin(2\epsilon)$ . Apply theorem 3. The cardinality statement is immediate.  $\square$

## B. Dispersed additive-update representation

**Definition 20** ( $r$ -round dispersed update stream). *The cyclic specialization of definition 13 with modulus  $Q \geq 3$ ,  $K = Q - 1$ ,  $\Delta = 2\pi/Q$ ,  $r \geq 2$  rounds, and  $\Sigma = \{+1\}$ : for  $x \in [K]_0^m$  the shares  $a^{(1)}, \dots, a^{(r)} \in \mathbb{Z}_Q^m$  satisfy*

$$\sum_{t=1}^r a_j^{(t)} = x_j \pmod{Q} \quad (j \in [m]), \quad (39)$$

*the flat input is the round-major record sequence  $\text{Update}(t, j, a_j^{(t)})$ , and the semantic input is the addressed aggregate vector  $x$ .*

**Proposition 9** (Semantic equivalence and token privacy). *Every valid flat stream of definition 20 implements the same  $U_x$  as the semantic input, up to global phase. Moreover, a single flat token never determines  $x_j$ : for every fixed  $(t, j, a)$  and every  $b \in [K]_0$ , there is a valid stream with  $a_j^{(t)} = a$  and  $x_j = b$ . Under the canonical sharing distribution in which the first  $r-1$  shares are independent uniform residues and the last is fixed by (39), every individual residue is uniform on  $\mathbb{Z}_Q$  and independent of  $x_j$ .*

*Proof.* For integer representatives there are  $k_j \in \mathbb{Z}$  with  $\sum_t a_j^{(t)} = x_j + Qk_j$ . Commutativity and  $Q\Delta = 2\pi$  give

$$\prod_t R_{P_j}(\Delta a_j^{(t)}) = R_{P_j}(\Delta x_j + 2\pi k_j) = (-1)^{k_j} R_{P_j}(\Delta x_j).$$

The product of the signs is global. For privacy, fix the named token and choose any one of the other  $r-1$  residues to enforce the desired sum; the remaining residues are arbitrary. The distributional claim follows because a sum containing at least one independent uniform residue is uniform, including the last share after conditioning on  $x_j$ .  $\square$

**Definition 21** (Forward-pass cut accounting). *Partition the round-major stream into Alice's prefix  $A$  (rounds  $1, \dots, r-1$ ) and Bob's suffix  $B$  (round  $r$ ). A forward pass compiler scans  $AB$  from left to right at most  $p = p(m)$  times. There are  $p$  crossings  $A \rightarrow B$  and at most  $p-1$  rewind crossings  $B \rightarrow A$ . At every crossing, the complete restart state is the five-field serialization of (3); rereadable input and every live IR object are therefore included. For a run on  $I$ , let  $A_p(I)$  be the number of bits irreversibly committed to the final output while the machine scans  $A$ , summed over all passes, and define*

$$\Sigma_p(I) := \sum_{t=1}^p B_{\text{cut}}(c_{A \rightarrow B, t}) + \sum_{t=1}^{p-1} B_{\text{cut}}(c_{B \rightarrow A, t}). \quad (40)$$

*Both are literal, self-delimiting transcript lengths. Unbarred resource symbols take maxima over valid inputs. For randomized compilers, write*

$$\bar{A}_p := \max_I \mathbb{E}_\rho A_p(I, \rho), \quad \bar{\Sigma}_p := \max_I \mathbb{E}_\rho \Sigma_p(I, \rho),$$

and let  $S := \max_I \mathbb{E}_\rho[\max_c B_{\text{cut}}(c; I, \rho)]$  over the  $2p - 1$  named crossings. Then  $\bar{\Sigma}_p \leq (2p - 1)S$ ; the deterministic version is identical without expectations.

**Theorem 4** (One-pass zero-error mask bound). *Let a deterministic one-pass compiler be correct to  $\epsilon < \sin(\pi/(2Q))$  on every valid stream of definition 20. Then*

$$A_1 + B_{\text{cut}}(c_{A \rightarrow B, 1}) \geq \lceil m \log_2 Q \rceil. \quad (41)$$

*Proof.* For each  $u \in \mathbb{Z}_Q^m$ , choose a canonical prefix whose aggregate is  $u$  (put  $u$  in its first round and zero in the other prefix rounds). The committed prefix and restart state after this prefix must distinguish all  $Q^m$  choices. Indeed, for  $u \neq u'$ , choose one common final share  $v$  coordinatewise. At each coordinate at most two residues of  $v_j$  would make either  $u_j + v_j$  or  $u'_j + v_j$  equal to the single excluded aggregate  $Q - 1$ ; since  $Q \geq 3$ , an allowed  $v_j$  exists. Both completions are valid, and they differ in every coordinate in which  $u$  and  $u'$  differ. By theorem 3, one circuit cannot be within  $\epsilon$  of both targets. The messages are prefix-free under the declared serializers, so Kraft's inequality gives (41).  $\square$

### C. Randomized multipass frontier

Put

$$L_\delta(m, K) := (1 - \delta)m \log_2 K - h_2(\delta), \quad 0 \leq \delta < \frac{1}{2}. \quad (42)$$

**Theorem 5** ( $K$ -ary randomized multipass recoverability frontier). *Let  $Q \geq 3$ ,  $K = Q - 1$ , and  $0 \leq \epsilon < \sin(\pi/(2Q))$ . Let a possibly randomized compiler be correct within projective error  $\epsilon$ , with probability at least  $1 - \delta$  for every  $x \in [K]_0^m$  and every valid  $r$ -round dispersed representation of  $x$ . Assume that it makes at most  $p = p(m)$  forward passes in the charged model of definition 21. Then:*

$$\bar{B}_{\text{rel}}^{\text{out}} \geq L_\delta(m, K), \quad \bar{B}_{\text{self}}^{\text{out}} \geq L_\delta(m, K), \quad (43)$$

$$\bar{A}_p + \bar{\Sigma}_p \geq L_\delta(m, K). \quad (44)$$

In particular, for the worst-input expected crossing cap  $S$  of definition 21,

$$\boxed{\bar{A}_p + (2p - 1)S \geq (1 - \delta)m \log_2 K - h_2(\delta).} \quad (45)$$

Here  $S$  includes RAM, parameters, cache/hash metadata, materialized input, and all IR fields from (3); it is not a RAM-only bound. For the tuned choice of corollary 7, both right-hand sides are  $\Omega((1 - \delta)m \log(1/\epsilon))$  for fixed  $\delta < 1/2$ .

*Proof.* For (43), take  $X$  uniform on  $[K]_0^m$  and fix one valid dispersed representation for each value of  $X$ . The packing separation is greater than  $2\epsilon$ , so nearest-packing decoding recovers  $X$  from a successful output. Fano's inequality gives

$$I(X; C) \geq m \log_2 K - h_2(\delta) - \delta \log_2(K^m - 1) \geq L_\delta(m, K).$$

Expected length in either separately fixed prefix-free output code is at least this mutual information; a worst-input expectation is at least the uniform average.

For (44), independently choose a uniform prefix aggregate  $U \in \mathbb{Z}_Q^m$ , put it in the first prefix round, put zero in any other prefix rounds, and set the last share  $V = X - U \pmod{Q}$ . Then  $V$  is uniform and independent of  $X$ , so  $H(X | V) = m \log_2 K$ . Convert the  $p$  scans into the standard alternating simulation across the  $A | B$  partition. The transcript  $T$  contains each complete crossing snapshot and all output chunks written on Alice's side. Given  $T$  and  $V$ , Bob can reproduce every Bob-side segment and the complete final output; random-tape state and dynamic control are present in the crossing serialization. The output decoder therefore recovers  $X$  with error at most  $\delta$ , and conditional Fano yields

$$I(X; T | V) \geq L_\delta(m, K).$$

The framed transcript is prefix-free conditional on  $V$ , so its expected literal length is at least this information. Its length is precisely  $A_p + \Sigma_p$  up to fields already included in the self-delimiting serializers. Maximizing over inputs proves (44); the cap and tuned- $\epsilon$  consequences follow.  $\square$

**Remark 14** (What is and is not optimal). *Equation (43) is an output-information bound and (45) is a transcript bound; neither converts self-contained circuit bits into gates. The factor  $2p - 1$  counts the actual alternating crossings of a forward scan. A random-access input oracle, an uncharged rereadable file, or reverse scans would be a different model. For  $p = 1$  and zero error, theorem 4 strengthens  $m \log_2 K$  to  $m \log_2(K + 1)$ . No such strengthening is claimed for randomized or general multipass protocols.*

### D. Memory-capped constructions

**Proposition 10** (Block hybrid). *Fix  $p, h \geq 1$  and let  $q = \min\{m, ph\}$ . There is a deterministic zero-error  $p$ -pass compiler for definition 20 that aggregates  $q$  coordinates and passes the other update records through in the public  $p$ -block hybrid mode of  $E_{\text{rel}}$ . It obeys*

$$S \leq \left\lceil \frac{q}{p} \right\rceil \log_2 Q + O(\log(mQrp)), \quad (46)$$

$$A_p \leq (r - 1)(m - q) \log_2 Q + O(r \log(mQrp)), \quad (47)$$

$$B_{\text{rel}}^{\text{out}} \leq q \log_2 K + r(m - q) \log_2 Q + O(r \log(mQrp)). \quad (48)$$

In the compact-output corner  $h = \lceil m/p \rceil$ , this becomes

$$B_{\text{rel}}^{\text{out}} \leq m \log_2 K + O(\log(mQrp)), \quad (49)$$

$$S \leq \left\lceil \frac{m}{p} \right\rceil \log_2 Q + O(\log(mQrp)). \quad (50)$$

For two rounds, the construction satisfies  $A_p + pS \leq m \log_2 Q + O(p \log(mQrp))$  after using a smaller final

block. Compared with (45) at  $\delta = 0$ , it is within the constant  $(2p - 1) \log_2(K + 1)/(p \log_2 K) < 2 \log_2 3$  for every  $K \geq 2$ .

*Proof.* Partition the first  $q$  coordinates into at most  $p$  consecutive blocks of size at most  $h$ . On pass  $t$ , keep only the residues of block  $t$  and add their first  $r - 1$  rounds modulo  $Q$ ; when the final-round mates arrive, emit the corresponding aggregate base- $K$  block. The live packed block, indices, and framing give (46). On the first pass, emit every share of each unselected coordinate in the declared round-major order; these rotations already have the correct product and require no retained coefficient table. The selected coordinates can be balanced across the  $p$  blocks, so the maximum block has size  $\lceil q/p \rceil$ . Their consecutive boundaries are a public function of  $(m, p, q)$ , so the header stores those integers rather than  $p$  separate delimiters. The prefix contains  $r - 1$  residues and the whole output contains  $r$  residues for each such coordinate, giving (47) and (48). If  $ph \geq m$  there are no passthrough coordinates. For  $r = 2$ ,  $(m - q) \log_2 Q + ph \log_2 Q$  is at most  $m \log_2 Q$  up to the last partial block and framing. The constant comparison uses  $(2p - 1)/p < 2$  and  $\log_2(K + 1)/\log_2 K \leq \log_2 3$ .  $\square$

**Proposition 11** (Semantic and randomized upper corners). *Under the same public dictionary, a one-pass semantic compiler receiving  $x$  emits the packed base- $K$  aggregate code with*

$$B_{\text{rel}}^{\text{out}} \leq \lceil m \log_2 K \rceil + O(\log(mK)), \quad S = O(\log(mK)), \quad (51)$$

and implements the same  $U_x$  as the flat stream. For any  $0 \leq \delta < 1/2$ , an input-independent abort coin that runs this exact compiler with probability  $1 - \delta$  and otherwise emits a fixed short circuit has worst-input success at least  $1 - \delta$  and expected relative output

$$(1 - \delta) \lceil m \log_2 K \rceil + O(\log(mK) + 1). \quad (52)$$

Thus the randomized family-relative output lower bound is matched up to framing and the binary-entropy additive term. A self-contained  $R_Z$ -circuit upper bound remains  $O(m(\log m + \log Q))$  bits and at most  $m$  rotations; no matching self-contained-bit or gate-count claim is made.

*Proof.* Pack the ordered aggregate vector as one base- $K$  integer and retain only its streaming index and framing state. Proposition 9 gives semantic equality with the flat representation. Prefix the output with the abort-coin branch tag for the randomized construction and take expectations. The self-contained construction writes each nonzero  $R_{P_j}(2\pi x_j/Q)$  with its address and rational parameter.  $\square$

## VII. REPRESENTATION-AWARE COMPILER

### A. Bounded Supported-Pattern Recognition

**Definition 22** (Supported-pattern region recognizer). *The implemented recognizer accepts only configured high-level forms: (i) explicit Pauli rotations or phase-gadget records that already carry their support, (ii) local phase-ladder/Fourier spans, (iii) expanded mirrored or self-inverse composite blocks, and (iv) repeated-run or preset semantic nodes. Within one supported contiguous span it maintains the declared generators  $P_1, \dots, P_k$ , coefficients  $\theta_1, \dots, \theta_k$ , qubit support, and source records. It may extend the span while the next record belongs to the supported grammar and passes the configured commutation check; an unsupported or certified noncommuting record closes the candidate. Rejected recognition leaves the original circuit available to the conservative fallback.*

**Remark 15** (Recognition, maximality, soundness, and complexity scope). *The artifact does not implement a general binary-symplectic extractor for arbitrary circuit DAGs and does not claim a maximal commuting region in that sense. Its only maximality notion is the maximal contiguous span accepted by the fixed supported-pattern grammar under the chosen scan order. Different generators remain separate records until aggregation. Soundness of a selected rewrite comes from the independent exact or symbolic equivalence certificate of definition 24; failure returns the unchanged candidate. No general worst-case linear-time claim is made: the runtime is the sum of the bounded scans, signature/hash bookkeeping, certificate work, and any configured backend probes, as detailed in Appendix B.*

### B. Aggregation and Certificate

**Definition 23** (Aggregation of an accepted supported region). *Given an accepted supported region with generators  $P_1, \dots, P_k$  and coefficients  $\theta_1, \dots, \theta_k$ , the aggregation emits a representative circuit implementing  $\exp(-i \sum_j \theta_j P_j/2)$  using at most  $c_0 + c_1 k$  gates in the target basis.*

Let

$$\mathcal{C}_n = \langle H, S, S^\dagger, X, Y, Z, CX, CZ, \text{SWAP} \rangle$$

be the Clifford gate language on addressed qubits. Fix a declared finite formal-symbol set  $\mathcal{S}$  and the exact coefficient module

$$\Lambda = \mathbb{Q}\pi \oplus \bigoplus_{s \in \mathcal{S}} \mathbb{Q}\theta_s.$$

For a signed Hermitian  $n$ -qubit Pauli  $P$  and  $\lambda \in \Lambda$ , write  $R_P(\lambda) = \exp(-i\lambda P/2)$ .

**Definition 24** (Exact commuting-Pauli/Clifford certificate domain). *The domain  $\mathcal{D}_{\text{CPC}}$  consists of finite addressed circuits over  $\mathcal{C}_n$ , exact records  $R_Z(\lambda)$  and  $R_P(\lambda)$ , identity and barrier records, and an arbitrary exact global phase. Every coefficient is a reduced rational  $\pi$  component plus a sorted sparse list of reduced rational formal-symbol components carried by the input; a floating-only parameter is not in the exact domain.*

*Scan a circuit in execution order while maintaining a Clifford frame  $F$  and the invariant*

$$U_{\text{prefix}} = FR_{P_t}(\lambda_t) \cdots R_{P_1}(\lambda_1) \quad \text{up to global phase.} \quad (53)$$

*The frame is represented by the complete signed inverse-conjugation tableau  $\{F^\dagger X_j F, F^\dagger Z_j F\}_{j=1}^n$ . Appending a physical rotation  $R_Q(\lambda)$  appends  $F^\dagger Q F$  to the Pauli list. The certificate predicate additionally requires all extracted  $P_j$  to commute pairwise.*

*Canonicalization moves a Pauli sign into its coefficient, discards an identity-axis rotation, adds coefficients on identical unsigned supports, reduces the rational  $\pi$  coefficient modulo 2 to the interval  $[0, 2)$ , keeps every formal-symbol coefficient exact, deletes zero records, and sorts the remaining records lexicographically by their binary  $(x, z)$  masks. Circuit and Clifford global phases are omitted. The resulting record is*

$$\begin{aligned} \beta_j &:= (b_{j,s})_{s \in \mathcal{S}} \in \mathbb{Q}^{\mathcal{S}}, & \bar{a}_j &\in \mathbb{Q} \cap [0, 2), \\ \text{can}(C) &= (n, T_F, ((x_j, z_j, \bar{a}_j, \beta_j))_{j=1}^\ell). \end{aligned} \quad (54)$$

*The symbolic checker accepts exactly when the two canonical byte records are equal.*

**Proposition 12** (Soundness of canonical symbolic acceptance). *For  $C, C' \in \mathcal{D}_{\text{CPC}}$  satisfying the pairwise-commutation predicate,*

$$\text{can}(C) = \text{can}(C') \implies U_C = e^{i\phi} U_{C'}$$

*for some  $\phi \in \mathbb{R}$ .*

*Proof.* Induction over the input records proves (53). For a Clifford record  $G$ , replace  $F$  by  $GF$ . For a rotation,

$$R_Q(\lambda)F = FR_{F^\dagger Q F}(\lambda),$$

which is exactly the tableau update used by the checker. Thus the extracted record implements the original circuit up to its explicitly ignored global phase.

Pairwise commutation permits deterministic reordering without changing the unitary. Sign normalization is exact because  $R_{-P}(\lambda) = R_P(-\lambda)$ , while coefficient reduction changes a factor only by

$$R_P(\lambda + 2k\pi) = (-1)^k R_P(\lambda).$$

All such signs multiply to one input-independent global phase; an identity-axis rotation is also a global phase. Equality of the sorted rotation records, including every formal-symbol component, therefore gives equal rotation products projectively.

Finally, equality of the signed tableaus means that  $F'F^\dagger$  commutes with every Pauli generator. The Pauli representation generates the full matrix algebra, so  $F'F^\dagger$  is a scalar. Hence the two Clifford frames also agree up to global phase, proving the claim.  $\square$

**Remark 16** (Status and completeness boundary). *Only status `completed_valid`, accompanied by the content-addressed canonical records and an independent certificate payload, is an acceptance. A canonical mismatch is `completed_invalid`; an unlisted gate or malformed address is `unsupported`; failure of pairwise commutation is `predicate_error`; and a floating coefficient without an exact sidecar, an unassigned formal symbol in the dense checker, or a dense distance inside the numerical guard band is `numerical_inconclusive`. Only `completed_valid` rows enter quality means. The proposition is a soundness statement, not a claim that every equivalent Clifford-plus-rotation circuit has the same record.*

For  $n \leq 6$ , the audit independently forms dense matrices and checks the normalized Frobenius projective distance

$$d_{\text{F,proj}}(U, V) = 2^{-n/2} \min_{|\omega|=1} \|U - \omega V\|_{\text{F}}.$$

The frozen suite covers both accepted and rejected pairs at every width  $1 \leq n \leq 6$  and changes, separately, an angle, a sign, one support bit, a qubit permutation, the Clifford frame, and a deleted or duplicated gate. The executable definition and status contract are `certificates/CERTIFICATE_SPEC.md`; the canonicalizer and checker are `certificates/canonicalize.py` and `certificates/symbolic_checker.py`.

Within the supported exact domain, the executable audit applies 56 adversarial mutations—angle, sign, support bit, qubit index/permutation, Clifford frame, deletion, and duplication—and observes 0 false accepts. It also performs 12 dense/symbolic cross-check cases for  $n \leq 6$  with no checker disagreement. The status contract distinguishes `completed_valid`, `completed_invalid`, `unsupported`, `predicate_error`, and `numerical_inconclusive`; only independently certified `completed_valid` rows enter quality means.

### C. Constructive Aggregation Proposition

**Proposition 13** (Constructive aggregation size). *For a supported semantic input accepted as a commuting diagonal phase layer with  $m$  independent phase terms, the semantic aggregation path emits an aggregate representative with at most  $c_0 + c_1 m$  phase records before final lowering, for constants  $c_0, c_1$  fixed by the representation and lowering conventions. Hence the constructive representative has size  $O(m)$  and is independent of the repetition count  $r$  at fixed width.*

*Proof.* By hypothesis the accepted layer is a product of  $m$  supported commuting diagonal terms. The constructive path keys terms by their diagonal generator and adds coefficients across repetitions. Repetition changes the coefficient on each key, not the number of keys. The aggregate contains at most one record per phase term, plus fixed boundary and lowering overhead.  $\square$

Proposition 13 shows that a semantic-first pipeline achieves an  $O(m)$ -record representative on the packing family. Under the public family dictionary it has  $B_{\text{rel}} = m + O(\log m)$ , matching the family-relative inequality of theorem 1; its separately reported self-contained size is  $O(m \log m)$ . On direct semantic input, the support mask can be streamed with  $B_{\text{cut}} = O(\log m)$ , which is the constructive side of the separation in theorem 2.

## VIII. CONTROLLED VALIDATION

### A. Matched-Representation Benchmark

**Definition 25** (Matched-representation benchmark). *For one sampled coefficient vector  $x$ , a matched cell fixes the same target  $U_x$ , target basis, hardware model, total error budget, timeout, memory cap, and seed, then crosses every compiler with nine representations: (1) semantic aggregate, (2) contiguous flat, (3) round-robin flat, (4) random commuting order, (5) masked-share/update stream, (6) locally folded flat, and (7–9) QASM, TKET, and ZX bridge records. Each representation has a unique identifier, serialized payload, generation seed, and unitary certificate. A cell is numerical only after the common certificate policy returns `completed_valid`; unsupported predicates and malformed or inconclusive records remain status outcomes.*

The executable campaign uses the no-wrap target

$$U_x = \prod_{j=1}^m \exp\left(-\frac{i\pi x_j}{8K} Z_j Z_{j+1}\right), \quad x_j \in \{0, \dots, K-1\}.$$

The active support is sampled at the requested density and each active coefficient is uniform on  $\{1, \dots, K-1\}$ . The frozen campaign contains 1,080 matched cells across the nine representations and five compilers. Of these, 1,065 are `completed_valid`; the remaining 15 are retained as `predicate_error`, all from the Qiskit masked-update path, and are not converted into numerical wins or losses. Every raw row records the generation parameters, sampled  $x$ , realized density, schedule, basis, timeout, memory cap, runtime, peak RSS, gates, depth, CX, serialized input and output/IR bytes, bridge sequence, and certificate kind.

The representation main effect, compiler main effect, their interaction, completion probability, and scaling exponents are computed from the complete matrix. The interaction is substantive: a compiler that reconstructs a phase polynomial or Pauli network before synthesis

belongs on the semantic-capable side even when its input arrived through a flat bridge.

### B. Semantic-Capable External Baselines and Bridge Loss

A separate 1,728-cell campaign freezes recommended-native and unified-target-basis configurations for the authors' exact phase-table reference, local Qiskit optimization, TKET PauliSimp, and PyZX full-reduce. It records tool version, available commit metadata, pass sequence, bridge, seed, basis, timeout, memory cap, pre/post-bridge gate/depth/CX and serialized IR, runtime, RSS, and certificate status. 1,695 cells are `completed_valid`; 33 are retained as `predicate_error`. TKET and PyZX are classified as semantic-capable, not as local baselines, whenever their configured global IR recovers the compact structure. Bridge loss is reported separately from compiler quality, and unsupported is never conflated with incorrect.

The stance of this campaign is explanatory, not competitive. TKET PauliSimp and PyZX full-reduce already implement the semantic-first principle on their global IRs; the contribution of this paper is the bound that says what any such tool must pay to do so on a dispersed stream (corollary 4). A preregistered scan tested this on a third-party compiler the authors do not control: with predictions and decision rule committed before execution (`PaperDraft/research/E7_preregistration.md`), TKET PauliSimp was run on the masked-share family at  $m \in \{8, 16, 32, 64, 128\}$ , three seeds per size, in the byte-identical pass sequence of this campaign. It achieved compact output in every cell (508 dispersed rotations aggregated to 127 non-Clifford  $R_z$  at  $m = 128$ ), and its measured pass-phase memory and serialized IR grew with log-log slope  $\approx 1.7$  in  $m$ —at least the linear growth that the  $\Omega(m \log K_e)$  floor requires, and in fact more—while the instrumented semantic-input control stayed flat (196  $\rightarrow$  199 bytes over the same  $16\times$  range of  $m$ ). The preregistered point prediction of *exactly* linear growth was thereby rejected under its own frozen decision rule, and we report it as rejected; the one-sided lower bound is consistent with, and slack against, the superlinear measurement. The frozen artifact is `data/frozen/e7_tket_state_scaling.json`. Two limits on what this scan is allowed to establish should be stated explicitly. First, the quantity measured at a third-party API boundary is a pass-phase memory and serialized-IR snapshot, which is a configured-pipeline diagnostic in the sense of table IV: it is not promoted to a measurement of  $B_{\text{cut}}$  and is not commensurable with the  $S$  of the frontier. The scan therefore tests only the qualitative dichotomy the theorem predicts—dispersed input forces growth in  $m$ , semantic input does not—and tests neither the constant nor the exponent. Second, superlinear growth means the floor is not the binding constraint on this tool: the bound establishes the floor, the observed cost exceeds it, and the excess is implementation structure




FIG. 4. Complete matched-representation interaction summary. Panel (a) reports certified completion probability for every representation–compiler pair; panel (b) reports serialized output size only among completed cells. Predicate errors remain status observations and never become numerical wins.

(the pairwise commutation structure of a Pauli graph is itself superlinear in the number of terms). No claim that the authors’ implementation outperforms these tools is made or implied.

**Lemma 10** (Dimension-preserving instrumentation). *For every accepted version-one artifact  $A$ , its reported size is exactly  $8 \text{stat}(A).st\_size$  bits. Address growth is charged by the literal varint bytes; in particular 127 and 128 have different address lengths. There is no representation-independent conversion from gate count, depth,  $cx$ , token count, or IR-node count to bits.*

*Proof.* The first statement is proposition 1. Gate or node count does not determine parameter numerators, denominators, qubit-address widths, edge endpoints, delimiters, or textual external payloads. Thus equal-count objects can have different artifact lengths. Conversely, family-relative dictionary codes can assign different self-contained circuits equal conditional length. Only the named artifact file supplies a bit value.  $\square$

### C. Packing-Family Scaling

The complete multifactor campaign of figure 4 is the primary matched-representation experiment. The earlier 5-by-5 sweep below is retained as a separately keyed legacy diagnostic because it reaches the same maximum width under a 600s protocol and preserves its historical `predicate_error` outcomes; it is not substituted for the new matrix. This legacy experiment links the growing- $m$  theorems to configured-pipeline behavior. Unlike the

fixed-width witnesses below, which exercise the exact-semantics corollary, the packing family grows the number of independent phase terms  $m$  at fixed tolerance.

*a. Instance.* We instantiate definition 14 with the chain of two-body characters  $a_j = e_j + e_{j+1}$  on  $n = m + 1$  qubits, which are  $\mathbb{F}_2$ -independent for every  $m$  (certified by rerunning the exact rank verifier), using the separate experimental setting  $K = 2$ ,  $\Delta = \pi/4$ . Each token is lowered as  $RZ_i(\theta) RZ_j(\theta) CP_{ij}(-2\theta)$  under the convention used here. The balanced round-robin encoding (definition 15) with dispersal  $r = 4$  interleaves the terms across shared wires, so per-wire locality does not trivially reveal the aggregate. Because these balanced tokens contain  $x_j/r$ , this is a pipeline diagnostic rather than an instance of the masked-share lower bound. The sweep is  $m \in \{4, 8, 16, 32, 64\}$  under the standard 600s protocol of this section. The exact rational- $\pi$  sidecar records  $\Delta/\pi = 1/4$  and per-token angle  $1/16$ . Correctness uses the dense certificate for  $n \leq 6$  and the sound exact symbolic checker of definition 24 and proposition 12 for  $n > 6$ ; any output outside  $\mathcal{D}_{CPC}$  is status-only.

*b. Results.* Table V reports completed output gates against  $m$  for semantic UCC, `qiskit opt3`, rebased TKET PauliSimp, staq rotation folding, and the phase-polynomial reference. All five paths completed at  $m = 4$ ; semantic UCC, `qiskit opt3`, staq rotation folding, and the phase-polynomial reference also completed at  $m = 8, 16, 32, 64$ , while rebased TKET PauliSimp returned `predicate_error` at those four sizes because its output fell outside the Hadamard-free symbolic-certificate predicate. Semantic UCC and the phase-polynomial reference produced the same linear resource sequence, from 25/21/8 gates/depth/ $cx$  at  $m = 4$  to 385/321/128 at  $m = 64$ .

| Dimension      | Recorded field                     | Executable serializer or definition                                                                    | Permitted comparison               |
|----------------|------------------------------------|--------------------------------------------------------------------------------------------------------|------------------------------------|
| bits           | <code>serialized_input_bits</code> | verified <code>SemanticStream</code> , <code>FlatUpdateStream</code> , or external-input artifact file | same codec/mode only               |
| bits           | <code>B_com_bits</code>            | verified literal committed-output artifact/prefix file                                                 | same codec only                    |
| bits           | five <code>B_*_bits</code> fields  | verified control/store/window/parameter/IR files                                                       | corresponding formal field only    |
| records/tokens | source records                     | exact count before serialization                                                                       | never converted by a constant      |
| IR nodes       | <code>ir_nodes</code>              | structural node count reported separately from <code>B_ir_bits</code>                                  | no implicit bit or gate conversion |
| gates/depth/CX | structural output                  | Qiskit circuit operations, DAG depth, and <code>cx</code> count                                        | like-for-like target basis         |

TABLE III. Instrumentation contract. Figures and tables compare within rows of one dimension. Cross-dimensional discussion cites the explicit codec or is stated only qualitatively.

| Evidence class                            | Included experiments                                                                          | Authorized inference                                                                   |
|-------------------------------------------|-----------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|
| instrumented theorem-model implementation | W3 traced reference compilers and restart verifier                                            | directly tests the serialized commitment-cut coordinates and reports theorem residuals |
| configured-pipeline diagnostic            | matched matrix, external semantic baselines, exact-semantics witnesses, and natural workloads | compares observed tool behavior; never proves a lower bound                            |
| model-specific consequence                | fixed-total-error synthesis and surface-code calculation                                      | reports consequences under the disclosed QRE; not a hardware-independent forecast      |

TABLE IV. Formal/evidentiary separation used throughout the paper. The first row is the instrumented reference implementation of the declared resource model; configured external tools are not retroactively placed in that class.

Among the completed external rows, `qiskit opt3` and `staq` produced  $17m$  and  $10m$  output gates, respectively, as the  $rm$  token stream grew. Status outcomes are recorded separately and omitted from numerical curves.

The earlier two-panel figure is deliberately omitted. Its gate panel is redundant with table V, while its family-relative bit panel was generated from a handwritten length formula and has no serialized per-point files. Under the present accounting contract such a curve may be restored only by emitting the exact conditional-code artifacts (together with the family-dictionary digest) and measuring their files.

*c. Scope.* The external-toolchain rows are configured-pipeline diagnostics in the sense of this section: they neither prove nor test theorem 2, whose serialized inter-round cut budget is not observable from output gate counts, and no bit curve is inferred from these gate-count rows. The separate reference-compiler implementation below measures theorem-model cut fields directly; it is not inferred from output gates.

#### D. Event-Level Commitment–Cut Tracing

The instrumented suite contains echo/pass-through, full aggregation, 25%, 50%, and 75% memory-capped hybrids, and an IR-materializing control. Each compiler follows the formal flat-update interface and records an event after every input update and commitment. The campaign scans  $m, r, K, \epsilon, p, S$ , schedule, and seed, producing 592 complete run summaries plus the compressed raw event log and content-addressed trace archive. At

each named cut it records input position, pass, control/store/window/parameter and IR bytes, committed-output bytes, current output length, RSS, and wall time. The literal measured definitions are

$$B_{\text{com}}(c) = 8|A_{\text{com},c}|_{\text{byte}},$$

$$B_{\text{cross}}(c) = 8 \sum_{f \neq \text{com}} |A_{f,c}|_{\text{byte}}.$$

RSS is recorded as an implementation metric and is not substituted for the serialized state cap  $S$ . The analyzer recomputes the theorem residual from the trace coordinates and retains every anomaly or failed run. All frozen points lie in the theorem-allowed region; no point was removed as an outlier.

Echo commits updates early with little retained state; full aggregation occupies the state-heavy compact-output corner; and the three capped hybrids interpolate through explicit block commitments. Their scaling fits and residuals are reported with the raw points, making this a direct measurement of the theorem resources rather than an inference from output gate count.

#### E. Width-Axis Scaling

To check the constructive aggregation prediction (proposition 13), we run a width-axis sweep for the `full_pair_cp` topology at  $n = 4, 5, 6$ .

The output is  $O(m_{\text{CP}})$  and independent of  $r$ , confirming the constructive proposition.



| $m$ | semantic UCC | qiskit opt3  | TKET PauliSimp (rebased) | staq folding | phase-poly ref |
|-----|--------------|--------------|--------------------------|--------------|----------------|
| 4   | 25/21/8      | 68/38/32     | 54/20/10                 | 40/27/32     | 25/21/8        |
| 8   | 49/41/16     | 136/54/64    | predicate_error          | 80/39/64     | 49/41/16       |
| 16  | 97/81/32     | 272/86/128   | predicate_error          | 160/63/128   | 97/81/32       |
| 32  | 193/161/64   | 544/150/256  | predicate_error          | 320/111/256  | 193/161/64     |
| 64  | 385/321/128  | 1088/278/512 | predicate_error          | 640/207/512  | 385/321/128    |

TABLE V. Packing-family scaling at growing  $m$  (chain two-body characters,  $n = m + 1$ , dispersal  $r = 4$ , 600 s budget). Completed cells report output gates/depth/cx; status outcomes are recorded separately and are not scored as structural wins. No external row is claimed to be a member of the transducer class of definition 9.




FIG. 5. Trace-derived commitment-crossing-state-IR surface. The black edge marks measured Pareto points and the theoretical lower envelope is evaluated in the same serialized-bit coordinates. Marker area in the raw slice reflects IR share; RSS is recorded separately.

## F. Exact-Semantics Witness Results

These experiments exercise the exact-semantics corollary (corollary 1), not the main resource tradeoff (theorem 2). They confirm that the constructive aggregation path works on fixed-width witnesses and that configured external pipelines behave predictably across the representation boundary. Full data tables are in section C.

The rational CP  $HD^rH$  witness is the canonical fixed-width family with 42-gate constant output. Figure 6 shows that semantic UCC remains at 42 gates while rebased TKET PauliSimp returns 52–73 gates across five scales; materialize-first baselines grow linearly. The irrational independent-Pauli witness (figure 7) shows the same pattern at 27-gate constant output. Key observations:

- Semantic UCC outputs a constant gate count (42 for rational, 27 for irrational) at all five tested sizes.
- Rebased TKET PauliSimp recovers a small validated form at all five scales on both witnesses, demonstrating that a generic pipeline can cross the formal boundary by reconstructing an aggregate Pauli representation.
- Public staq rotation folding, qiskit opt3, and TKET GuidedPauliSimp grow linearly with input size.




FIG. 6. Fixed-width non-inverse Fourier-layer scaling on the rational  $HD^rH$  witness. Semantic UCC remains at 42 gates and rebased TKET PauliSimp returns 52–73 gates across the five scales. Public staq rotation folding, qiskit opt3, and TKET GuidedPauliSimp grow with input size; PyZX full\_reduce completes at 4k and reaches the time budget at larger scales. Status-only cells are omitted from numerical curves.

- PyZX full\_reduce completes at small sizes but reaches time or correctness limits at larger scales.

| Width ( $n$ ) | CP pairs ( $m_{CP}$ ) | Output Gates | CX Count |
|---------------|-----------------------|--------------|----------|
| 4             | 6                     | 42           | 12       |
| 5             | 10                    | 65           | 20       |
| 6             | 15                    | 93           | 30       |

TABLE VI. Width-axis scaling for the semantic Fourier-layer witness. For each fixed width  $n$ , the semantic output is constant across all tested repetition counts  $r$ . Across widths, the output scales as  $O(m_{CP})$  where  $m_{CP} = n(n-1)/2$ .




FIG. 7. Configured-pipeline scaling on the irrational independent-Pauli witness. Semantic UCC remains at 27 gates. Rebased TKET PauliSimp returns a small validated form at all five scales, demonstrating that a generic pipeline can cross the formal boundary by reconstructing an aggregate Pauli representation. Public staq rotation folding, `qiskit opt3`, and TKET GuidedPauliSimp grow with input size. PyZX `full_reduce` completes at 4k and 10k, fails the numerical tolerance at 20k, and reaches the time budget at larger scales. Status-only cells are omitted from numerical curves.

These measurements exercise the constructive side of corollary 1 and disclose the behavior of configured tools; they do not prove the resource bound for any real compiler. The rebased PauliSimp result is the classification’s positive prediction in action: once a generic pass sequence reconstructs a compact Pauli representation, the output need not grow with the flat input. Conversely, staq’s growing output and PyZX’s status outcomes are not impossibility results—they are diagnostics of configured-pipeline behavior on this specific witness family.

### G. Mechanism Ablation

The ablation strengthens the mechanism interpretation. With Fourier-layer IR disabled, the compiler completes at 4k, 10k, and 20k but materializes the same large non-aggregated outputs as `qiskit opt3`; at 50k and 100k the disabled-IR rows exceed the 600s budget (figure 8). In contrast, the Fourier-enabled semantic UCC path and the phase-polynomial reference recover the same 42-gate canonical circuit at every tested size. Thus the constant form is tied to semantic Fourier-layer lifting and coefficient aggregation, not merely to a favorable ordering of flat post-lowering passes.




FIG. 8. Mechanism ablation for the fixed-width  $HD^rH$  witness. Disabling the Fourier-layer semantic path yields completed but linearly large outputs through 20k and timeout rows at 50k and 100k, while Fourier-enabled semantic UCC and the phase-polynomial reference recover the canonical 42-gate circuit.

### H. Fixed-Total-Error End-to-End Synthesis and QRE

The new campaign scans  $\epsilon_{\text{total}} \in \{10^{-3}, 10^{-6}, 10^{-9}\}$  and assigns the same additive fractions to every pipeline:  $0.10\epsilon_{\text{total}}$  for the certified algorithmic stage,  $0.45\epsilon_{\text{total}}$  for rotation synthesis, and  $0.45\epsilon_{\text{total}}$  for logical failure. For equal allocation and  $N_{\text{rot}}$  output rotations, it chooses  $p = \lceil \log_{10}(N_{\text{rot}}/\epsilon_{\text{synth}}) \rceil$  and  $\epsilon_i = 10^{-p}$ , hence  $\sum_i \epsilon_i \leq \epsilon_{\text{synth}}$ . A second measured-cost greedy allocation spends only this decimal slack; it is an upper construction, not an optimality claim. Each IEEE-serialized angle is actually synthesized by the frozen deterministic GMP build of staq `grid_synth` [8, 31]. Across 670 checked angle-precision pairs, the reported error never exceeds the allocated radius; exact multiples of  $\pi/4$  use staq’s checked exact branch.

The full  $2 \times 3 \times 3 \times 2 \times 2 \times 4 = 288$  matrix crosses two target sizes, three pipelines, three total errors, two allocation policies, two QEC profiles, and  $1/4/16/32$  factories. All 288 rows are `completed_valid`, carry a content-addressed exact-symbolic or dense  $n \leq 6$  certificate, and satisfy the requested total-error bound. The two surface-code profiles [32–34] use  $p_{\text{phys}} = 10^{-3}$  and  $10^{-4}$ , respectively, with separately frozen data/factory coefficients, factory periods [14], and cycle times. They choose the least odd  $d$  meeting the disclosed union bound. These are executable sensitivity scenarios rather than vendor estimates, and follow the accounting conventions of published end-to-end

| Requested | semantic UCC | artifact UCC (Fourier disabled) | qiskit opt3       | phase-poly ref |
|-----------|--------------|---------------------------------|-------------------|----------------|
| 4,000     | 42 / 2.9 s   | 9,588 / 18.2 s                  | 9,588 / 1.0 s     | 42 / 0.2 s     |
| 10,000    | 42 / 8.7 s   | 23,988 / 94.2 s                 | 23,988 / 5.1 s    | 42 / 0.3 s     |
| 20,000    | 42 / 28.5 s  | 47,988 / 339.3 s                | 47,988 / 19.5 s   | 42 / 0.3 s     |
| 50,000    | 42 / 20.7 s  | timeout                         | 119,988 / 120.9 s | 42 / 0.5 s     |
| 100,000   | 42 / 28.7 s  | timeout                         | 239,988 / 485.8 s | 42 / 0.8 s     |

TABLE VII. Fourier-layer semantic-path ablation on the fixed-width  $HD^rH$  witness under the 600 s budget. Disabling the Fourier-layer IR completes at 4k–20k but materializes large outputs; 50k–100k exceed the budget. Fourier-enabled semantic UCC and the phase-polynomial reference return 42 gates at all sizes.

| pipeline                         | $N_{\text{rot}}$ | $T$       | physical qubits | factories | cycles     | qubit-s  |
|----------------------------------|------------------|-----------|-----------------|-----------|------------|----------|
| External semantic-capable (TKET) | 29               | 798       | 258,152         | 16        | 13,317     | 3.44e+03 |
| Materialize-first                | 23,992           | 2,679,110 | 468,968         | 16        | 57,966,528 | 2.72e+07 |
| Semantic-first reference         | 22               | 1,800     | 305,000         | 16        | 31,675     | 9.66e+03 |

TABLE VIII. Predeclared 20k representative slice of the fixed-total-error campaign:  $\epsilon_{\text{total}} = 10^{-6}$ , equal synthesis allocation, conservative  $p_{\text{phys}} = 10^{-3}$  scenario, and 16 factories. Each compiled output has an independent certificate, and every unique-angle staq synthesis returns **Check flag** = 1. The frozen full table additionally records code distance, data and factory qubits separately, and modeled quantum runtime.




FIG. 9. Physical-qubit/runtime Pareto grid at the representative 20k,  $\epsilon_{\text{total}} = 10^{-6}$  slice. Labels give factory counts. The external TKET path receives a materialized CX- $R_Z$ -H bridge but reconstructs a compact phase form before synthesis, so it is classified as semantic-capable.

estimates [12, 35].

The physical point is the ordering of semantic aggregation relative to synthesis. In a fault-tolerant pipeline, arbitrary rotations are eventually approximated by discrete resources such as Clifford+ $T$  sequences. If the compiler first materializes  $r$  separate rotations and only later synthesizes them, the resource estimator sees  $r$  independent approximation tasks. A resource-rich pipeline instead combines the coefficients first and presents one aggregated rotation per phase term. An FT toolchain that reconstructs the phase polynomial before synthesis collapses this multiplicity itself and should be classified on the semantic-capable side of the comparison (see section X for the model-specific resource discussion).

## IX. APPLICATION WORKLOADS

### A. Workload Specification

The application campaign contains seven structured families: Ising/MaxCut cost evolution [1], diagonal Hamiltonian simulation, commuting-block Trotterization [4], QFT arithmetic, QPE controlled powers, controlled-power ladders, and phase-polynomial blocks [18]. Family sizes and naming follow the low-level benchmark conventions of [36, 37]. Random Clifford+ $R_Z$  and routing-dominated [38, 39] circuits are negative controls. The separate controlled packing family (definition 14) with growing  $m$ , fixed  $\epsilon$ , is reported in section VIII C.

### B. Strong Baseline Comparison

The canonical comparison includes Qiskit `optimization_level=3` [29], TKET PauliSimp [17, 40], PyZX `full_reduce` [22], and the authors' exact phase-table reference. Each external tool has a recommended-native and a unified  $\{CX, R_Z, H\}$  configuration, with pass sequence, bridge, version, seed, timeout, memory cap, and certificate frozen. TKET supplies the Pauli-network/global-semantic class and PyZX the graph/global-IR class. The complete W4 interaction is figure 4; the independent W5 campaign and bridge-loss accounting are described in section VIII B. Historical staq, GuidedPauliSimp, and phase-polynomial diagnostic rows remain in their original immutable campaigns and are not silently mixed into the canonical matrix. This comparison is diagnostic, not a contest: the external tools are treated as independent implementations of the semantic-first principle whose internal costs the bounds




FIG. 10. Total output gate counts on the three official real instances. The semantic UCC artifact matches `qiskit opt3`-level output quality while substantially improving over the upstream UCC v0.4.12 pre-semantic default.

of sections IVC and VI explain, gate-count parity is reported as parity, and no headline ranks the authors’ pipeline above them.

### C. Real-Task Panel

These results support the recoverability-parity regime (figure 10): the semantic artifact matches `qiskit opt3`-level output quality while improving over the upstream UCC v0.4.12 pre-semantic default of table IX by approximately 65%, 73%, and 78% gate reduction on the three instances. The selector cost in this panel is the implementation’s conservative structural score  $\kappa = \text{gates} + \text{depth} + 10 \cdot \text{multi-qubit}$ , with candidate-control guards; the factor 10 is a conservative lexicographic-style heuristic favoring lower two-qubit exposure, used only for portfolio selection. This explains the **REAL24-phase-est.** row (source key `phase_estimation_real`) in which translation-only has fewer total gates than the selected candidate but more `cx` gates and a higher score (957,132 versus 870,919). This quality-recovery pattern confirms that semantic lifting behaves as an anti-regression mechanism on workloads beyond the minimal Fourier construction.

### D. Natural-Kernel Factorial Ablation

The frozen W8 campaign crosses nine workload families, scales 1/2/4, three seeds, densities 0.5/1.0, and all-to-all/line hardware topologies. Seven are structured families: QAOA/Ising, diagonal-Hamiltonian simulation, commuting-block Trotterization, QPE controlled powers, QFT arithmetic, controlled-power ladders, and phase-polynomial blocks. Random Clifford+ $R_Z$  and routing-dominated circuits are the two negative controls. For every instance the authors’ reference is evaluated on the

complete  $2^6$  toggle matrix

semantic lift  $\times$  aggregation  $\times$  selector  $\times$  cache/reuse  
 $\times$  preset recognizer  $\times$  projected-block selection.

The authors factorial contributes 20,736 cells and the matched strongest available external semantic baseline, TKET PauliSimp, contributes 324, for 21,060 cells total. Every cell is `completed_valid`, passes the independent dense projective certificate at six qubits, and carries the same fixed-total-error W7 QRE fields. Size increases repeat and densify structure rather than relaxing the certificate width.

Five structured families satisfy the predeclared all-size/all-seed criterion; QFT arithmetic and phase-polynomial blocks do not, and remain visible in table X and figure 11. The largest isolated gate-count contribution is the preset recognizer, with mean reduction 364.9 and bootstrap interval [349.8, 380.5]. Aggregation and semantic lift also have positive mean effects; selector, cache/reuse, and projected-block selection have zero isolated gate effect in this controlled implementation. Both negative controls have zero all-on-minus-all-off differences in gates, depth, CX, logical  $T$ , and spacetime, hence 0 systematic quality regressions in the frozen grid. These are finite reference measurements, not empirical proofs of the streaming theorem or universal compiler superiority.

## X. MODEL-SPECIFIC FAULT-TOLERANT RESOURCE CONSEQUENCES

The representation gap propagates into the fault-tolerant resources of the explicit model in table VIII and figures 9 and 13. Over all 96 matched target/error/allocation/QEC/factory settings, materialize-first uses 385–1,753 times the logical  $T$  states and 538–5,573 times the spacetime volume of semantic-first. No reversal occurs within the scanned error and factory grid. This is not a claim that the authors’ implementation dominates other semantic tools: external TKET PauliSimp reconstructs the phase structure before synthesis and uses 0.443–0.536 times the semantic-first  $T$  states on this witness. Correctly crediting that recovery is part of the comparison.

These values are outputs of actual checked angle synthesis followed by the disclosed scheduling models, not universal hardware predictions. The absolute physical numbers still depend on routing, correlated-error, factory layout, feed-forward, and code-distance assumptions outside this campaign. The complete sensitivity plot and all preflight anomalies are retained in the root QRE artifact package.

## XI. RELATED WORK

a. *Phase-polynomial, Pauli-gadget, and ZX optimization.* Commuting-rotation aggregation is a standard

| Instance          | Method                       | Gates   | Depth   | CX     | Runtime  |
|-------------------|------------------------------|---------|---------|--------|----------|
| REAL24-phase-est. | translation only             | 159,838 | 121,214 | 67,608 | 0.080 s  |
|                   | qiskit opt3                  | 166,805 | 119,204 | 58,491 | 1.298 s  |
|                   | upstream UCC v0.4.12 default | 471,819 | 328,271 | 58,491 | 80.923 s |
|                   | semantic UCC                 | 166,805 | 119,204 | 58,491 | 1.364 s  |
| REAL24-grover     | translation only             | 85,403  | 55,892  | 33,408 | 0.104 s  |
|                   | qiskit opt3                  | 79,029  | 54,163  | 32,544 | 0.552 s  |
|                   | upstream UCC v0.4.12 default | 287,047 | 152,298 | 32,544 | 1.915 s  |
|                   | semantic UCC                 | 79,029  | 54,163  | 32,544 | 0.556 s  |
| REAL24-qaoba      | translation only             | 36,512  | 2,416   | 23,808 | 0.022 s  |
|                   | qiskit opt3                  | 36,512  | 2,416   | 23,808 | 0.272 s  |
|                   | upstream UCC v0.4.12 default | 167,833 | 7,133   | 23,808 | 1.010 s  |
|                   | semantic UCC                 | 36,512  | 2,416   | 23,808 | 0.273 s  |

TABLE IX. Official real algorithm-family instances. The upstream UCC v0.4.12 row is the frozen pre-semantic default defined in the provenance map of Appendix B. Parity with `qiskit opt3` reflects conservative portfolio selection; this panel is an anti-regression check, not evidence of semantic recovery.

| Family                   | external-minus-authors gates [95% CI] | strict success |
|--------------------------|---------------------------------------|----------------|
| controlled_power_ladder  | 9.0 [5.7, 12.7]                       | yes            |
| diagonal_hamiltonian     | 38.0 [25.9, 51.8]                     | yes            |
| phase_polynomial_blocks  | -24.3 [-33.8, -16.3]                  | no             |
| qaoba_ising              | 44.7 [33.7, 56.6]                     | yes            |
| qft_arithmetic           | 8.1 [1.9, 13.8]                       | no             |
| qpe_controlled_powers    | -12.0 [-22.0, -3.0]                   | yes            |
| trotter_commuting_blocks | 33.0 [17.6, 48.2]                     | yes            |

TABLE X. Complete structured-family summary against the matched external semantic-capable baseline. The predeclared strict criterion requires valid outputs, no worse gate and CX counts at every primary cell, and lower measured compiler runtime at every primary scale and seed. Non-winning families are retained.

optimization target [8, 15–17, 20–22, 40]. The phase-polynomial line makes the generator-coefficient structure explicit and then re-synthesizes the parity network, with or without a connectivity constraint [18, 19, 25]; the ZX line keeps a global diagrammatic normal form in which phases migrate and fuse across syntactic boundaries [23, 41]. All of these are *upper-bound* constructions: they exhibit a compiler that recovers structure on the inputs where their rewrite system applies. Our contribution is complementary and in the opposite direction: we formalize a representation-dependent recoverability boundary with an explicit  $\epsilon$ -resource tradeoff that no such pass can evade without paying in the charged channels, connect it to communication complexity, and measure model-specific fault-tolerant consequences under one disclosed QRE. Where these tools do recover the aggregate—as rebased TKET PauliSimp does on our witnesses—the bound predicts, and section VIII D measures, the state they must carry to do it.

*b. Gate synthesis and  $T$ -count accounting.* The exchange rate in our statement is set by the synthesis grid. Solovay–Kitaev decomposition [5], exact Clifford+ $T$  synthesis [6], and the number-theoretic approximation algorithms of Ross–Selinger type [7, 8] all realize a  $Z$ -rotation as a word of length  $\Theta(\log(1/\epsilon))$  whose letters individually carry no part of the angle; multi-qubit  $T$ -count lower bounds sharpen the same accounting at the unitary level [42]. We use these results as the quantitative source

of  $K_\epsilon$  and of the repair cost per mis-committed coordinate, not as objects of study.

*c. Compiler IRs, verified compilation, and full-stack tooling.* Quantum software stacks use intermediate representations to preserve information across lowering [10, 11, 24, 29, 40, 43–45]. Recent designs are explicit about what the IR retains: OpenQASM 3 admits timing and classical control in the instruction stream [46], while SSA- and MLIR-based representations expose dataflow so that passes can be composed and verified [47, 48]. Our budget-explicit model charges exactly this retained information— $B_{\text{IR}}$  in equation (3)—rather than treating a richer IR as a free side channel, which is what makes the semantic and materialized pipelines comparable at all.

*d. Routing, mapping, and benchmark suites.* Architecture-directed passes reshape the stream that a downstream compiler sees [24, 38, 39], and the platform-specific compilers for shuttling-based trapped ions [49] and dynamically field-programmable neutral-atom arrays [50] show how far that reshaping can go before the stream reaches a downstream stage at all. In our model the SWAP and routing records they insert are Clifford and therefore interior to a block, so they normalize away (lemma 6) without weakening the bound; the routing-dominated family is used as a negative control in section IX for precisely this reason. Benchmark suites supply the workload conventions we follow [36, 37, 45, 51].




FIG. 11. Matched natural-family comparison and negative controls. Error bars are paired bootstrap 95% intervals; all structured families and both controls are retained.




FIG. 12. All six factorial gate-count main effects and pairwise interactions. Zero-effect factors are displayed rather than omitted.

*e. Fault-tolerant resource estimation.* The downstream cost of a  $T$  gate is set by distillation and by the scheduling model [13, 14, 32–34]. End-to-end estimates for concrete algorithms [35] and the estimation frameworks that generalize them [11, 12] make the same point we rely on in section X: constant-factor differences upstream of synthesis become order-of-magnitude differences in spacetime volume, and the margin available to a quadratic speedup is thin enough that they matter [52]. Our QRE campaign is an instance of this methodology under a disclosed model, not a new estimator.

*f. Noise tailoring and error mitigation.* Randomized compiling and Pauli twirling [9, 27, 28] are the mechanism that produces the dispersed stream we analyze, and their sign randomization is exactly the transformation our lemma 6 normalizes. The broader mitigation litera-

ture [53], including learned sparse noise models used for probabilistic error cancellation [54], shares the structural feature that matters here: the physically executed circuit is a randomized ensemble whose individual members do not display the quantities the source program wrote down.

*g. Classical lower-bound methodology.* The proof imports fooling-set and communication-complexity reasoning into a quantum-compilation representation question [30, 55, 56]. The pass-explicit form of the frontier follows the classical space-passes tradition for one-way and multi-pass computation over a stream [57, 58], where the object charged is the state serialized at a cut rather than total work.






FIG. 13. Sensitivity to total error, QEC profile, and factory count on the 20k witness. Top: logical  $T$  demand and the conservative-profile spacetime volume. Bottom: ratios to semantic-first, with bands spanning both QEC profiles and all four factory counts. The dashed line is equality. No materialize-first reversal occurs in the measured grid; the strong external TKET path is reported on the semantic-capable side.

## XII. DISCUSSION

### A. Representation as a Resource

The central message is that representation is a computational resource in quantum compilation. A compiler receiving a semantic representation preserving generator-coefficient structure can aggregate in  $O(m)$  records. The cyclic packing gives  $\Omega(m \log(1/\epsilon))$  bits separately for family-relative and self-contained output codes and is compatible with an  $r$ -round modular share stream. In the stated forward-scan model its pass-explicit frontier is  $\bar{A}_p + (2p - 1)S \geq (1 - \delta)m \log_2 K - h_2(\delta)$  (theorem 5). Thus a deferred flat compiler needs  $S = \Omega(m \log K/p)$ . The block hybrid attains  $S = O((m/p) \log(K + 1))$  with compact family-relative output, within the proved constant factor, while the direct semantic vector uses only  $O(\log(mK))$  state (propositions 10 and 11). The older

binary two-round theorem remains a sharper zero-error, one-pass special case.

This boundary is channel-explicit. Equation (3) serializes control, store, window structure, every live parameter, every live IR object, cache/hash layout, dynamic-pass state, and random seed. A compiler that reconstructs a global aggregate—as rebased TKET PauliSimp does on the fixed-width witnesses—may succeed, but its reconstruction table or graph is charged in  $B_{\text{cut}}$  rather than treated as a free semantic side channel.

### B. Scope and Limitations

1. The results have different scopes. The unified theorem (theorem 5) covers  $r$ -round additive shares, randomized failure, and  $p = p(m)$  forward scans of the round-major stream, but charges all  $2p - 1$  serialized



crossings and claims only a constant-factor state upper envelope. Reverse scans, an uncharged random-access input oracle, or a different output code are outside that theorem. The older representation-dependent tradeoff (theorem 2) is scoped to the two-round masked-share stream and gives the stronger deterministic one-pass mask count. The exact-semantics corollary is scoped to fixed width, constant pass/cut budget, and certified noncollision. The growing-width theorems instead use the complete bit codec and allow random access and dynamic  $p = p(m, I, \rho)$ . The decoder is representation-invariant, but all syntax lengths retain their declared encoding semantics.

2. The tuned cyclic packing is stated for  $0 < \epsilon \leq 1/8$  and chooses sharing modulus  $Q_\epsilon = K_\epsilon + 1$ . Its coefficient interval can cross  $\pi$ ; theorem 3 handles circular distance and global phase exactly. The prior no-wrap packing and the binary/ $\mathbb{Z}_3$  theorem remain separate valid specializations. The configured scaling experiment uses  $K = 2$ ,  $\Delta = \pi/4$  and is not claimed to instantiate the masked-share theorem.
3. The fixed-total-error experiment performs actual Clifford+ $T$  synthesis and executable surface-code calculations under two disclosed physical/QEC profiles, but these profiles share one phenomenological model family. Absolute physical counts still require cross-validation under an independent estimator and architecture-specific routing/factory layouts.
4. The packing-family scaling experiment (section VIII C) links the growing- $m$  theorems to configured-pipeline behavior. The instrumented reference compilers measure the masked-share cut fields directly (section VIII D), but they are controlled formal-model objects, not instrumented versions of every external compiler.
5. Dense independent certification fixes the W8 natural-workload width at six qubits; the size axis repeats and densifies structure rather than increasing Hilbert-space width. Runtime and RSS are local-machine measurements. The two negative controls and three zero gate-effect ablation factors are retained as finite null results, not safety theorems.
6. The clean-room reconstruction creates an independent working directory and fresh virtual environment bound to locked project packages. It proves directory-level reconstruction of selected theorem, benchmark, and QRE results, not binary or container independence from the host.

### XIII. CONCLUSION

We have formalized approximate recoverability with a growing-width token codec and a bit-level cut budget

whose five executable serialization fields include parameter, RAM/cache, random-seed, dynamic-pass, and IR channels. The cyclic  $K_\epsilon$ -ary family has exact minimum separation  $2 \sin(\pi/(2(K_\epsilon + 1)))$ , satisfies  $K_\epsilon = \Theta(1/\epsilon)$ , and admits semantic and  $r$ -round flat representations of the same  $U_x$  even though no individual update determines  $x_j$ . For randomized failure  $\delta$  and  $p = p(m)$  forward passes, each named output code and the prefix-side commitment plus crossing transcript require at least  $(1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$  expected bits. Hence a deferred compiler satisfies  $(2p - 1)S \geq (1 - \delta)m \log_2 K_\epsilon - h_2(\delta)$ . A block hybrid reaches compact family-relative output with  $S = O((m/p) \log(K_\epsilon + 1))$ , within the proved constant factor, and the randomized relative-output bound is matched up to framing. The semantic vector achieves compact output with logarithmic state. The exact-semantics corollary retains the fixed-width final-output lower bounds of the prior version. The traced reference-compiler suite measures the named commitment, state, and IR channels directly; the matched and external campaigns show when global semantic tools recover compact structure; and the natural-workload factorial records positive, null, and negative-control outcomes without selection. Fixed-total-error synthesis/QRE exposes a downstream physical consequence under an explicit model. Every W3–W9 headline measurement is generated from the normalized frozen-data registry, while retained legacy diagnostics are labelled historical. Recoverability should be treated as a controlled variable in compilation and resource estimation.

### REPRODUCIBILITY

For scale, the headline campaigns comprise 592 traced reference-compiler runs; a 1,080-cell matched-representation design and a separate 1,728-cell external campaign containing TKET PauliSimp and PyZX full-reduce, with completed metrics accepted only through the stated symbolic or dense certificate; 288 complete fixed-total-error fault-tolerant estimates under a common algorithmic, synthesis, and logical-failure budget; and 21,060 natural-workload and negative-control cells crossing a six-factor ablation, in which 5 structured families satisfy the predeclared all-size/all-seed criterion and the two negative controls show 0 systematic regressions. Explicitly labelled historical diagnostics retain their original artifact provenance.

The canonical registry is `data/manifest.yaml`; the normalized schema is `schema/experiment.schema.json`. Its complete key is

```
experiment_id, instance_id, representation_id,
method_id, seed, backend_hash, tool_version,
artifact_commit, config_hash.
```

The normalized Parquet index contains 24,748 unique records and retains 48 predicate-error rows outside quality means. The data audit checks CSV/Parquet mirrors, conflicting numeric cells, duplicate TeX labels, source

hashes, row counts, orphaned frozen files, and historical versus canonical campaign disposition. It finds no duplicate primary keys, registered-file omissions, conflicting mirror cells, hash drift, or unregistered frozen artifacts in the final registry.

All measurement macros, tables, and figures used by the W3–W9 sections are generated from registered frozen sources by `scripts/build_all_tables.py` and `scripts/build_all_figures.py`; the operation and unique source row for each printed number are recorded in `data/provenance_map.yaml` and tested by `tests/test_paper_number_provenance.py`. The correctness language, symbolic checker, and mutation/dense audit are in `certificates/`; the QRE schemas and two physical profiles are in `qre/`; and raw W3–W8 run directories, manifests, logs, certificates, seeds, and package versions remain registered rather than collapsed into the summary tables.

The one-command entry point is `reproducibility/run_submission.sh`. It reruns data audits, regenerates all tables and figures, executes the full test suite, builds this manuscript and the supplement, and rejects unresolved references. The clean-room runner copies a minimal artifact to an independent temporary directory, creates a fresh environment bound to the locked project packages, reruns 9 headline-theorem tests, compiles and densely certifies a matched QAOA/Ising instance through the authors and TKET paths, and exercises 17 deterministic synthesis pairs in a fixed-total-error QRE subset. This is working-directory isolation, not a claim of host-binary independence. Complete audit and limitation records are `DATA_AUDIT_REPORT.md`, `CLEAN_ROOM_REPORT.md`, `QRE_REPORT.md`, and `NATURAL_WORKLOAD_REPORT.md`.

## ACKNOWLEDGMENTS

J.Y.’s deepest thanks go to two teachers who are also coauthors of this paper, Yangyang Li and Xiu-Hao Deng. Their help was immense; without them this paper would not exist.

J.Y. thanks Ke Lin of Xidian University for opening the door to research, and for the recommendation that led to joining Prof. Li’s group.

J.Y. thanks Brad Chase: the first ideas behind this paper arose while fixing an issue in his `ucc` repository.

J.Y. also thanks the teachers who were willing to offer guidance along the way, and the senior members of Prof. Li’s group, who have been generous with their help throughout this work.

Thanks are due, finally, to everyone who supported this paper, too many to name individually here.

Y.L. acknowledges support from the National Natural Science Foundation of China under Grant No. 62476209, the Key Research and Development Program of Shaanxi under Grant No. 2024CY2-GJHX-18, and the Fundamental Research Funds for the Central Universities under Grant No. QTZX26097. X.-

H.D. acknowledges support from the Science, Technology and Innovation Commission of Shenzhen Municipality (Grant Nos. JCYJ20170412152620376 and KYTDPT20181011104202253), the Innovation Program for Quantum Science and Technology (Grant No. 2021ZD0301703), the Guangdong Major Project of Basic Research (Grant No. 2025B0303000007), and the Shenzhen Science and Technology Program (Grant No. KQTD20200820113010023).

AI assistants were used for drafting and revision, repository and code assistance, literature discovery, and consistency checks. Their output was not treated as a scientific source or as independent proof verification. The authors retain full responsibility for the manuscript, proofs, code, data, figures, citations, and conclusions.

This paper takes some trouble over provenance: every printed number is traceable to a unique source row, every experiment to a frozen registry, every prediction to a criterion written down before the run. That being so—*what one begins with, one must end with* (Zuo Zhuan, Xuan 12)—recording the provenance of the framing as well is only seeing the matter through, even though that provenance does not lie in the physics literature.

Two passages from the classical Chinese canon shaped parts of this work. Their bearing is plain and their neighborhood fitting, and so they are quoted here.

The first is from the *Xici* commentary to the *Yijing*: “What is above form is called the Way; what is below form is called the vessel; to transform and tailor it is called change; to carry it forward and put it into practice is called continuity.” What is useful here is not an analogy but a set of distinctions. The passage names four things separately, and ranks none of them above another: a description that carries structure, a description that carries only instructions, the act of lowering the one into the other, and the act of execution. A compiler is precisely the engine that performs the third and hands off the fourth; and the vocabulary of compilation itself—semantic level, lowering, commitment—reproduces the same fourfold division without, so far as we know, having borrowed it. The suggestion taken from the passage is that the level at which a computation is described is not a notational convenience to be optimized away, but a variable that can be charged for. This paper charges it, in bits.

The second is Chapter 64 of the *Laozi*: “What is at rest is easy to hold; what has not yet shown a sign is easy to plan for; what is brittle is easy to shatter; what is minute is easy to scatter. Act on it before it comes to be; order it before it falls into disorder.” Every clause here prices an action by the state of the object acted upon rather than by the action itself—which is exactly this paper’s claim about representation: the same aggregation is cheap or expensive depending only on the form in which the layer arrives. The verb of the first clause, *to hold*, is the operation charged here as compiler state. The chapter further separates planning early, which it recommends, from grasping early, of which it warns that “one who

grasps it loses it”; that separation is, here, the difference between a compiler that defers and one that commits, and the price of the difference is the commitment toll. And the rule on which the chapter closes is the design rule of our abstract, read from the other side.

Neither passage contains a theorem, and neither anticipates one; no result reported here depends on either. A reader who finds these readings strained may discard them entirely and judge the paper by whatever in it proves useful. They are recorded because the question did come from there, and because a framing, like a datum, may rightly be asked by a reader where it came from. Translations are J.Y.’s own.

J.Y. conceived the study, developed the theory and proofs, implemented the compiler instrumentation and experimental pipeline, performed the majority of the measurements and analyses, and wrote the manuscript. Y.L. and X.-H.D. carried out part of the measurements and analyses, supervised the work, contributed to the framing and interpretation of the results, and revised the manuscript. All authors discussed the results and approved the final version.

## Appendix A: Supplementary Theory and Full Proofs

### 1. Operational Gap Notation

For reference, the experimental gap can be written without adding another formal theorem hypothesis. Let:

- $L_B(C)$  be exact lowering into target basis  $B$ ,
- $\kappa$  be a conservative structural cost,
- $\mathcal{T}_K(C)$  be the bounded set of semantic candidates produced by the semantic lift under budget  $K$ ,
- $\mathfrak{A}$  be a configured post-lowering pipeline family used for an experiment.

Define the  $(K, w)$ -representation gap of  $C$  by

$$\text{Gap}_{\text{rep}}^{K,w}(C) := \inf_{A \in \mathfrak{A}} \kappa(A(L_B(C))) - \inf_{T \in \mathcal{T}_K(C)} \kappa(L_B(T)).$$

Define the operational recoverability gap by

$$\text{Gap}_{\text{rec}}^{K,w}(C) := \kappa(C_{\text{flat}}^*) - \kappa(C_{\text{term}}^*),$$

where  $C_{\text{flat}}^*$  is the best candidate recovered from the flat basis-lowered representation and  $C_{\text{term}}^*$  is the best candidate recovered from the bounded semantic term representation.

The meaning is simple: semantics may survive lowering, while cheap recoverability of the opportunity may not. This notation is operational and experimental; the formal lower bound in the main text is the budget-explicit final-output theorem for finite-window transducers.

**Definition 26** (Aggregate recovery). *Consider the Fourier-layer family  $C_{m,r} = H_\Lambda D_m(\theta)^r H_\Lambda$  with  $D_m(\theta) = \prod_{j=1}^m e^{-i\theta_j P_j/2}$ . A compiler performs aggregation-equivalent recovery on this family if, for each non-removable term  $P_j$ , its successful output can be associated with an  $O(1)$ -size record, subcircuit, or algebraic object whose coefficient is equivalent modulo  $2\pi$  to the aggregate  $r\theta_j$ . The object may be explicit, as in a phase-polynomial or commuting-diagonal IR, or implicit, as in a specialized rewrite, lookup, ZX transformation, or synthesis routine that computes the same aggregate diagonal action without storing all  $r$  lowered occurrences separately.*

**Proposition 14** (Operational recovery channels). *For the Fourier-layer family, a pipeline that recovers the constant fixed-width output can be classified by the channel through which it carries the aggregate coefficient information. Examples include:*

1. store enough persistent summary information to carry the aggregate;
2. lift enough lowered volume into a richer IR to reconstruct the aggregate;
3. use nonlocal provenance information from before basis lowering that already contains the aggregate;
4. reconstruct a global semantic representation equivalent to a phase-polynomial, ZX, stabilizer-phase, or commuting-diagonal Hamiltonian object.

Consequently, a successful future flat-looking implementation would not by itself refute the main theorem. It would instead classify the implementation by identifying where the aggregate information crosses the finite-state boundary, or by showing that the formal transducer model is missing a cheaper channel.

*Justification.* The fixed-width theorem lower-bounds final output length for constant-budget finite-window transducers on a certified noncolliding family. A pipeline that achieves the semantic constant output must still represent the aggregate coefficients that distinguish the noncolliding cases. The listed channels are implementation-level ways to carry that information. This is a classification statement rather than an optimality claim.  $\square$

### Additional Structural Families

The artifact also targets phase-ladder, mirrored, and conjugation patterns that appear in QPE/QFT-style workloads and Grover-style workloads. These families are not used here as additional theorem instances. They are best read as structural stress families: the semantic controller preserves repeated phase-heavy ladders, mirrored shells, and conjugation shells long enough to compare projected costs before materializing the selected candidate.

For phase-ladder workloads, the relevant empirical question is whether bounded semantic terms preserve a repeated commuting-diagonal opportunity that is harder to recover after early lowering. For mirrored and conjugation workloads, the relevant empirical question is whether the artifact recognizes repeated shell structure and avoids materializing every candidate before selection. The Grover and mirrored-shell experiments in Appendix C are therefore empirical positive cases and robustness checks, not support for a separate theorem-level family separation.

## 2. Self-Contained Rational CP Period Certificate

The canonical four-qubit rational witness has diagonal block

$$D = \prod_{j=0}^3 RZ_j(\theta_j) \prod_{0 \leq a < b \leq 3} CP_{ab}(\theta_{ab}),$$

$$\theta_{RZ}/\pi = (1/7, 1/8, 1/9, 1/10),$$

$$(\theta_{01}, \theta_{02}, \theta_{03})/\pi = (1/12, 1/13, 1/14),$$

$$(\theta_{12}, \theta_{13}, \theta_{23})/\pi = (1/14, 1/15, 1/16).$$

For a reduced rational angle  $\theta/\pi = p/q$ , both  $RZ(\theta)$  and  $CP(\theta)$  have projective single-term order

$$\text{ord}(p/q) = \frac{2q}{\gcd(|p|, 2q)},$$

because the relevant relative phase is  $p/q$  modulo 2. The complete single-term certificate is

| term      | support | $\theta/\pi$ | $T_j$ |
|-----------|---------|--------------|-------|
| $RZ_0$    | 0       | 1/7          | 14    |
| $RZ_1$    | 1       | 1/8          | 16    |
| $RZ_2$    | 2       | 1/9          | 18    |
| $RZ_3$    | 3       | 1/10         | 20    |
| $CP_{01}$ | 01      | 1/12         | 24    |
| $CP_{02}$ | 02      | 1/13         | 26    |
| $CP_{03}$ | 03      | 1/14         | 28    |
| $CP_{12}$ | 12      | 1/14         | 28    |
| $CP_{13}$ | 13      | 1/15         | 30    |
| $CP_{23}$ | 23      | 1/16         | 32    |

Hence

$$T_{\text{terms}} = \text{lcm}(14, 16, 18, 20, 24, 26, 28, 28, 30, 32)$$

$$= 2^5 \cdot 3^2 \cdot 5 \cdot 7 \cdot 13 = 131040,$$

which proves that  $D^{131040}$  is a global phase. To certify minimality for the product rather than merely obtain this common upper bound, subtract the phase of 0000. For  $z = z_0 z_1 z_2 z_3$ , the exact relative phase is

$$w(z) = z_0/7 + z_1/8 + z_2/9 + z_3/10$$

$$+ z_0 z_1/12 + z_0 z_2/13 + z_0 z_3/14$$

$$+ z_1 z_2/14 + z_1 z_3/15 + z_2 z_3/16 \pmod{2}.$$

In particular,

$$w(1011) = \frac{37007}{65520},$$

$$\text{ord}(w(1011)) = \frac{2 \cdot 65520}{\gcd(37007, 2 \cdot 65520)}$$

$$= 131040.$$

Thus no smaller positive power of  $D$  is a global phase, and  $T = 131040$  exactly. It follows that  $D^1, \dots, D^R$  are pairwise distinct up to global phase for every  $R < T$ .

The machine certificate uses only Python `fractions.Fraction` and prints all 16 relative numerators, denominators, and orders. The frozen integer output contains

```
term_order_lcm=131040,
relative_order_lcm=131040,
minimality_state_1011_order=131040,
exact_global_phase_period_T=131040.
```

The reproducibility files and hashes are listed below; each digest is split at its 32-hex midpoint only for typesetting.

- verifier: `research/rational_witness/verify_period.py`, SHA256 `c7c90f26e00e515383e94e4bbd707574eb3aabf94ad71c3cd5e2cdb8d5db7e76`;
- integer output: `research/rational_witness/period_certificate.txt`, SHA256 `375295d776f21ea2a82ee4f2bab5205d25bfef85ee864900baacfbdb80be9bd6e`;
- manifest: `research/rational_witness/SHA256SUMS`.

From the repaired-manuscript root, the byte-for-byte check is

```
cd research/rational_witness
python3 verify_period.py \
-check period_certificate.txt
```

## 3. Full Proofs for the Transducer Lower Bound

This section supplies the full finite-state proof layer underlying Corollaries 1 to 3. The main text states the shorter version; the appendix keeps the pumpability data explicit. The following statements restate the main-text theorem and corollaries with full hypotheses and complete proofs; they are not additional independent results.

### a. Pumpable Pair Families

Persistent store can depend on  $r$ , so the induction tracks both a store value and a tape word. In this subsection,  $D$  in pumpability data denotes a period modulus, not

the diagonal witness layer in  $C_r = H_\Lambda D^r H_\Lambda^{-1}$ . The next definition restates the main-text pumpable-pair definition with explicit residue data for the full proof.

**Definition 27** (Pumpable pair family, restated with explicit residue data). *Let  $\Gamma$  be a finite alphabet and  $P$  a finite set of store values. A family*

$$X(r) = (\beta(r), W(r)) \in P \times \Gamma^*, \quad r \in \mathbb{N},$$

*is pumpable if there exist integers  $h \geq 0$  and  $D \geq 1$  such that for every residue  $\rho \in \{0, \dots, D-1\}$  there are*

$$\beta_\rho \in P, \quad A_\rho, B_\rho, C_\rho \in \Gamma^*,$$

*with*

$$X(h + \rho + kD) = (\beta_\rho, A_\rho B_\rho^k C_\rho)$$

*for every  $k \geq 0$ . The word  $B_\rho$  is the pump word on residue  $\rho$ . The residue is live if  $B_\rho \neq \epsilon$  and collapsed if  $B_\rho = \epsilon$ .*

**Lemma 11** (Input pumpability). *For the fixed-width witness  $T_0(r) = \Pi M^r \Sigma$ , the pair family consisting of the initial persistent store and  $T_0(r)$  is pumpable with  $h = 0$ ,  $D = 1$ , pump word  $M$ , and fixed store value.*

*Proof.* Take  $A_0 = \Pi$ ,  $B_0 = M$ , and  $C_0 = \Sigma$ .  $\square$

### b. One-Pass Propagation

Fix a precision bound  $q$  and a charged alphabet  $\Gamma = \Gamma_B(q)$ . Under this bound, a pass of  $\mathcal{A}$  can be viewed as an ordinary deterministic subsequential transducer over a finite alphabet. Its finite state consists of the ephemeral state, the persistent store, and the  $w$ -token window. The end-of-pass persistent store is one component of the final finite state.

**Lemma 12** (One-pass pumpability propagation). *Let  $X(r) = (\beta(r), W(r))$  be a pumpable pair family over a finite input alphabet. Let  $\tau$  be one deterministic finite-window, finite-burst pass whose output alphabet and store set are finite. Let*

$$\tau(\beta(r), W(r)) = (\beta'(r), W'(r))$$

*be the resulting final store and output tape. Then  $X'(r) = (\beta'(r), W'(r))$  is pumpable.*

*Proof.* Fix pumpability data  $h, D, \{\beta_\rho, A_\rho, B_\rho, C_\rho\}_\rho$  for  $X$ . We first treat one residue  $\rho$ , then recombine all residues into  $r$ -indexed pumpability data for  $X'$ .

For a fixed residue  $\rho$ , write

$$W_k = AB^kC, \quad \beta_k = \beta,$$

where  $k \geq 0$  and  $A, B, C, \beta$  abbreviate the residue data. Absorb the ephemeral state, persistent store, and window into the pass's finite state set  $S$ . Starting from the state

determined by  $\beta$ , process  $A$ . This emits a fixed word  $O_A$  and reaches a state  $s_A \in S$ .

Let  $f : S \rightarrow S$  be the state map obtained by reading the block  $B$ , and let  $o(s) \in (\Gamma')^*$  be the output emitted while reading  $B$  starting in state  $s$ . The sequence

$$s_A, f(s_A), f^2(s_A), \dots$$

in the finite set  $S$  is ultimately periodic. Thus there exist  $u \geq 0$  and  $\ell \geq 1$  with  $u + \ell \leq |S| + 1$  such that

$$f^{t+\ell}(s_A) = f^t(s_A) \quad \text{for all } t \geq u.$$

For  $k = u + v\ell + \eta$  with  $0 \leq \eta < \ell$ , the output while reading  $B^k$  is

$$\begin{aligned} & (o(s_A) o(f(s_A)) \cdots o(f^{u-1}(s_A))) \\ & \cdot \left( o(f^u(s_A)) \cdots o(f^{u+\ell-1}(s_A)) \right)^v \\ & \cdot (o(f^u(s_A)) \cdots o(f^{u+\eta-1}(s_A))). \end{aligned}$$

After this, the pass reads the fixed suffix  $C$ , emits a word depending only on the state  $f^{u+\eta}(s_A)$ , and ends in a persistent store value depending only on the same state.

Restore the residue subscripts. For each original residue  $\rho \in \{0, \dots, D-1\}$ , let  $u_\rho \geq 0$  and  $\ell_\rho \geq 1$  be the preperiod and cycle length obtained above, let  $f_\rho$ ,  $s_{A_\rho}$ , and  $O_{A_\rho}$  be the corresponding state map, post-prefix state, and prefix output. For  $0 \leq \eta < \ell_\rho$ , let  $F_{\rho,\eta}$  be the complete word emitted while reading the fixed suffix  $C_\rho$  from state  $f_\rho^{u_\rho+\eta}(s_{A_\rho})$ , including the end-of-tape finalizer output, and define

$$\begin{aligned} P_\rho &:= o(s_{A_\rho}) o(f_\rho(s_{A_\rho})) \cdots o(f_\rho^{u_\rho-1}(s_{A_\rho})), \\ Q_\rho &:= o(f_\rho^{u_\rho}(s_{A_\rho})) \cdots o(f_\rho^{u_\rho+\ell_\rho-1}(s_{A_\rho})), \\ R_{\rho,\eta} &:= o(f_\rho^{u_\rho}(s_{A_\rho})) \cdots o(f_\rho^{u_\rho+\eta-1}(s_{A_\rho})) \cdot F_{\rho,\eta}. \end{aligned}$$

Let  $\beta'_{\rho,\eta}$  be the end-of-pass persistent store value reached from that state. Define

$$\begin{aligned} u^* &:= \max_\rho u_\rho, & \ell^* &:= \text{lcm}_\rho \ell_\rho, \\ h' &:= h + Du^*, & D' &:= D\ell^*. \end{aligned}$$

Each new residue  $\rho' \in \{0, \dots, D' - 1\}$  decomposes uniquely as

$$\rho' = \rho + D\tilde{\eta}, \quad \rho \in \{0, \dots, D-1\}, \quad \tilde{\eta} \in \{0, \dots, \ell^* - 1\}.$$

For  $r = h' + \rho' + k'D'$  with  $k' \geq 0$ , the old representation  $r = h + \rho + kD$  has

$$k = u^* + \tilde{\eta} + k'\ell^*.$$

Since  $u^* \geq u_\rho$ , we have  $k \geq u_\rho$  for every  $k' \geq 0$ ; and since  $\ell_\rho \mid \ell^*$ , the residue of  $k - u_\rho$  modulo  $\ell_\rho$  does not depend on  $k'$ . Set

$$\eta := (u^* + \tilde{\eta} - u_\rho) \bmod \ell_\rho, \quad v_0 := \left\lfloor \frac{u^* + \tilde{\eta} - u_\rho}{\ell_\rho} \right\rfloor \geq 0,$$

so that

$$k = u_\rho + (v_0 + k'\ell^*/\ell_\rho)\ell_\rho + \eta.$$

By the output decomposition above, for every  $k' \geq 0$ ,

$$X'(h' + \rho' + k'D') = (\beta'_{\rho,\eta}, A'_{\rho'}(B'_{\rho'})^{k'}C'_{\rho'}),$$

where

$$A'_{\rho'} := O_{A_\rho} P_\rho Q_\rho^{v_0}, \quad B'_{\rho'} := Q_\rho^{\ell^*/\ell_\rho}, \quad C'_{\rho'} := R_{\rho,\eta}.$$

This is exactly pumpability of  $X'$  with data  $(h', D')$  and residue data  $\{\beta'_{\rho,\eta}, A'_{\rho'}, B'_{\rho'}, C'_{\rho'}\}_{\rho'}$ .  $\square$

**Corollary 8** (Per-tape induction, for proof). *Assume  $\mathcal{A}$  has constant budget, so all executions use one finite alphabet  $\Gamma_B(q_*)$ . For every pass index  $i = 0, 1, \dots, p$ , the pair family consisting of the persistent store after pass  $i$  and the tape  $T_i(r)$  is pumpable.*

*Proof.* The base case is Lemma 11. Apply Lemma 12 successively to passes  $1, \dots, p$ .  $\square$

**Remark 17** (Unary scratch tapes). *If pass 1 writes one marker per input period, then  $T_1(r) = \text{marker}^r$  is pumpable with live pump word **marker**. Lemma 12 applies to every later pass. A later constant-budget pass may output one marker every  $c$  input markers, permute markers, or add a fixed periodic decoration, but it remains pumpable. It cannot turn  $\text{marker}^r$  into a family of  $O(1)$  many correct aggregate circuits for all noncolliding  $r$ ; if the final pump collapses, the final output repeats on a residue class.*

#### c. Collapse Versus Final-Output Length

**Lemma 13** (Pumpable outputs are bounded-or-linear on each residue). *Let  $W(r)$  be a pumpable word family with data  $(h, D, A_\rho, B_\rho, C_\rho)$ . For  $r = h + \rho + kD$ :*

- (a) if  $B_\rho = \epsilon$ , then  $W(h + \rho + kD)$  is independent of  $k$ ;
- (b) if  $B_\rho \neq \epsilon$ , then

$$|W(h + \rho + kD)| \geq k.$$

Consequently, on a live residue,

$$|W(r)| \geq \frac{r - h - D + 1}{D}$$

for all  $r \geq h + D$  in that residue.

*Proof.* Both statements are immediate from  $W(h + \rho + kD) = A_\rho B_\rho^k C_\rho$ . If  $B_\rho \neq \epsilon$ , then  $|B_\rho| \geq 1$ , so  $|A_\rho B_\rho^k C_\rho| \geq k$ . Since  $r = h + \rho + kD$  and  $0 \leq \rho < D$ ,  $k \geq (r - h - D + 1)/D$ .  $\square$

**Lemma 14** (Correct noncolliding final outputs cannot collapse). *Assume finite-range noncollision (Assumption 1) with  $N(R) = R$  on the interval  $1 \leq r \leq R$ . Suppose the final output family  $T_p(r)$  is pumpable with data  $(h, D, A_\rho, B_\rho, C_\rho)$  and  $\mathcal{A}$  is correct on every  $r \in \{1, \dots, R\}$ . Without loss of generality take  $h \geq 1$ : if  $h = 0$ , replace  $h$  by  $h + D$ ; this preserves pumpability after discarding the finite prefix of indices  $k = 0$  on each residue and absorbing one copy of  $B_\rho$  into  $A_\rho$ . Then every residue class appearing twice in the tail  $\{h, h + 1, \dots, R\}$  is live. In particular, for every  $r \in \{h + D, \dots, R\}$ ,*

$$|T_p(r)| \geq \frac{r - h - D + 1}{D}.$$

*Proof.* Let  $r = h + \rho + kD$  with  $k \geq 1$  and  $r \leq R$ . Then  $r - D = h + \rho + (k - 1)D$  also lies in the tail, and by the convention  $h \geq 1$  it satisfies  $r - D \geq h \geq 1$ , so both  $r$  and  $r - D$  lie in the correctness domain  $\{1, \dots, R\}$ . If  $B_\rho = \epsilon$ , Lemma 13 gives

$$T_p(r) = T_p(r - D)$$

as literal circuit strings. Hence the two outputs have the same unitary. Correctness at both  $r$  and  $r - D$  would imply

$$U(\mathcal{L}_B(C_r)) = U(\mathcal{L}_B(C_{r-D}))$$

up to global phase, which by exactness of the lowering means  $D^r \equiv D^{r-D}$  up to global phase, contradicting Assumption 1 with  $N(R) = R$ . Therefore  $B_\rho \neq \epsilon$ , and Lemma 13 gives the stated length bound.  $\square$

#### d. Final-Output Theorem and Corollaries

**Theorem 6** ( $p$ -pass final-output lower bound, restated with full hypotheses). *Fix a fixed-width diagonal witness family  $C_r = H_\Lambda D^r H_\Lambda^{-1}$  with exact lowering  $\mathcal{L}_B(C_r) = \Pi M^r \Sigma$ , and assume unbounded noncollision (Assumption 2): the powers  $D^r$  are pairwise distinct up to global phase for all  $r \geq 1$ , so  $N(R) = R$  on every range. Let  $\mathcal{A}$  be a uniform deterministic  $p$ -pass finite-window, finite-burst transducer with constant budget that is correct on the input  $\mathcal{L}_B(C_r) = \Pi M^r \Sigma$  for every  $r \geq 1$ . Here uniform means one fixed deterministic transducer, with the same passes, transition rules, finite state, window, burst bound, and store bound for all repetition counts  $r$ .*

*Then there exist constants  $h \geq 1$  and  $D \geq 1$ , depending on  $\mathcal{A}$  and the fixed witness but not on  $r$ , such that for every  $r \geq h + D$ ,*

$$|T_p(r)| \geq \frac{r - h - D + 1}{D} = \Omega_{\mathcal{A}}(r).$$

*Equivalently, no constant-budget transducer in this model that is correct on the whole noncolliding fixed-width family can emit  $o(r)$ -length final outputs on that family.*

*Proof.* By constant budget, all executions use a finite charged token alphabet. Corollary 8 gives pumpability of the final pair family and hence of  $T_p(r)$ , with data  $(h, D)$  depending only on  $\mathcal{A}$  and the witness. By the convention of Lemma 14, take  $h \geq 1$ , so every comparison point  $r - D \geq h \geq 1$  used below lies in the correctness domain.

Fix any  $r \geq h + D$  and any  $R \geq r$ . By Assumption 2,  $N(R) = R$ , and  $\mathcal{A}$  is correct on every  $r' \in \{1, \dots, R\}$ . Thus Lemma 14 applies and gives the displayed bound at  $r$ . In particular every final residue is live: if any final residue collapsed, two sufficiently large inputs in that residue would receive the same final output string, contradicting correctness and noncollision.  $\square$

*a. Correctness quantifier.* Correctness in Theorem 6 is required at every  $r \geq 1$ ; the conclusion then holds at every  $r \geq h + D$ , without exception. This global quantifier is not decorative: the collapse argument in Lemma 14 compares the outputs at  $r$  and  $r - D$  and needs correctness at both points. Correctness at isolated points is not sufficient; if  $\mathcal{A}$  were correct at  $r$  but not at  $r - D$ , identical final strings at those two inputs would yield no contradiction.

**Corollary 9** (Finite-range tail form, restated with full hypotheses). *Let  $\mathcal{A}$  be a constant-budget transducer, and let  $(h, D)$  be pumpability data for its final pair family as provided by Corollary 8. Suppose finite-range noncollision (Assumption 1) holds with  $N(R) = R$  on the range  $1 \leq r \leq R$  and  $\mathcal{A}$  is correct on every  $r \in \{1, \dots, R\}$ . Then*

$$|T_p(r)| \geq \frac{r - h - D + 1}{D} \quad \text{for every } h + D \leq r \leq R.$$

*The constants  $(h, D)$  are determined by  $\mathcal{A}$  and the witness family alone; in particular they do not depend on  $R$ . No statement is made for transducers outside the constant-budget regime: there, pumpability data in the sense of Definition 27 need not exist, and any finite-range inequality would require explicit control of pumping constants not supplied by this theorem.*

*Proof.* Constant budget puts all executions over one finite charged alphabet, so Corollary 8 supplies pumpability data  $(h, D)$  for the final tape. The construction of this data uses only the fixed transducer, its finite alphabets and state sets, and the fixed witness words  $\Pi, M, \Sigma$ ; it is obtained before any finite endpoint  $R$  is chosen. Hence  $(h, D)$  are independent of  $R$ . Lemma 14 then applies verbatim on the range  $1 \leq r \leq R$ .  $\square$

**Corollary 10** (Bit-level output counting toll, restated for proof). *Assume finite-range noncollision supplies  $N(R)$  pairwise distinct target unitaries, up to global phase, on a range. Suppose a deterministic exact compiler emits final outputs  $T_p(r)$ . Then*

$$\max_{1 \leq r \leq R} B_{\text{self}}(T_p(r)) \geq \lceil \log_2 N(R) \rceil.$$

*The same inequality holds for a fixed family-relative code when one public dictionary is shared by the whole range.*

*For the rational-angle CP witness, only the displayed per-range inequality is claimed on fixed ranges  $R < T = 131040$ .*

*Proof.* A circuit string determines its unitary, so a deterministic exact compiler that is correct on two inputs whose target unitaries are distinct up to global phase must emit two distinct output strings. The noncollision hypothesis therefore forces at least  $N(R)$  pairwise distinct codewords. If their prefix-free lengths are  $\ell_1, \dots, \ell_{N(R)}$ , Kraft's inequality gives  $\sum_i 2^{-\ell_i} \leq 1$ , so  $N(R)2^{-\max_i \ell_i} \leq 1$ . Hence the largest length is at least  $\lceil \log_2 N(R) \rceil$ . The argument applies independently to either fixed code and does not convert token count or gate count into bits.  $\square$

#### 4. Bookkeeping for the Masked-Share Cut

This subsection records the restart property used by the masked-share tradeoff (theorem 2). It also makes explicit why a materialized update log, phase table, graph, RAM/cache layout, old-input spool, dynamic-pass state, or random seed is part of the cut message rather than an uncharged reconstruction channel.

**Lemma 15** (Restart from the serialized cut). *Fix  $\mathcal{A}$  as in definition 9. At the inter-round cut  $c_{\text{MS}}$ , the five fields in (3), together with any prefix already committed to the final write-once output, determine the continuation of that run for every supplied second-round continuation  $v$ . For a randomized run this statement is conditional on the serialized seed/tape and cursor.*

*Proof.* The control field restores the exact transition, all head positions, dynamic pass/termination schedule, and random seed/tape cursor. The store and window fields restore RAM, stack, hash/cache metadata, nonparameter machine data, and ordered live token structures. The parameter field restores every omitted coefficient by its location tag. The IR field restores every other live intermediate tape, DAG, table, addressed record, and any consumed input block retained for rescanning. These disjoint fields therefore reconstruct the complete machine configuration. Conditional on the serialized random tape, the receiver can append the supplied second round and execute the same dynamically chosen continuation. The separately transmitted final-output prefix is prepended to the suffix produced after restart. No additional state, old-input file, or public-family side channel is used.  $\square$

## Appendix B: Artifact Engineering

The implementation described in this appendix is a bounded recognizer and candidate controller, not a general maximal-commuting-region extractor for arbitrary DAGs. Its semantic nodes are limited to explicit supported Pauli/phase records, local phase-ladder or Fourier spans, mirrored/self-inverse composite blocks, and repeated-run



or preset semantic references. The word “region” below means a contiguous span accepted by that configured grammar under its scan order. Soundness of a selected rewrite is supplied by the independent certificate and conservative fallback, not by a claim of global maximality.

### 1. Candidate Control and Bounded Pipeline

The semantic path is advisory, not authoritative: semantic terms are lowered late, semantic subterms are compiled once and reused, and a semantic candidate loses to a flat reference whenever it does not clearly improve the conservative structural cost. The term caches, compiled-subterm caches, and repeated-run block caches in the implementation follow from treating the semantic term rather than the materialized circuit as the unit of comparison. The whole method is the following bounded pipeline.

1. **Input.** A high-level circuit  $C$ , target basis  $B$ , and iteration budget  $K$ .
2. Compute stable signatures for instructions and candidate blocks in  $C$ .
3. Run bounded preprocessing  $\mathcal{S}_K(C)$ : repeated-prefix and repeated-run detection and reuse; adjacent block-inverse(block) cancellation; commutative inverse cancellation; semantic lifting into instruction-span, phase-ladder, and mirrored/self-inverse nodes where applicable; early stop if one full round does not improve the selected cost.
4. Construct the candidate set

$$\mathcal{C}(C) = \{\mathcal{L}_B(C), \mathcal{P}_{\text{ucc}}(\mathcal{L}_B(C)), \mathcal{P}_{\text{ucc}}(\mathcal{L}_B(\mathcal{S}_K(C))), \mathcal{R}(C)\},$$

where  $\mathcal{R}(C)$  is an optional semantic or preset-reference candidate for large repeated structures, itself produced from the semantic term algebra (Fourier-layer lowering, mirrored/self-inverse lowering, or a bounded repeated-run rebuild).

5. If the circuit factors as  $C = P \cdot B^r \cdot S$ , form the bounded block-candidate set  $\mathcal{B}(B) = \{B_1, \dots, B_q\}$  and select  $i^* = \arg \min_i \kappa_{\text{proj}}(P, B_i, r, S)$  on the projected full-circuit cost, rebuilding only  $B_{i^*}$ . This is what lets the repeated-run path avoid materializing several giant candidates merely to compare them.
6. Evaluate all candidates under the conservative structural cost  $\kappa$ , with fixed lexicographic tie-breaking, and return  $\hat{C} = \arg \min_{X \in \mathcal{C}(C)} \kappa(X)$ .

The role of candidate control is not to dominate all downstream flows. It is to prevent avoidable structured-circuit regressions by preserving the best available output among a small family of semantically equivalent candidates.

### 2. Complexity Discussion

The method is intentionally bounded and does not attempt global optimization. Let  $n := |C|$  be the number of high-level instructions. Suppose one outer iteration consists of  $r$  bounded scans, and scan  $j$  has cost  $T_j(n)$ . Then the total runtime of preprocessing satisfies

$$T_{\text{pre}}(n) \in O\left(K \cdot \sum_{j=1}^r T_j(n)\right).$$

The dominant scans are near-linear in  $n$  up to implementation-dependent hashing, block-comparison, and bookkeeping costs, so practical behavior stays near-linear up to a bounded constant on repeat-heavy workloads. A large-circuit fast path handles repeat-dominant cases, source-level short-circuiting avoids downstream work on Qiskit-native all-to-all composites, and backend-aware portfolio expansion is restricted to cases where instability matters. The method is therefore a bounded systems heuristic, not a globally optimizing pass.

**Proposition 15** (Expected regime of benefit). *The preprocessing stage is most likely to help on circuits whose useful structure is already explicit at the high-level instruction stream or in the bounded semantic lift, especially exact inverse blocks, repeated blocks, phase-ladder layers, mirrored/self-inverse shells, and dominant repeated prefixes. It is not expected to uniformly improve families whose performance is dominated by later synthesis or routing decisions that are not captured by the preprocessing signatures or semantic nodes.*

*Justification.* This follows from the construction of the pass. Every preprocessing rule operates on high-level structural regularities detected through stable signatures, exact inverse relations, or the bounded semantic lift into semantic terms such as Fourier layers and mirrored/self-inverse shells. Therefore the pass can only exploit structure that survives in this representation. Conversely, if output quality is dominated by later basis-specific synthesis or hardware-routing effects not visible at this stage, the preprocessing pass may improve little or only preserve parity with stronger downstream baselines. The representation-aware compilation principle is therefore selective rather than universal: it claims that some opportunities are materially easier to preserve and compare before lowering, not that all useful optimization should be moved to the semantic layer.  $\square$

### 3. Integration Into UCC

The semantic artifact changes the UCC flow in six places:

1. it performs structural simplification before basis lowering,

2. it lifts suitable subcircuits into a bounded semantic term algebra before candidate generation,
3. it compares candidate outputs conservatively to avoid structured-circuit regressions,
4. it uses source-level short-circuiting on suitable all-to-all Qiskit-native cases,
5. it restricts backend-aware portfolio expansion to cases where extra probing is justified,
6. it memoizes bounded semantic references and repeated-run block selections for reuse.

The organizing distinction is that some circuit growth is unavoidable basis-translation overhead while some is structural regression caused by losing recoverable semantic evidence before candidate selection.

#### 4. Method and Provenance Map

Every frozen numerical row is addressed by the complete primary key

`(experiment, instance, representation, method, seed, backend_hash, tool_version, config_hash)`.

The backend and configuration hashes are SHA256 digests of canonical JSON (sorted keys and no insignificant whitespace). A row collision on all eight fields is rejected. In particular, the reorder-probe hardware rows and the canonical hardware-aware rows have different experiment and configuration keys, and therefore different reader labels even when the underlying MQT instance name is the same. The frozen source is `research/experiment_registry/frozen_hardware_rows.json`; `scripts/generate_keyed_tables.py` validates key uniqueness and generates the three affected table fragments plus `expanded_primary_keys.json`. The generated labels expose the campaign, instance, seed, and the first eight configuration-hash digits. Hand-edited numeric rows are not accepted by this path.

The canonical validation data are grouped by experiment family in the artifact. The external-baseline method map immediately below uses the disclosed `timeout_s = 600` second wall-clock budget per method per instance; it does not override the distinct RP23, HW24, or REAL24 campaign manifests. The table maps paper-facing names to JSON keys, tool/pass sequences, artifact flags, and source classes. The older real-panel UCC default row is included only to define that frozen-control label in the main text and Appendix C; it is not one of the canonical methods used in the 600s external-baseline comparison.

The source classes in table XI expand as follows: Here “main” denotes membership in the canonical comparison campaign and does not necessarily imply inclusion in the two headline tables.

- main: canonical 600s external-baseline comparison.

- abl: Fourier-layer IR ablation.
- res: resource-consequence check.
- corr: large-instance correctness extension.
- irr: irrational independent-Pauli unbounded witness suite.
- status: provenance-only configuration excluded from numerical quality comparison when its predicate fails.
- real: official real-instance frozen panel.
- appx: older frozen-control appendix panels.

The new executable protocol and provenance are in `research/round6_major_revision/`: the witness declaration and exact-rank verifier, `unbounded_witness_results.json`, `rational_cp_round6_results.json`, `public_phase_folding_results.json`, `tket_paulisimp_results.json`, per-cell logs, and the environment snapshot. After global-phase alignment, the implemented check computes the maximum elementwise error  $\epsilon$  and requires  $\epsilon \leq \text{atol} + \text{rtol} \max_{i,j} |U_{ij}|$ . Thus  $\text{atol} = \text{rtol} = 10^{-9}$  gives an effective upper threshold of approximately  $2 \times 10^{-9}$  because unitary entries have magnitude at most one. Recorded maximum errors increase with repetition count  $r$ , while every row marked completed and valid satisfies this criterion. The irrational-witness PyZX 10k result passes at approximately  $1.05 \times 10^{-9}$ ; the 20k result at approximately  $2.68 \times 10^{-9}$  is a correctness failure. Timeout, predicate-error, unsupported-conversion, and correctness-failure statuses are kept distinct. The staq provenance records public commit `a2acd39` (with the full hash in the artifact) and the generic double-precision OpenQASM patch used for continuous-angle round trips.

#### 5. Certificate, Cut-Trace, and QRE Artifacts

The exact canonicalizer and symbolic checker are `certificates/canonicalize.py` and `certificates/symbolic_checker.py`. The frozen audit data/`frozen/certificate_audit.json` contains 12 accepted/rejected dense cross-check cases over widths one through six and 56 adversarial mutations. The checker and dense evaluator use separate implementations: the former propagates complete signed Clifford tableaus and exact rational- $\pi$  Pauli coefficients, while the latter constructs full numerical matrices and a projective distance. Only `completed_valid` records carrying an independent certificate enter numerical means; unsupported, predicate-error, and numerical-inconclusive records remain status-only.

The theorem-model reference compiler is `scripts/trace_cut_budget.py`. Its event schema records after every input update `committed_output_bytes`, the five cut fields from (3), pass index, RSS, elapsed nanoseconds, artifact paths, and digests. It deletes the explicit live

| Paper name                               | JSON key                 | Tool/pass sequence                                                                                                               | Artifact flags                                                               | Budget                     | Source classes       |
|------------------------------------------|--------------------------|----------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------|----------------------------|----------------------|
| semantic UCC (Fourier-layer IR enabled)  | semantic_ucc             | ucc.compile to {cx, rx, ry, rz, h} with Fourier-layer IR enabled                                                                 | UCC_DISABLE_FOURIER_LAYER_IR unset                                           | 600 s                      | main, abl, res, corr |
| upstream UCC v0.4.12                     | baseline_ucc             | ucc.compile loaded from the legacy baseline repository in research/compare_real_instances.py; target basis {cx, rx, ry, rz, h}   | legacy baseline repo; semantic candidate controls absent                     | 240 s in real-panel runner | real, appx           |
| pre-semantic default                     |                          |                                                                                                                                  |                                                                              |                            |                      |
| artifact UCC (Fourier-layer IR disabled) | no_fourier_ucc           | same artifact compile path with Fourier-layer IR disabled                                                                        | UCC_DISABLE_FOURIER_LAYER_IR=1                                               | 600 s                      | abl, res             |
| qiskit opt3                              | qiskit_opt3              | qiskit.transpile(..., optimization_level=3, layout_method="trivial", routing_method="none")                                      | none                                                                         | 600 s                      | main, abl, res       |
| PyZX configured bridge                   | pyzx_configured          | Qiskit QASM2 to PyZX, to_basic_gates, basic_optimization, Qiskit target-basis cleanup                                            | none                                                                         | 600 s                      | main                 |
| PyZX full_reduce                         | pyzx_full_reduce         | Qiskit QASM2 to PyZX graph, full_reduce, circuit extraction, Qiskit target-basis cleanup                                         | none                                                                         | 600 s                      | main                 |
| TKET FullPeephole                        | tket_fullpeephole        | Qiskit-to-TKET bridge when available, otherwise QASM fallback, then FullPeepholeOptimise                                         | none                                                                         | 600 s                      | main                 |
| TKET PauliSimp (unrebased configuration) | tket_paulisimp_unrebased | direct Qiskit-TKET bridge, then PauliSimp; retained only to record the predicate error                                           | none                                                                         | 600 s                      | status               |
| TKET PauliSimp (rebased)                 | tket_paulisimp_rebased   | direct bridge, DecomposeBoxes, AutoRebase({CX, RZ, RZ}), RemoveRedundancies, PauliSimp, RemoveRedundancies, target-basis cleanup | generic identical preamble                                                   | 600 s                      | main, irr            |
| TKET GuidedPauliSimp                     | tket_guided_paulisimp    | same TKET bridge/fallback sequence, then GuidedPauliSimp                                                                         | none                                                                         | 600 s                      | main                 |
| staq rotation folding                    | staq_rotation_folding    | Qiskit opt0 target-basis QASM2, public staq fold_rotations, QASM2 parse, target-basis cleanup                                    | no family labels; generic double-precision parser/serializer patch disclosed | 600 s                      | main, irr            |
| phase-polynomial reference               | phase_poly_reference     | recognize the $HD^TH$ witness, aggregate commuting rz/cp coefficients, then lower through Qiskit opt0                            | none                                                                         | 600 s                      | main, abl, res, corr |

TABLE XI. Method/provenance map for the canonical validation protocol. The direct Qiskit-to-TKET bridge was available in the validation environment; the QASM fallback is documented for reproducibility when that bridge is unavailable.

compiler-state object and resumes from only the committed prefix, five serialized fields, and unread input suffix. The frozen `research/cut_trace/` bundle was produced by the earlier serializer and lacks the per-event field files, so its 1,320 events and 24 summaries are historical only and are not accepted by the version-one verifier. The checked three-cut codec fixture and its manifests are in the project-root `instrumentation/example_trace/`; a full empirical curve requires rerunning this revised tracer.

The W7 runner and allocator are the project-root `scripts/run_qre_campaign.py` and `scripts/allocate_synthesis_error.py`; their total-error schema, two QEC profiles, and factory grid are under `qre/`. The 288-row Parquet/CSV, full angle records, analysis, figures, and manifests are under `data/frozen/`, `data/runs/w7-fixed-total-error-20260803-local/`, and `figures/`. The earlier single-profile runner and its six-row result are retained as historical artifacts under `research/fixed_budget_qre/` but are not the W7 main result. The grid synthesizer is built from public staq commit `a2acd39e60ed0e5f1978fa530df7e1e0c7cedf7a`; the disclosed artifact change only replaces the non-deterministic candidate-search seed with 20260801. The runner omits staq’s `-time` fast path because that path bypasses its internal detail/check branch; it instead measures subprocess time externally and requires `Check flag = 1` for every unique angle. The QRE is the explicit surface-code model family stated in the main text, not a vendor service or an independent architecture implementation.

## 6. Immutable Campaign Hierarchy

The machine-readable index `research/campaign_manifests/campaign_index.json` defines two base profiles and five noninterchangeable campaigns. RP23 is the pass-order probe, HW24 is the canonical line-backend cam-

paign, REAL24 is the official real-instance anti-regression panel, MR25 is the matched-representation matrix, and NF25 is the natural-kernel factorial ablation. Each prefix is unique and must appear in a reader-facing label before the source instance name.

Every campaign manifest records or explicitly marks unavailable the backend topology, native gate set, seed, timeout policy, memory cap, tool versions, artifact/freeze commit, bridge sequence, certificate policy, immutable config hash, and frozen dataset hash. Historical campaigns did not enforce a memory cap or independent certificate; their manifests say `not_enforced_historical` and “diagnostic only” rather than imputing a modern policy. In particular, HW24 uses 120 s for QPE/QAOA and 240 s for Grover, whereas RP23 uses 600 s. Those rows may share an MQT source-instance string but cannot share a short label or be averaged.

The configuration hierarchy is single-parent: a campaign inherits only its named base profile and then supplies explicit overrides. Backend and config hashes use canonical JSON. The frozen source snapshots `frozen_hardware_aware_rows.json` and `frozen_real_panel_rows.json` live beside the manifests, so later changes to a runner cannot silently change a displayed row.

## 7. Claim–Evidence Matrix

Every headline statement in the abstract, introduction, resource section, and conclusion is registered with one primary theorem, proposition, table, figure, or artifact test and an explicit scope. The machine-readable source is `research/claim_evidence/claim_evidence.json`; the compact reader map is table XII.

| ID  | claim location          | primary evidence              | authorized scope                                                                     |
|-----|-------------------------|-------------------------------|--------------------------------------------------------------------------------------|
| C01 | abstract/introduction   | definition 9 and equation (3) | formal model only                                                                    |
| C02 | abstract/conclusion     | theorem 1                     | declared family-relative and self-contained codes                                    |
| C03 | abstract/conclusion     | theorem 2                     | two-round masked-share stream                                                        |
| C04 | introduction/discussion | proposition 8                 | ordered-share family-relative codec                                                  |
| C05 | validation              | figure 4                      | 1,080 frozen cells; 24 cells per representation-compiler pair                        |
| C06 | validation              | figure 5                      | 592 formal-model reference-compiler runs                                             |
| C07 | validation/workloads    | table X and figure 12         | 9 families; $2^6$ factors; 3 scales and 3 seeds                                      |
| C08 | resource consequences   | table VIII                    | one disclosed surface-code model                                                     |
| C09 | compiler                | proposition 12                | declared symbolic domain only                                                        |
| C10 | validation              | figure 4 and section VIII B   | native and unified configurations; bridge loss reported separately                   |
| C11 | application workloads   | table IX                      | REAL24 campaign                                                                      |
| C12 | conclusion              | table IV                      | evidence-class separation                                                            |
| C13 | abstract/conclusion     | theorem 5                     | round-major forward scans; all 2p-1 crossings charged                                |
| C14 | abstract/discussion     | proposition 10                | family-relative p-block codec; compact-output corner for every r; full curve for r=2 |

TABLE XII. Claim–evidence matrix. The full JSON also contains the exact claim text and evidence class. A claim outside the displayed scope is not made by this paper.

## Appendix C: Supplementary Experiments

In this appendix, rows labeled “baseline UCC” are frozen-control rows: they denote the upstream UCC v0.4.12 pre-semantic default path defined in the provenance map of Appendix B, not a separate semantic method.

### 1. Exact-Semantics Witness Data

The following tables contain the full uniform-600s comparison data for the exact-semantics witnesses discussed in corollary 1 of the main text.

### 2. Fourier Witness Diagnostics

The tables in this subsection support the main-text Fourier witness without serving as additional theorem instances. The canonical theorem-linked CP witness is the finite-range case in table XIII; the generalized topology/angle rows are robustness diagnostics, and the external stress rows use the budgets recorded in their frozen source files.

### 3. Controlled Fixed-Basis Results

*a. Repeated QFT + QFT<sup>-1</sup> (100k gates)*

*b. Repeated QFT + QFT control (100k gates)*

These results support a representation-level claim: after controlling for the final target basis, there is still a large structural opportunity before basis lowering.

### *c. Systematic QFT + QFT<sup>-1</sup> external scaling*

The inverse-QFT scaling control is not the primary theorem witness, because `qiskit opt3` also finds the zero-gate circuit at smaller sizes. Its role is to test scalability when both paths return the same zero-gate output. As shown in table XXIII, the semantic artifact returns the zero-gate circuit at all tested sizes, is faster than `qiskit opt3` wherever `qiskit opt3` completes, and avoids the `qiskit opt3` timeouts at 50k and 100k. As an additional control, it also matches `qiskit_commutative_inverse` in output quality and is faster at every tested size.

### 4. Why Simple Pass Reordering Is Not Enough

To test whether the observed gains are mainly explained by a missing pass or pass-ordering issue inside `UCCDefaults`, we ran a controlled comparison between baseline UCC, several minimal `UCCDefaults`-style alternatives, and the semantic artifact. The minimal alternatives were `defaults_preinverse_only`, `defaults_reordered_only`, and `defaults_minimal_bundle`; the all-to-all cases were normalized to the same final target basis  $B = [\text{cx}, \text{rx}, \text{ry}, \text{rz}, \text{h}]$ , and the backend-aware cases used fixed transpiler seed 12345. This is a separate reorder-probe experiment, so its runtime numbers should be read independently from the frozen mainline tables below. Its hardware rows have reader labels `RP23-qpeexact-s12345-c89463b96` and `RP23-qaqa-s12345-c89463b96`; the canonical hardware campaign below instead uses the `HW24-...-c0f32747f` labels. The differing labels are different full experiment keys, not repeated measurements of one row.

Two qualifications are not visible in table XXIV itself. On `RP23-qaqa-s12345-c89463b96` the minimal bundle already matches the semantic artifact, so the reorder-only explanation is adequate there. On `RP23-qpeexact-s12345-c89463b96` it is not: the minimal bundle gives 1,893 gates, depth 576, `cx` 1,311, against 1,593, depth 420, `cx` 772 for the semantic artifact.

| Requested | Actual | semantic UCC | qiskit opt3          | PyZX full_reduce | TKET PauliSimp (rebased) | staq rotation folding | TKET GuidedPauliSimp |
|-----------|--------|--------------|----------------------|------------------|--------------------------|-----------------------|----------------------|
| 4,000     | 3,998  | 42/24/12     | 9588/5593/4788       | 107/59/31        | 52/29/19                 | 9188/4996/4788        | 13574/6392/4788      |
| 10,000    | 9,998  | 42/24/12     | 23988/13993/11988    | timeout          | 62/36/28                 | 22988/12496/11988     | 33974/15992/11988    |
| 20,000    | 19,998 | 42/24/12     | 47988/27993/23988    | timeout          | 59/27/20                 | 45988/24996/23988     | 67974/31992/23988    |
| 50,000    | 49,998 | 42/24/12     | 119988/69993/59988   | timeout          | 73/44/35                 | 114988/62496/59988    | 169974/79992/59988   |
| 100,000   | 99,998 | 42/24/12     | 239988/139993/119988 | timeout          | 66/37/28                 | 229988/124996/119988  | 339974/159992/119988 |

TABLE XIII. Uniform-600s comparison on the rational CP finite-range witness. Completed cells report gates/depth/cx; other entries are status outcomes and are not scored as structural wins. The family is certified only for  $R < T = 131040$ , with per-transducer tail constants. No external row is a formal theorem instance.

| Requested | Actual | semantic UCC | qiskit opt3        | PyZX full_reduce    | TKET PauliSimp (rebased) | staq rotation folding | TKET GuidedPauliSimp |
|-----------|--------|--------------|--------------------|---------------------|--------------------------|-----------------------|----------------------|
| 4,000     | 3,998  | 27/18/6      | 4799/2801/2394     | 38/19/8             | 36/14/6                  | 3402/2201/2394        | 8786/4000/2394       |
| 10,000    | 9,998  | 27/18/6      | 12000/7001/5994    | 48/33/12            | 35/15/8                  | 8502/5501/5994        | 21986/10000/5994     |
| 20,000    | 19,998 | 27/18/6      | 24000/14001/11994  | correctness failure | 38/16/8                  | 17002/11001/11994     | 43986/20000/11994    |
| 50,000    | 49,998 | 27/18/6      | 60000/35001/29994  | timeout             | 36/15/8                  | 42502/27501/29994     | 109986/50000/29994   |
| 100,000   | 99,998 | 27/18/6      | 120000/70001/59994 | timeout             | 37/16/8                  | 85002/55001/59994     | 219986/100000/59994  |

TABLE XIV. Uniform-600s comparison on the irrational independent-Pauli witness. Completed cells report gates/depth/cx; other entries are status outcomes and are not scored as structural wins. The exact rank and irrational-angle certificate, rather than this numerical table, instantiate corollary 1.

## 5. Public Benchmark Suite Results

To move beyond official library examples, we added two public benchmark sources: **MQT Bench** and **SupermarQ**.

*a. MQT Bench.* Representative results are:

- **mqt\_qpeexact\_32**: qiskit opt3 = 1,891 gates, baseline UCC = 5,494, semantic UCC = 1,891;
- **mqt\_qpeinexact\_24**: qiskit opt3 = 1,206 gates, baseline UCC = 3,682, semantic UCC = 1,206;
- **mqt\_ae\_8**: qiskit opt3 = 207 gates, baseline UCC = 542, semantic UCC = 207;
- **mqt\_draper\_qft\_adder\_32**: qiskit opt3 = 1,630 gates, baseline UCC = 5,278, semantic UCC = 1,630;
- **mqt\_qaoa\_32**: qiskit opt3 = 1,422 gates, baseline UCC = 6,310, semantic UCC = 1,422;
- **mqt\_grover\_20**: qiskit opt3 = 3,848,810 gates in 38.043 s, baseline UCC times out at  $> 120$  s, and semantic UCC reaches the same output quality in 39.123 s.

Thus, on **MQT Bench**, the semantic artifact repeatedly recovers qiskit opt3-level output quality and clearly improves over baseline UCC. The added **qpeinexact**, **ae**, and **draper\_qft\_adder** cases are especially relevant because they extend the phase-ladder/Fourier-layer story beyond one exact-QPE benchmark into inexact QPE, amplitude estimation, and QFT-based arithmetic. Runtime still remains weaker on the hardest Grover-style case.

Additional mirrored/conjugation positive cases make the mirrored-shell heuristic more concrete. On the official **grover\_real** instance, baseline UCC returns 287,047 gates while the semantic artifact and qiskit opt3 both return 79,029. On **MQT Bench** Grover instances, the same pattern persists across scale: **mqt\_grover\_8** gives

18,992  $\rightarrow$  5,170, **mqt\_grover\_12** gives 259,797  $\rightarrow$  71,706, and **mqt\_grover\_16** gives 2,112,285  $\rightarrow$  581,098, where in each case the semantic artifact records the same gate count as qiskit opt3 while clearly improving on baseline UCC. The large fixed-basis mirrored Grover benchmark shows the same parity pattern: baseline UCC gives 4,122,006 gates, while the semantic artifact and qiskit opt3 both give 1,158,011. These cases make the mirrored/conjugation line less dependent on the routed timeout-recovery interpretation of **hw\_mqt\_grover\_20**.

A dedicated mirrored/conjugation scaling sweep makes this support more systematic. On public **MQT Bench** Grover instances, the semantic artifact records the same gate count as qiskit opt3 at sizes 8, 12, and 16, reducing baseline UCC from 18,992 to 5,170 gates, from 259,797 to 71,706 gates, and from 2,112,285 to 581,098 gates, respectively. On the fixed-basis **grover\_mirrored** repeated-shell scaling sweep, the same parity pattern holds at 10k, 20k, 50k, and 100k input gates: baseline UCC gives 412,206, 824,406, 2,061,006, and 4,122,006 output gates, while the semantic artifact and qiskit opt3 give 115,811, 231,611, 579,011, and 1,158,011 gates. Thus this second structural family targeted by the artifact has empirical positive cases in both public Grover scaling and large repeated-shell fixed-basis scaling.

*b. SupermarQ.* Representative results are:

- **supermarq\_mermin\_bell\_8**: baseline UCC = 386 gates, depth 135; semantic UCC = 76 gates, depth 44; qiskit opt3 = 76 gates, depth 44;
- **supermarq\_qaoa\_vanilla\_12**: baseline UCC = 947 gates, depth 225; semantic UCC = 222 gates, depth 80; qiskit opt3 = 222 gates, depth 80;
- **supermarq\_hamiltonian\_sim\_8**: baseline UCC = 71 gates, depth 36; semantic UCC = 51 gates, depth 24; qiskit opt3 = 51 gates, depth 24.

This second public benchmark source strengthens the

| Topology      | Angles | Points | qiskit ok/TO | Semantic completed | Semantic output signatures                     |
|---------------|--------|--------|--------------|--------------------|------------------------------------------------|
| chain CP      | 3      | 12     | 9 / 3        | all 12             | 23–27 gates, 17–18 depth, 6 cx                 |
| ring CP       | 3      | 12     | 9 / 3        | all 12             | 32 gates, 23 depth, 8 cx                       |
| sparse CP 0.5 | 3      | 12     | 6 / 6        | all 12             | 25–27 gates, 11–12 depth, 6 cx                 |
| full-pair CP  | 2      | 6      | 4 / 2        | all 6              | 42 gates, 24 depth, 12 cx                      |
| Total points  | –      | 42     | 28 / 14      | all rows           | repetition-independent at fixed topology/angle |

TABLE XV. Generalized fixed-width Fourier witness suite. Each point is a paired run of semantic UCC and `qiskit opt3` on an  $HD^rH$  circuit at  $n = 4$ , with requested sizes up to 50k gates. The rows are filtered from the recorded 284-row generalized suite by the  $n = 4$  topology/angle/size slice. Recorded `qiskit opt3` timeout cells in that source use legacy 60s or 90s budgets, so the table is a robustness diagnostic rather than a uniform-600s comparison.

| Width ( $n$ ) | Checks  | Output | Gates | Depth | CX Count | Seed Variance |
|---------------|---------|--------|-------|-------|----------|---------------|
| 4             | 15 / 15 | 42     | 24    | 12    | 0        |               |
| 5             | 15 / 15 | 65     | 32    | 20    | 0        |               |
| 6             | 15 / 15 | 93     | 40    | 30    | 0        |               |

TABLE XVI. Correctness and seed-robustness certificate for the full-pair Fourier witness. The correctness experiment checks 3 widths  $\times$  3 requested sizes  $\times$  5 angle seeds using `Qiskit Operator.equiv`; all 45 configurations are equivalent to the original circuit up to global phase. The seed-robustness experiment contains 90 rows including `qiskit opt3`; semantic UCC has zero output gate-count variance across the five tested nonresonant angle seeds at each width.

paper in two ways. It shows clear anti-regression behavior on additional nontrivial families, and it also exposes a realistic limitation: the semantic artifact is not uniformly better on every benchmark family.

## 6. Hardware-Aware Results

We added a backend-aware comparison using a fixed 20-qubit bidirectional line backend with native operations `u`, `sx`, `p`, `cx`, `measure`, and `id`. The benchmark source is `MQT Bench`. The reader-facing rows are `HW24-qpeexact-s12345-c0f32747f`, `HW24-qaoa-s12345-c0f32747f`, and `HW24-grover-s12345-c0f32747f`; the full source instance remains part of each registry primary key.

*a. HW24-qpeexact-s12345-c0f32747f.* This is a clean fixed-seed parity case: the semantic artifact matches `qiskit opt3` in total gates, depth, and `cx` count, while still remaining much better than baseline UCC on all structural metrics. The tradeoff is runtime rather than output quality.

*b. HW24-qaoa-s12345-c0f32747f.* This is the clearest hardware-aware external-baseline improvement in the paper: the semantic artifact improves on `qiskit opt3` in total gates, depth, and `cx` count, while also improving on baseline UCC in total gates, depth, and `cx` count.

*c. HW24-grover-s12345-c0f32747f.* `qiskit opt3` and baseline UCC still time out under this setting, but the semantic artifact returns a result in 1.999s via a stronger mirrored/self-inverse repeated-run backend fallback. The resulting circuit is still extremely large,

but the selector lowers total gates, depth, and `cx` simultaneously (to 6,467,857 total gates, depth 4,094,196, and `cx` count 3,045,048), so we treat this as a robustness result rather than a main competitive routed-quality claim.

## 7. Stability and Seed Sensitivity

We measured repeated-run stability on the frozen artifact in two ways:

1. five repeated runs for the official real instances and the canonical hardware-aware seed,
2. a five-seed sweep for the backend-aware `MQT Bench` cases.

The main conclusions are:

- the official real-instance outputs are deterministic across repeated runs,
- the hardware-aware outputs are deterministic for each fixed seed,
- `hw_mqt_qaoa_20` remains better than `qiskit opt3` in total gates, depth, and `cx` count across all five tested seeds 0, 1, 42, 12345, and 54321,
- `hw_mqt_qpeexact_20` is now a fixed-seed parity case rather than a strict improvement.

Thus, the hardware-aware claim is a stable workload-family result across a small fixed seed portfolio, rather than a single favorable seed.

## 8. Scaling and Ablation

The main scaling trends are summarized in figure 14. The size sweep covers 4k, 10k, 20k, 50k, and 100k inputs for QFT-inverse, QPE-style, QAOA-ring, and mirrored-Grover families; the QAOA-ring block uses 20 qubits. The most important observations are stable quality behavior across scale and improving repeated-family runtime once the semantic representation and dispatch controller are enabled. The method is most effective on exact or

| Method                                  | Requested | Actual | Status    | Gates/depth/cx | Certificate                      |
|-----------------------------------------|-----------|--------|-----------|----------------|----------------------------------|
| semantic UCC (Fourier-layer IR enabled) | 100,000   | 99,998 | completed | 42/24/12       | <code>Operator.equiv=True</code> |
| phase-polynomial reference              | 100,000   | 99,998 | completed | 42/24/12       | <code>Operator.equiv=True</code> |

TABLE XVII. Large-instance correctness extension on the canonical  $n = 4$ ,  $r = 9999$  Fourier witness. The check is a finite-instance Qiskit `Operator.equiv` certificate for the two completed compressed outputs only; it does not certify the generalized topology/angle suite or prove an algebraic identity for all  $r$ .

| Width | Requested | semantic UCC | qiskit | opt3   | PyZX | bridge | TKET | FullPeephole |
|-------|-----------|--------------|--------|--------|------|--------|------|--------------|
| 5     | 4,000     | 65           | 11,711 | 14,640 |      |        |      | 30,369       |
| 5     | 10,000    | 65           | 29,315 | 36,640 |      |        |      | timeout      |
| 6     | 4,000     | 93           | 12,108 | 15,321 |      |        |      | 33,881       |
| 6     | 10,000    | 93           | 30,417 | 38,487 |      |        |      | timeout      |

TABLE XVIII. Representative external stress tests on the width-axis full-pair Fourier witness. Entries report output gates after the configured bridge pipeline and final target-basis compilation. PyZX completes all four runs but does not recover the bounded semantic forms 65 and 93; TKET completes the 4k runs with larger outputs and times out at 10k. The recorded budget in this stress file is 300s.

highly repeated structure: `qft_inverse` remains maximally strong across all tested sizes, `qpe_style` remains a strong positive fixed-basis case, `qaoa_ring` consistently avoids UCC inflation, and the mirrored-Grover family continues to match Qiskit quality at 100k while remaining more expensive. Direct spot checks on the semantic artifact give `qpe_style` = 82,200 gates at 50k in 2.504 s and 164,400 gates at 100k in 4.636 s, reflecting the phase-ladder plus repeated-run controller improvements.

The non-monotonic semantic runtime between the 20k and 50k points is a dispatch effect visible in the implementation, not an inferred hardware effect. The 20,000-gate full-preset limit `_MAX_FULL_PRESET_REFERENCE_SIZE` allows the 19,998-gate input to enter the full preset-reference comparison, whereas the 49,998-gate input exceeds that threshold and uses the cheaper repeated-prefix/hierarchical route. The corresponding canonical

witness runtimes are 28.314 s and 20.934 s. This threshold explains the local drop; it is not evidence of asymptotically decreasing compilation time.

The ablation picture is narrow. The dominant practical component is conservative candidate control on recoverability-sensitive families, most clearly on the 10,000-input-gate `qaoa_ring` row, where the full configuration returns 10,000 gates while the no-candidate-selection variant returns 30,055. By contrast, `qft_inverse`, `qft_control`, `qpe_style`, and `grover_mirrored` are less sensitive to the removal of individual structural rules. This reinforces the paper’s current framing: the main contribution is a bounded semantic representation and dispatch controller that preserve recoverability before flat materialization, rather than a single local identity.

- 
- [1] E. Farhi, J. Goldstone, and S. Gutmann, A quantum approximate optimization algorithm, arXiv preprint arXiv:1411.4028 (2014), arXiv:1411.4028 [quant-ph].
  - [2] A. Peruzzo, J. McClean, P. Shadbolt, M.-H. Yung, X.-Q. Zhou, P. J. Love, A. Aspuru-Guzik, and J. L. O’Brien, A variational eigenvalue solver on a photonic quantum processor, *Nature Communications* **5**, 4213 (2014).
  - [3] M. Cerezo, A. Arrasmith, R. Babbush, S. C. Benjamin, S. Endo, K. Fujii, J. R. McClean, K. Mitarai, X. Yuan, L. Cincio, and P. J. Coles, Variational quantum algorithms, *Nature Reviews Physics* **3**, 625 (2021).
  - [4] A. M. Childs, Y. Su, M. C. Tran, N. Wiebe, and S. Zhu, Theory of trotter error with commutator scaling, *Physical Review X* **11**, 011020 (2021).
  - [5] C. M. Dawson and M. A. Nielsen, The solovay–kitaev algorithm, *Quantum Information and Computation* **6**, 81 (2006).
  - [6] V. Kliuchnikov, D. Maslov, and M. Mosca, Fast and efficient exact synthesis of single-qubit unitaries generated by clifford and T gates, *Quantum Information and Computation* **13**, 607 (2013).
  - [7] P. Selinger, Efficient clifford+T approximation of single-qubit operators, *Quantum Information and Computation* **15**, 159 (2015).
  - [8] N. J. Ross and P. Selinger, Optimal ancilla-free clifford+t approximation of z-rotations, *Quantum Information and Computation* **16**, 901 (2016), arXiv:1403.2975 [quant-ph].
  - [9] J. J. Wallman and J. Emerson, Noise tailoring for scalable quantum computation via randomized compiling, *Physical Review A* **94**, 052325 (2016), arXiv:1512.01098.
  - [10] M. Fellous-Asiani, J. H. Chai, Y. Thonnart, H. K. Ng, R. S. Whitney, and A. Auffèves, Optimizing resource efficiencies for scalable full-stack quantum computers, *PRX Quantum* **4**, 040319 (2023).
  - [11] G. Watkins, H. M. Nguyen, K. Watkins, S. Pearce, H.-K. Lau, and A. Paler, A high performance compiler for very large scale surface code computations, *Quantum* **8**, 1354 (2024), arXiv:2302.02459 [quant-ph].
  - [12] M. E. Beverland, P. Murali, M. Troyer, K. M. Svore, T. Hoeffler, V. Kliuchnikov, G. H. Low, M. Soeken,



| Topology      | semantic UCC | qiskit opt3     | PyZX bridge     | TKET              | FullPeephole |
|---------------|--------------|-----------------|-----------------|-------------------|--------------|
| chain CP      | 27           | 6,851 / 17,136  | 10,838 / 27,121 | 20,555 / timeout  |              |
| ring CP       | 32           | 7,996 / 19,996  | 11,984 / 29,984 | 23,987 / 59,987   |              |
| sparse CP 0.5 | 27           | 6,851 / 17,136  | 10,838 / 27,121 | timeout / timeout |              |
| full-pair CP  | 42           | 10,783 / 26,983 | 13,574 / 33,974 | timeout / timeout |              |

TABLE XIX. Topology-axis external stress test at  $n = 4$  with nonresonant seed 0. Each external-baseline cell reports output gates at requested sizes 4k / 10k. The deterministic sparse-CP draw uses the same number of CP edges as chain CP and is graph-isomorphic to a chain at  $n = 4$ , which explains the matching `qiskit opt3` and PyZX gate totals for those two rows. The recorded budget in this stress file is 300s.

| Pipeline                                    | Representation tested                                                                                           | Observed behavior in the disclosed suites                                                                                                                                                                                        | Interpretation                                                                                          |
|---------------------------------------------|-----------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| semantic UCC<br>(Fourier-layer IR enabled)  | pre-basis Fourier-layer term                                                                                    | completed 5/5; 42 gates, depth 24, 12 cx at every tested scale                                                                                                                                                                   | semantic side                                                                                           |
| phase-polynomial<br>reference               | same Qiskit input; only<br>commuting-diagonal<br>coefficient aggregation                                        | completed 5/5; 42 gates, depth 24, 12 cx at 4k–100k                                                                                                                                                                              | authors' protocol-specific<br>constructive control                                                      |
| <code>qiskit opt3</code>                    | preset target-basis pipeline                                                                                    | completed 5/5; grows from 9,588 gates at 4k to 239,988 gates at 100k                                                                                                                                                             | strong completed baseline,<br>but not bounded                                                           |
| PyZX configured bridge                      | Qiskit–QASM–PyZX basic<br>optimization–Qiskit                                                                   | completed 4/5; grows through 169,974 gates at 50k and times out at 100k                                                                                                                                                          | configured bridge does not<br>recover bounded form                                                      |
| PyZX <code>full_reduce</code>               | QASM to PyZX graph;<br><code>full_reduce</code> ; extract and<br>lower                                          | rational witness: completes only at 4k with 107 gates, depth 59, 31 cx; timeout at 10k–100k. Irrational witness: completes and validates at 4k and 10k with 38 and 48 gates; correctness failure at 20k; timeout at 50k and 100k | strong small-scale<br>simplification; correctness<br>failure remains distinct<br>from timeout and error |
| TKET FullPeephole                           | direct Qiskit–TKET bridge<br>to FullPeepholeOptimise                                                            | timeout on all five canonical sizes under the 600s budget                                                                                                                                                                        | no completed bounded<br>recovery in this protocol                                                       |
| TKET PauliSimp<br>(unrebased configuration) | direct Qiskit–TKET bridge<br>to PauliSimp                                                                       | predicate error on all five sizes for both witness families                                                                                                                                                                      | provenance-only<br>configuration; not a<br>structural loss                                              |
| TKET PauliSimp<br>(rebased)                 | generic DecomposeBoxes,<br>AutoRebase to {CX,Rz,Rx},<br>RemoveRedundancies,<br>PauliSimp,<br>RemoveRedundancies | completes and validates at all five scales: 52–73 gates on the rational witness and 35–38 on the irrational witness                                                                                                              | generic aggregate<br>reconstruction crosses the<br>formal boundary                                      |
| staq rotation folding                       | generic target-basis QASM2<br>through public<br><code>fold_rotations</code>                                     | completes and validates at all five scales, but grows from 9,188 to 229,988 gates on the rational witness and from 3,402 to 85,002 on the irrational witness                                                                     | completed public baseline;<br>growing output is not an<br>impossibility result                          |
| TKET GuidedPauliSimp                        | direct Qiskit–TKET bridge<br>to GuidedPauliSimp                                                                 | completed 5/5 on each witness; grows from 13,574 to 339,974 gates on the rational witness and from 8,786 to 219,986 on the irrational witness                                                                                    | successful pipeline, but not<br>bounded                                                                 |
| artifact UCC<br>(Fourier-layer IR disabled) | same artifact with<br>Fourier-layer semantic path<br>disabled                                                   | in ablation/resource runs, completes at 4k, 10k, and 20k with 9,588, 23,988, and 47,988 gates; times out at 50k and 100k                                                                                                         | semantic-path causality<br>check; completed rows<br>materialize large outputs                           |

TABLE XX. External-pipeline status matrix and constructive controls for the rational CP and irrational independent-Pauli witnesses. Completed outputs and status outcomes are separated. The rebased and unrebased PauliSimp configurations are distinct, and no row is treated as a formal transducer instance.

| Method           | Output Gates | Output Depth |
|------------------|--------------|--------------|
| translation only | 400,000      | 142,500      |
| baseline UCC     | 968,784      | 276,266      |
| semantic UCC     | 0            | 0            |

TABLE XXI. Controlled fixed-basis results for repeated QFT + QFT<sup>-1</sup>.

| Method           | Output Gates | Output Depth |
|------------------|--------------|--------------|
| translation only | 400,000      | 140,000      |
| baseline UCC     | 1,207,504    | 317,501      |
| semantic UCC     | 347,500      | 125,000      |

TABLE XXII. Controlled fixed-basis results for repeated QFT + QFT.

A. Sundaram, and A. Vaschillo, Assessing requirements

- to scale to practical quantum advantage, arXiv preprint arXiv:2211.07629 (2022), arXiv:2211.07629 [quant-ph].
- [13] S. Bravyi and A. Kitaev, Universal quantum computation with ideal clifford gates and noisy ancillas, *Physical Review A* **71**, 022316 (2005).
- [14] D. Litinski, Magic state distillation: Not as costly as you think, *Quantum* **3**, 205 (2019), arXiv:1905.06903 [quant-ph].
- [15] M. Amy, D. Maslov, and M. Mosca, Polynomial-time t-depth optimization of clifford+t circuits via matroid partitioning, *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* **33**, 1476 (2014).
- [16] Y. Nam, N. J. Ross, Y. Su, A. M. Childs, and D. Maslov, Automated optimization of large quantum circuits with continuous parameters, *npj Quantum Information* **4**, 23 (2018).
- [17] A. Cowtan, S. Dilkes, R. Duncan, W. Simmons, and S. Sivarajah, Phase gadget synthesis for shallow circuits, in *Proceedings 16th International Conference on Quantum Physics and Logic (QPL 2019)*, Electronic Proceedings

| Input gates | qiskit opt3 |          | qiskit_commutative_inverse |         | semantic UCC |
|-------------|-------------|----------|----------------------------|---------|--------------|
|             | Gates       | Runtime  | Gates                      | Runtime | Runtime      |
| 4,000       | 0           | 1.052 s  | 0                          | 0.082 s | 0.063 s      |
| 10,000      | 0           | 6.090 s  | 0                          | 0.182 s | 0.124 s      |
| 20,000      | 0           | 24.596 s | 0                          | 0.341 s | 0.279 s      |
| 50,000      | timeout     | > 90 s   | 0                          | 0.979 s | 0.829 s      |
| 100,000     | timeout     | > 90 s   | 0                          | 2.271 s | 1.651 s      |

TABLE XXIII. External scaling on repeated  $\text{QFT} + \text{QFT}^{-1}$ . The semantic UCC artifact returns zero gates at all tested sizes; the table reports runtime because output quality is identical to the zero-gate reference wherever the references complete. Timeout cells in this source use the recorded 90 s budget.

| Case                              | Method                   | Gates   | Depth   | CX count | Runtime  |
|-----------------------------------|--------------------------|---------|---------|----------|----------|
| phase_estimation_real             | baseline UCC             | 471,819 | 328,271 | 58,491   | 75.538 s |
|                                   | defaults_preinverse_only | 485,780 | 332,924 | 82,925   | 80.045 s |
|                                   | defaults_reordered_only  | 485,780 | 332,924 | 82,925   | 80.005 s |
|                                   | defaults_minimal_bundle  | 485,780 | 332,924 | 82,925   | 80.084 s |
|                                   | semantic UCC             | 166,805 | 119,204 | 58,491   | 1.989 s  |
|                                   | qiskit opt3              | 166,805 | 119,204 | 58,491   | 1.254 s  |
| qaoa_ring_100k                    | baseline UCC             | 300,055 | 163,055 | 40,000   | 2.431 s  |
|                                   | defaults_preinverse_only | 434,104 | 217,962 | 40,000   | 11.352 s |
|                                   | defaults_reordered_only  | 445,216 | 217,962 | 40,000   | 10.120 s |
|                                   | defaults_minimal_bundle  | 434,104 | 217,962 | 40,000   | 10.818 s |
|                                   | semantic UCC             | 100,000 | 62,000  | 40,000   | 3.131 s  |
|                                   | qiskit opt3              | 100,000 | 62,000  | 40,000   | 0.307 s  |
| RP23-qpeexact<br>s12345-c89463b96 | baseline UCC             | 2,077   | 626     | 1,262    | 0.126 s  |
|                                   | defaults_preinverse_only | 1,893   | 576     | 1,311    | 0.292 s  |
|                                   | defaults_reordered_only  | 1,900   | 582     | 1,312    | 0.289 s  |
|                                   | defaults_minimal_bundle  | 1,893   | 576     | 1,311    | 0.287 s  |
|                                   | semantic UCC             | 1,593   | 420     | 772      | 0.296 s  |
|                                   | qiskit opt3              | 1,572   | 407     | 812      | 0.033 s  |
| RP23-qaoa<br>s12345-c89463b96     | baseline UCC             | 2,088   | 522     | 1,439    | 0.098 s  |
|                                   | defaults_preinverse_only | 2,050   | 487     | 1,458    | 0.270 s  |
|                                   | defaults_reordered_only  | 2,050   | 487     | 1,458    | 0.269 s  |
|                                   | defaults_minimal_bundle  | 2,050   | 487     | 1,458    | 0.271 s  |
|                                   | semantic UCC             | 2,050   | 487     | 1,458    | 0.278 s  |
|                                   | qiskit opt3              | 2,091   | 532     | 1,530    | 0.048 s  |

TABLE XXIV. Controlled comparison of baseline UCC, minimal UCCDefaults-style alternatives, the semantic artifact, and qiskit opt3. Small default-pipeline fixes explain part of the improvement on some cases, but they do not account for the main gains on the real and structured workloads. The overlapping phase\_estimation\_real rows use the per-method output cx counts from the recorded real-instance data. The keyed hardware rows are generated from the frozen registry; their RP23 labels distinguish this reorder-probe configuration from the HW24 campaign below.

| Method       | Output Gates | Output Depth | CX Count | Runtime |
|--------------|--------------|--------------|----------|---------|
| qiskit opt3  | 1,572        | 407          | 812      | 0.056 s |
| baseline UCC | 1,976        | 648          | 1,350    | 0.174 s |
| semantic UCC | 1,572        | 407          | 812      | 0.235 s |

TABLE XXV. Canonical hardware-aware results for reader label HW24-qpeexact-s12345-c0f32747f, generated from the frozen keyed dataset.

| Method       | Output Gates | Output Depth | CX Count | Runtime |
|--------------|--------------|--------------|----------|---------|
| qiskit opt3  | 2,091        | 532          | 1,530    | 0.078 s |
| baseline UCC | 2,057        | 513          | 1,503    | 0.141 s |
| semantic UCC | 2,050        | 487          | 1,458    | 0.112 s |

TABLE XXVI. Canonical hardware-aware results for reader label HW24-qaoa-s12345-c0f32747f, generated from the frozen keyed dataset.

- in Theoretical Computer Science, Vol. 318 (2020) pp. 213–228, arXiv:1906.01734 [quant-ph].
- [18] M. Amy, P. Azimzadeh, and M. Mosca, On the controlled-NOT complexity of controlled-NOT-phase circuits, *Quantum Science and Technology* **4**, 015002 (2018).
- [19] N. de Beaudrap, X. Bian, and Q. Wang, Techniques to reduce  $\pi/4$ -parity-phase circuits, motivated by the ZX calculus, in *Proceedings 16th International Conference on Quantum Physics and Logic (QPL 2019)*, Electronic Proceedings in Theoretical Computer Science, Vol. 318

- (2020) pp. 131–149.
- [20] R. Duncan, A. Kissinger, S. Perdrix, and J. van de Wetering, Graph-theoretic simplification of quantum circuits with the zx-calculus, *Quantum* **4**, 279 (2020), arXiv:1902.03178 [quant-ph].
- [21] A. Kissinger and J. van de Wetering, Reducing the number of non-clifford gates in quantum circuits, *Physical Review A* **102**, 022406 (2020), arXiv:1903.10477 [quant-ph].
- [22] A. Kissinger and J. van de Wetering, Pyzx: Large scale automated diagrammatic reasoning, in *Proceedings 16th*




FIG. 14. Scaling behavior on representative structured families. Left: output gate count versus input gate count for upstream default (baseline UCC) and the semantic UCC artifact. Right: runtime versus input gate count. The semantic path preserves clear quality advantages across scale, while runtime behavior remains family-dependent.

- International Conference on Quantum Physics and Logic (QPL 2019)*, Electronic Proceedings in Theoretical Computer Science, Vol. 318 (2020) pp. 229–241.
- [23] J. van de Wetering, ZX-calculus for the working quantum computer scientist, arXiv preprint arXiv:2012.13966 (2020), arXiv:2012.13966 [quant-ph].
- [24] S. Martiel and T. Goubault de Brugière, Architecture aware compilation of quantum circuits via lazy synthesis, *Quantum* **6**, 729 (2022), arXiv:2012.09663 [quant-ph].
- [25] A. Meijer-van de Griend and R. Duncan, Architecture-aware synthesis of phase polynomials for NISQ devices, in *Proceedings 19th International Conference on Quantum Physics and Logic (QPL 2022)*, Electronic Proceedings in Theoretical Computer Science, Vol. 394 (2023) pp. 116–140.
- [26] L. Rieseboos, X. Fu, S. Varsamopoulos, C. G. Almudever, and K. Bertels, Pauli frames for quantum computer architectures, in *Proceedings of the 54th Annual Design Automation Conference (DAC 2017)* (2017) pp. 76:1–76:6.
- [27] A. Hashim, R. K. Naik, A. Morvan, J.-L. Ville, B. Mitchell, J. M. Kreikebaum, M. Davis, E. Smith, C. Iancu, K. P. O’Brien, I. Hincks, J. J. Wallman, J. Emerson, and I. Siddiqi, Randomized compiling for scalable quantum computing on a noisy superconducting quantum processor, *Physical Review X* **11**, 041039 (2021), arXiv:2010.00215.
- [28] E. Knill, Quantum computing with realistically noisy devices, *Nature* **434**, 39 (2005).
- [29] A. Javadi-Abhari, M. Treinish, K. Krsulich, C. J. Wood, J. Lishman, J. Gacon, S. Martiel, P. D. Nation, L. S. Bishop, A. W. Cross, B. R. Johnson, and J. M. Gambetta, Quantum computing with qiskit, arXiv preprint arXiv:2405.08810 (2024), arXiv:2405.08810 [quant-ph].
- [30] E. Kushilevitz and N. Nisan, *Communication Complexity* (Cambridge University Press, 1997).
- [31] M. Amy and V. Gheorghiu, staq—a full-stack quantum processing toolkit, *Quantum Science and Technology* **5**, 034016 (2020), arXiv:1912.06070 [quant-ph].
- [32] A. G. Fowler, M. Mariantoni, J. M. Martinis, and A. N. Cleland, Surface codes: Towards practical large-scale quantum computation, *Physical Review A* **86**, 032324 (2012).
- [33] D. Horsman, A. G. Fowler, S. Devitt, and R. Van Meter, Surface code quantum computing by lattice surgery, *New Journal of Physics* **14**, 123011 (2012).
- [34] D. Litinski, A game of surface codes: Large-scale quantum computing with lattice surgery, *Quantum* **3**, 128 (2019), arXiv:1808.02892 [quant-ph].
- [35] C. Gidney and M. Ekerå, How to factor 2048 bit rsa integers in 8 hours using 20 million noisy qubits, *Quantum* **5**, 433 (2021), arXiv:1905.09749 [quant-ph].
- [36] A. Li, S. Stein, S. Krishnamoorthy, and J. Ang, QASM-Bench: A low-level quantum benchmark suite for NISQ evaluation and simulation, *ACM Transactions on Quantum Computing* **4**, 10 (2023).
- [37] N. Quetschlich, L. Burgholzer, and R. Wille, Mqt bench: Benchmarking software and design automation tools for quantum computing, *Quantum* **7**, 1062 (2023), arXiv:2204.13719 [quant-ph].
- [38] A. Cowtan, S. Dilkes, R. Duncan, A. Krajenbrink, W. Simmons, and S. Sivarajah, On the qubit routing problem, in *14th Conference on the Theory of Quantum Computation, Communication and Cryptography (TQC 2019)*, Leibniz International Proceedings in Informatics, Vol. 135 (2019) pp. 5:1–5:32.
- [39] G. Li, Y. Ding, and Y. Xie, Tackling the qubit mapping problem for NISQ-era quantum devices, in *Proceedings of the 24th International Conference on Architectural Support for Programming Languages and Operating Systems*

- (*ASPLOS 2019*) (2019) pp. 1001–1014.
- [40] S. Sivarajah, S. Dilkes, A. Cowtan, W. Simmons, A. Edgington, and R. Duncan, t|ket>: A retargetable compiler for nisq devices, *Quantum Science and Technology* **6**, 014003 (2021).
  - [41] A. Kissinger and J. van de Wetering, Simulating quantum circuits with ZX-calculus reduced stabiliser decompositions, *Quantum Science and Technology* **7**, 044001 (2022).
  - [42] V. Gheorghiu, M. Mosca, and P. Mukhopadhyay, T-count and T-depth of any multi-qubit unitary, *npj Quantum Information* **8**, 141 (2022).
  - [43] D. S. Steiger, T. Häner, and M. Troyer, Projectq: An open source software framework for quantum computing, *Quantum* **2**, 49 (2018), arXiv:1612.08091 [quant-ph].
  - [44] K. Hietala, R. Rand, S.-H. Hung, X. Wu, and M. Hicks, A verified optimizer for quantum circuits, *Proceedings of the ACM on Programming Languages* **5**, 37 (2021).
  - [45] D. Mills, S. Sivarajah, T. L. Scholten, and R. Duncan, Application-motivated, holistic benchmarking of a full quantum computing stack, *Quantum* **5**, 415 (2021), arXiv:2006.01273 [quant-ph].
  - [46] A. Cross, A. Javadi-Abhari, T. Alexander, N. de Beaudrap, L. S. Bishop, S. Heidel, C. A. Ryan, P. Sivarajah, J. Smolin, J. M. Gambetta, and B. R. Johnson, OpenQASM 3: A broader and deeper quantum assembly language, *ACM Transactions on Quantum Computing* **3**, 12 (2022).
  - [47] A. Peduri, S. Bhat, and T. Grosser, QSSA: An SSA-based IR for quantum computing, in *Proceedings of the 31st ACM SIGPLAN International Conference on Compiler Construction (CC 2022)* (2022) pp. 2–14.
  - [48] A. McCaskey and T. Nguyen, A MLIR dialect for quantum assembly languages, in *2021 IEEE International Conference on Quantum Computing and Engineering (QCE)* (2021) pp. 255–264.
  - [49] F. Kreppel, C. Melzer, D. Olvera Millán, J. Wagner, J. Hilder, U. Poschinger, F. Schmidt-Kaler, and A. Brinkmann, Quantum circuit compiler for a shuttling-based trapped-ion quantum computer, *Quantum* **7**, 1176 (2023), arXiv:2207.01964 [quant-ph].
  - [50] D. B. Tan, D. Bluvstein, M. D. Lukin, and J. Cong, Compiling quantum circuits for dynamically field-programmable neutral atoms array processors, *Quantum* **8**, 1281 (2024), arXiv:2306.03487 [quant-ph].
  - [51] T. Tomesh, P. Gokhale, V. Omole, G. S. Ravi, K. N. Smith, J. Viszlai, X.-C. Wu, N. Hardavellas, M. R. Martonosi, and F. T. Chong, Supermarq: A scalable quantum benchmark suite, in *2022 IEEE International Symposium on High-Performance Computer Architecture (HPCA)* (2022) pp. 587–603.
  - [52] R. Babbush, J. R. McClean, M. Newman, C. Gidney, S. Boixo, and H. Neven, Focus beyond quadratic speedups for error-corrected quantum advantage, *PRX Quantum* **2**, 010103 (2021).
  - [53] Z. Cai, R. Babbush, S. C. Benjamin, S. Endo, W. J. Hugrins, Y. Li, J. R. McClean, and T. E. O’Brien, Quantum error mitigation, *Reviews of Modern Physics* **95**, 045005 (2023).
  - [54] E. van den Berg, Z. K. Mineev, A. Kandala, and K. Temme, Probabilistic error cancellation with sparse Pauli-Lindblad models on noisy quantum processors, *Nature Physics* **19**, 1116 (2023).
  - [55] M. Sipser, *Introduction to the Theory of Computation*, 3rd ed. (Cengage Learning, 2012).
  - [56] R. Alur and P. Černý, Expressiveness of streaming string transducers, in *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (FSTTCS 2010)*, Leibniz International Proceedings in Informatics, Vol. 8 (2010) pp. 1–12.
  - [57] J. I. Munro and M. S. Paterson, Selection and sorting with limited storage, *Theoretical Computer Science* **12**, 315 (1980).
  - [58] N. Alon, Y. Matias, and M. Szegedy, The space complexity of approximating the frequency moments, *Journal of Computer and System Sciences* **58**, 137 (1999).